

Hardware Acceleration for Realtime Robotics



Brian Plancher
12/14/21



Hi I'm Brian!



Hi I'm Brian!

- I'm obsessed with my dog Alvin



Hi I'm Brian!

- I'm obsessed with my dog Alvin and my daughter Tess



Hi I'm Brian!

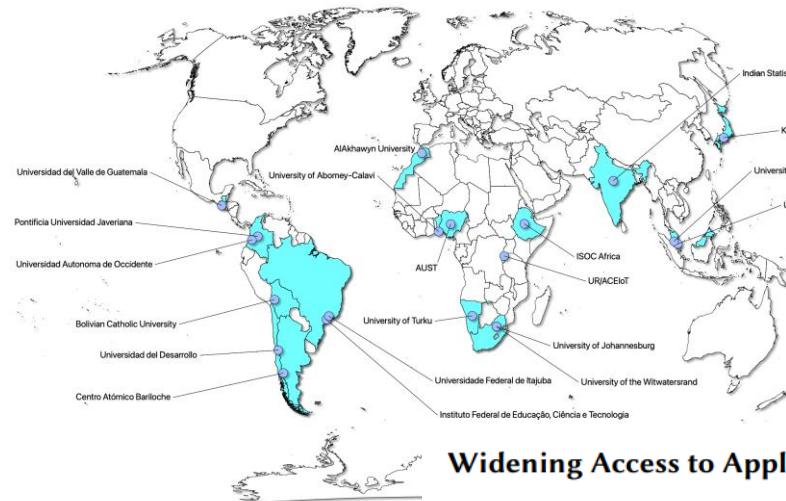
- I'm obsessed with my dog Alvin and my daughter Tess
- **I am passionate about teaching accessible, interdisciplinary hands-on, project-based courses**



Hi I'm Brian!

- I'm obsessed with my dog Alvin and my daughter Tess
- **I am passionate about teaching accessible, interdisciplinary hands-on, project-based courses**

<https://TinyMLedu.org>



Widening Access to Applied Machine Learning with TinyML

Vijay Janapa Reddi*, Brian Plancher*, Susan Kennedy*, Laurence Moroney†
Pete Warden† Anant Agarwal*‡ Colby Banbury* Massimo Banzi§ Matthew Bennett*
Benjamin Brown* Sharad Chitlangia¶ Radhika Ghosal* Sarah Grafman* Rupert Jaeger||
Srivatsan Krishnan* Maximilian Lam* Daniel Leiker|| Cara Mann* Mark Mazumder*
Dominic Pajak§ Dhilan Ramaprasad* J. Evan Smith* Matthew Stewart* Dustin Tingley*

edX Courses ▾ Programs & Degrees ▾ Schools & Partners ▾ What do you want to learn?

Catalog > Data Analysis & Statistics Courses > HarvardX's Tiny Machine Learning (TinyML)

HARVARD UNIVERSITY

Fundamentals of TinyML

Focusing on the basics of machine learning and embedded systems, such as smartphones, this course will introduce you to the "language" of TinyML.

Hi I'm Brian!

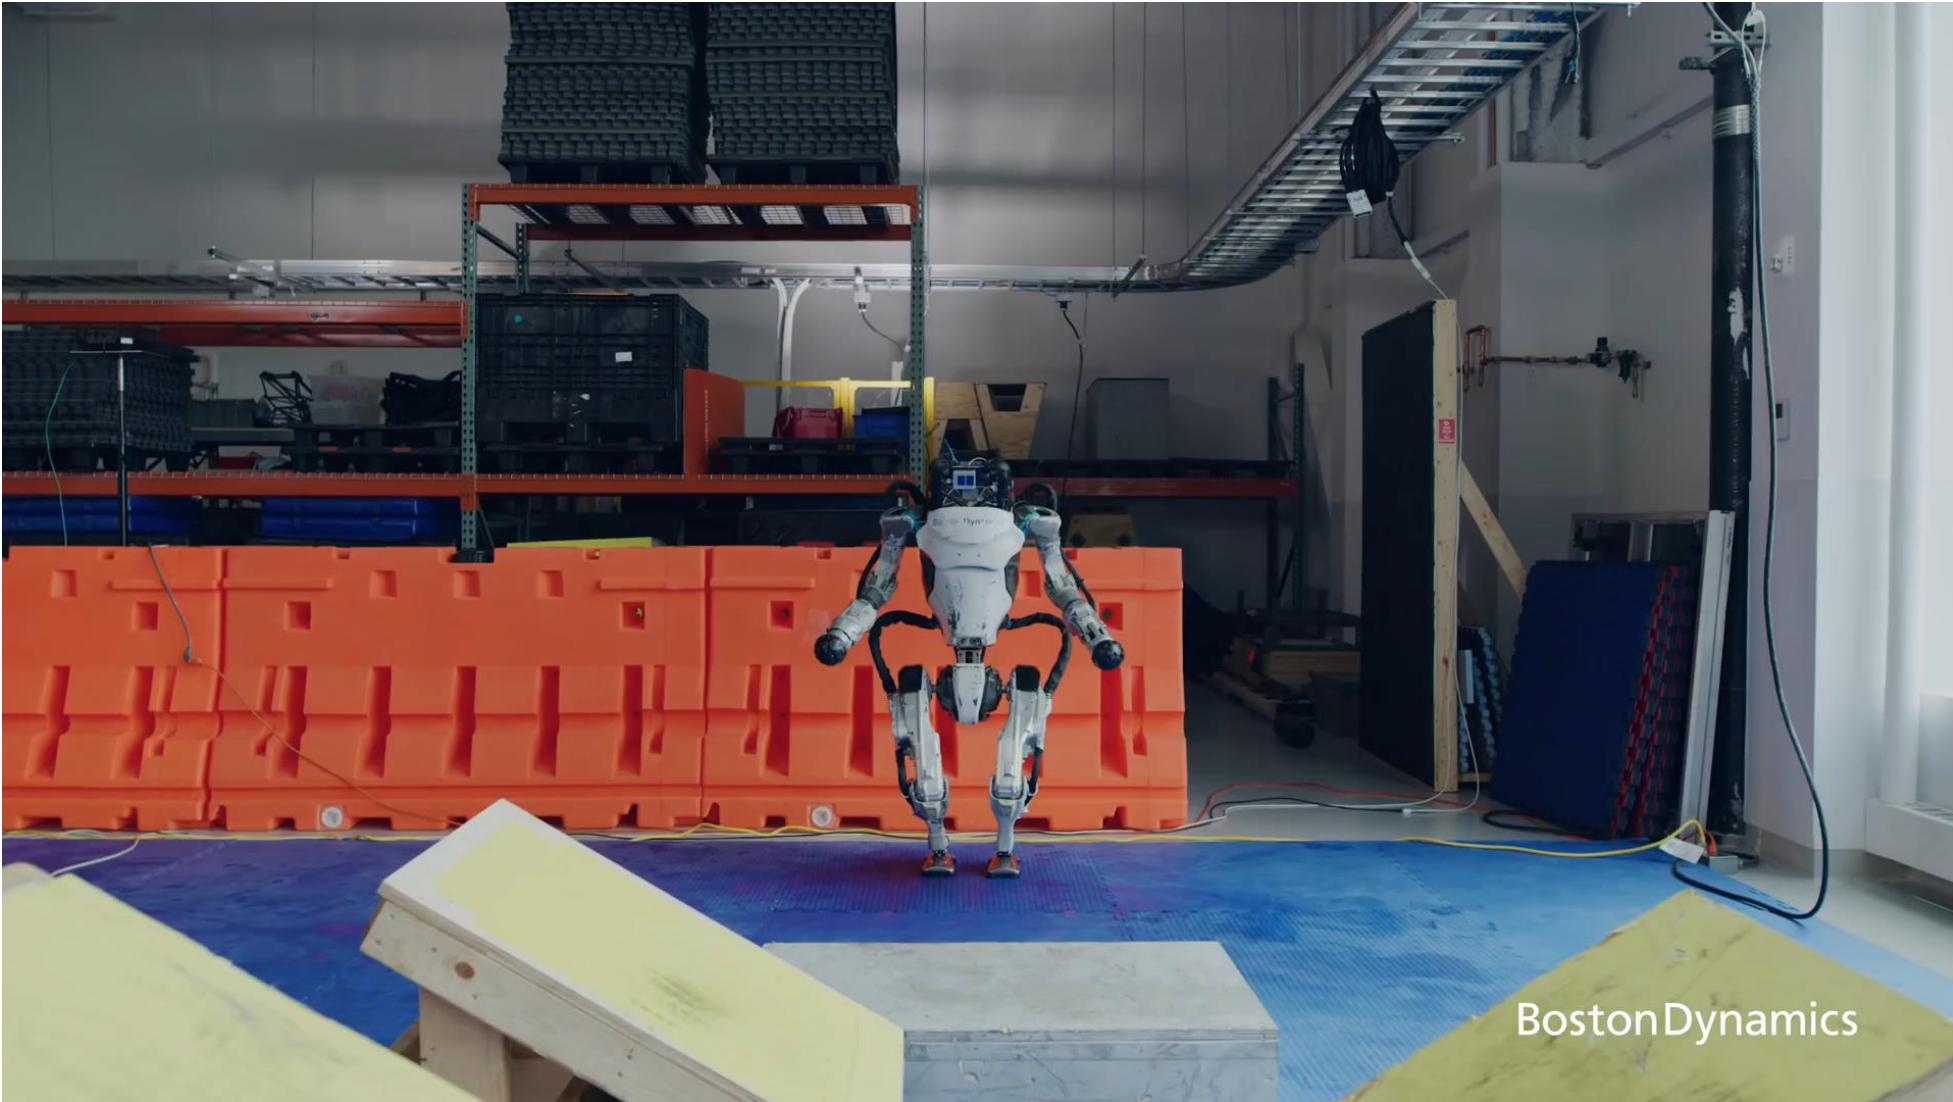
- I'm obsessed with my dog Alvin and my daughter Tess
- I am passionate about teaching accessible, interdisciplinary hands-on, project-based courses
- **My research focuses on developing open-source planning and control algorithms that enable robots to operate in the real world and help people!**



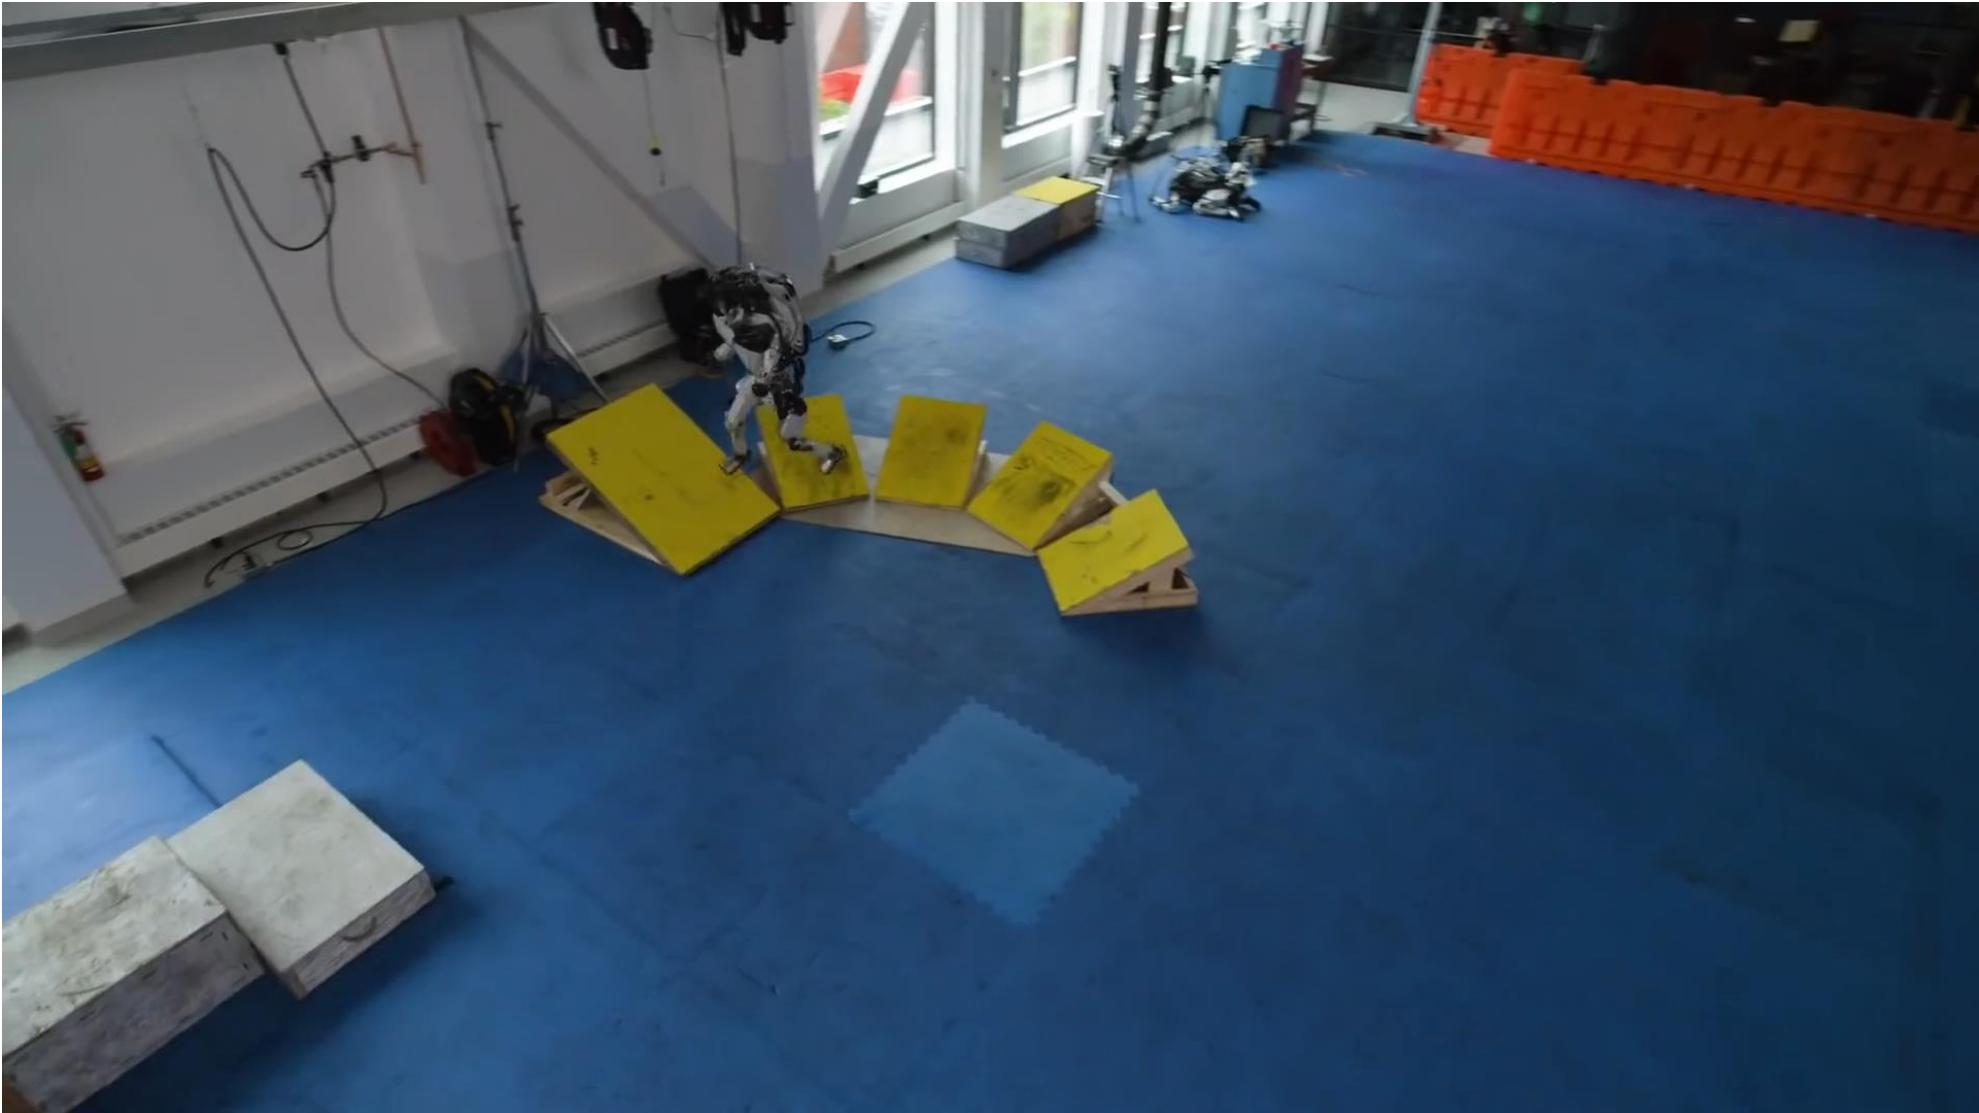
How about you?

1. How many people in the audience have a degree in computer science? (e.g., Faculty, Staff)
2. How many people have taken a handful of CS courses? (e.g., Junior or Senior CS Majors)
3. How about one or two CS courses?
4. How many people have no CS background?

Robots can do amazing things ...



... but they still have a long way to go!



How can we fix this?

How can we fix this?

Be Smarter = Better Algorithms

Think Faster = Hardware Acceleration

Learning Goals for Today

1. Learn some of the **language of robotics**
(and computer architecture)
2. Understand the important of **parallelism**
3. Gain practice in exploring the **process** and
opportunities of **hardware acceleration**

Hardware Acceleration for Realtime Robotics



A Quick Overview of (And My Approach To) Robotics

A Crash Course in (GPU) Architecture

A Case Study in Accelerating Rigid Body Dynamics

The Future of Accelerating Robotics

Hardware Acceleration for Realtime Robotics



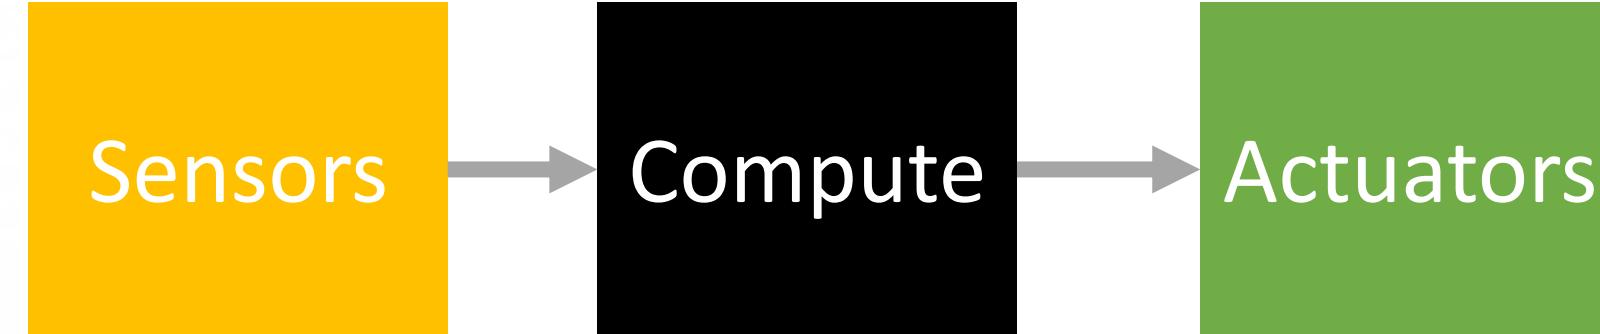
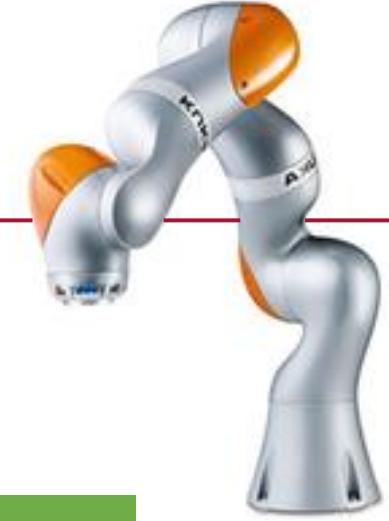
A Quick Overview of (And My Approach To) Robotics

A Crash Course in (GPU) Architecture

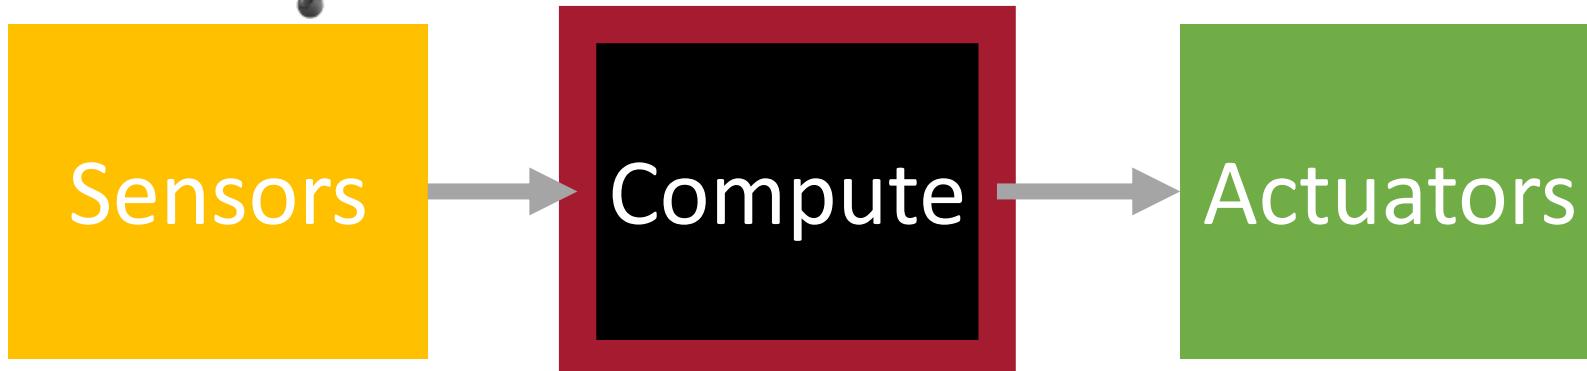
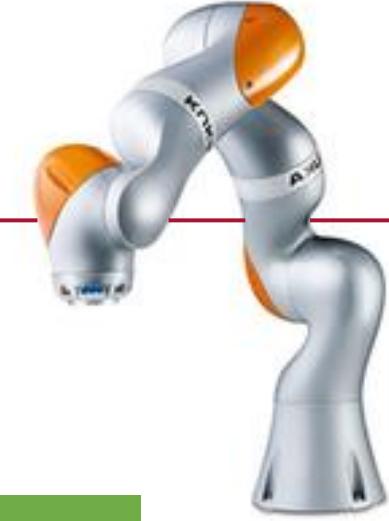
A Case Study in Accelerating Rigid Body Dynamics

The Future of Accelerating Robotics

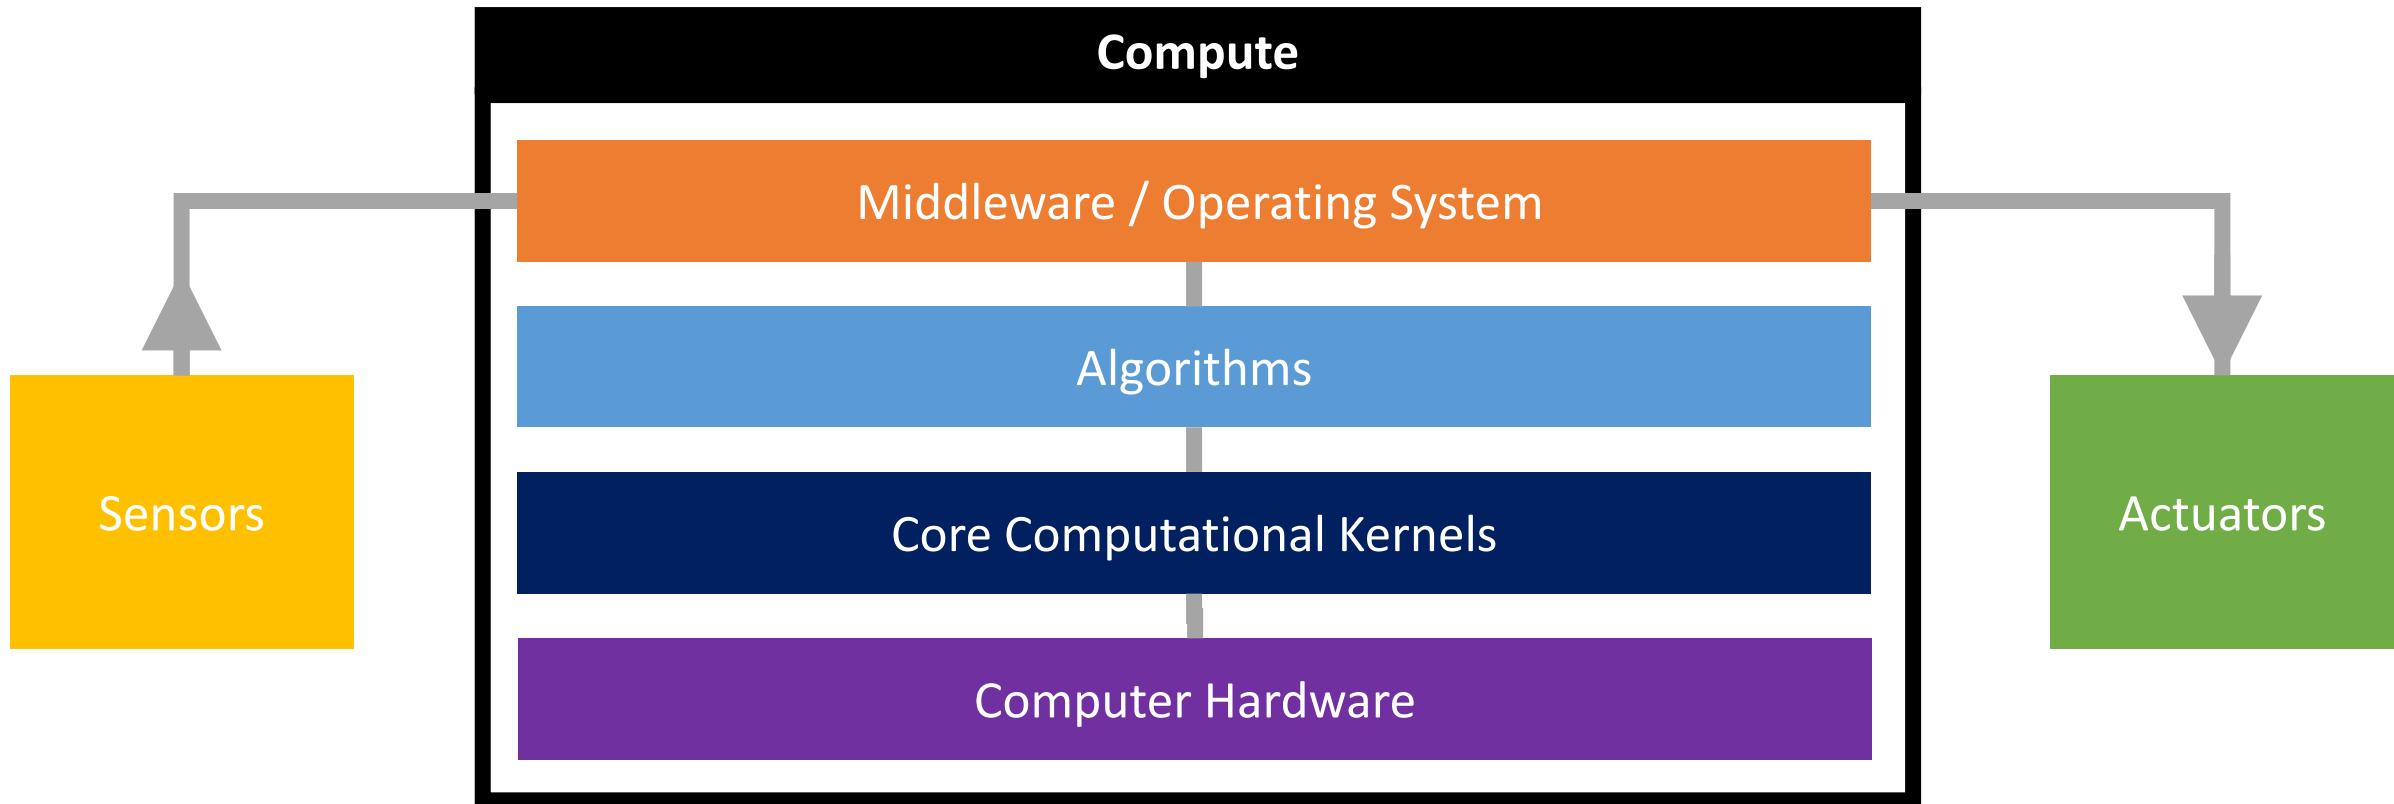
Robotics is a BIG space



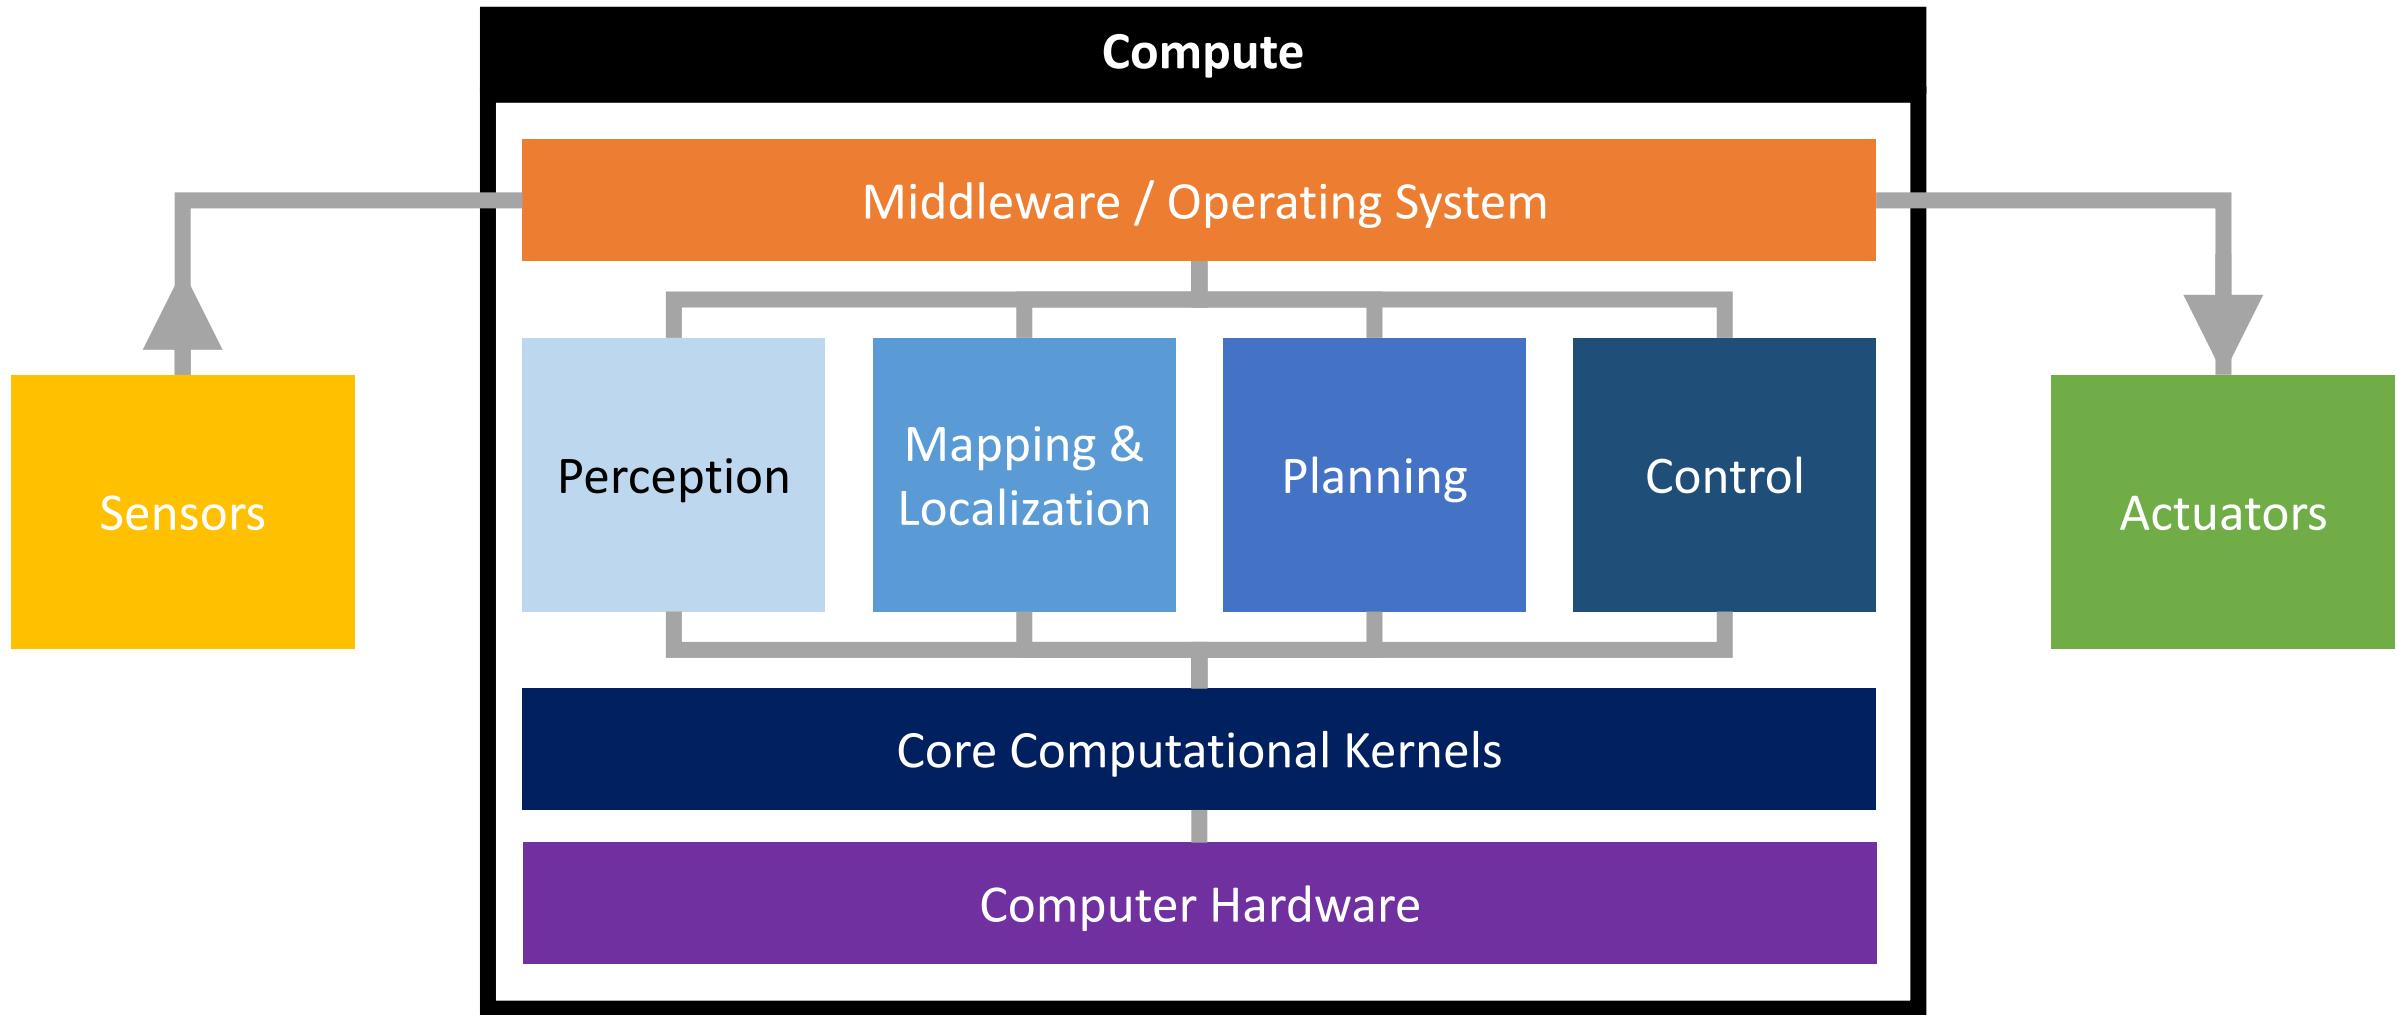
Robotics is a BIG space



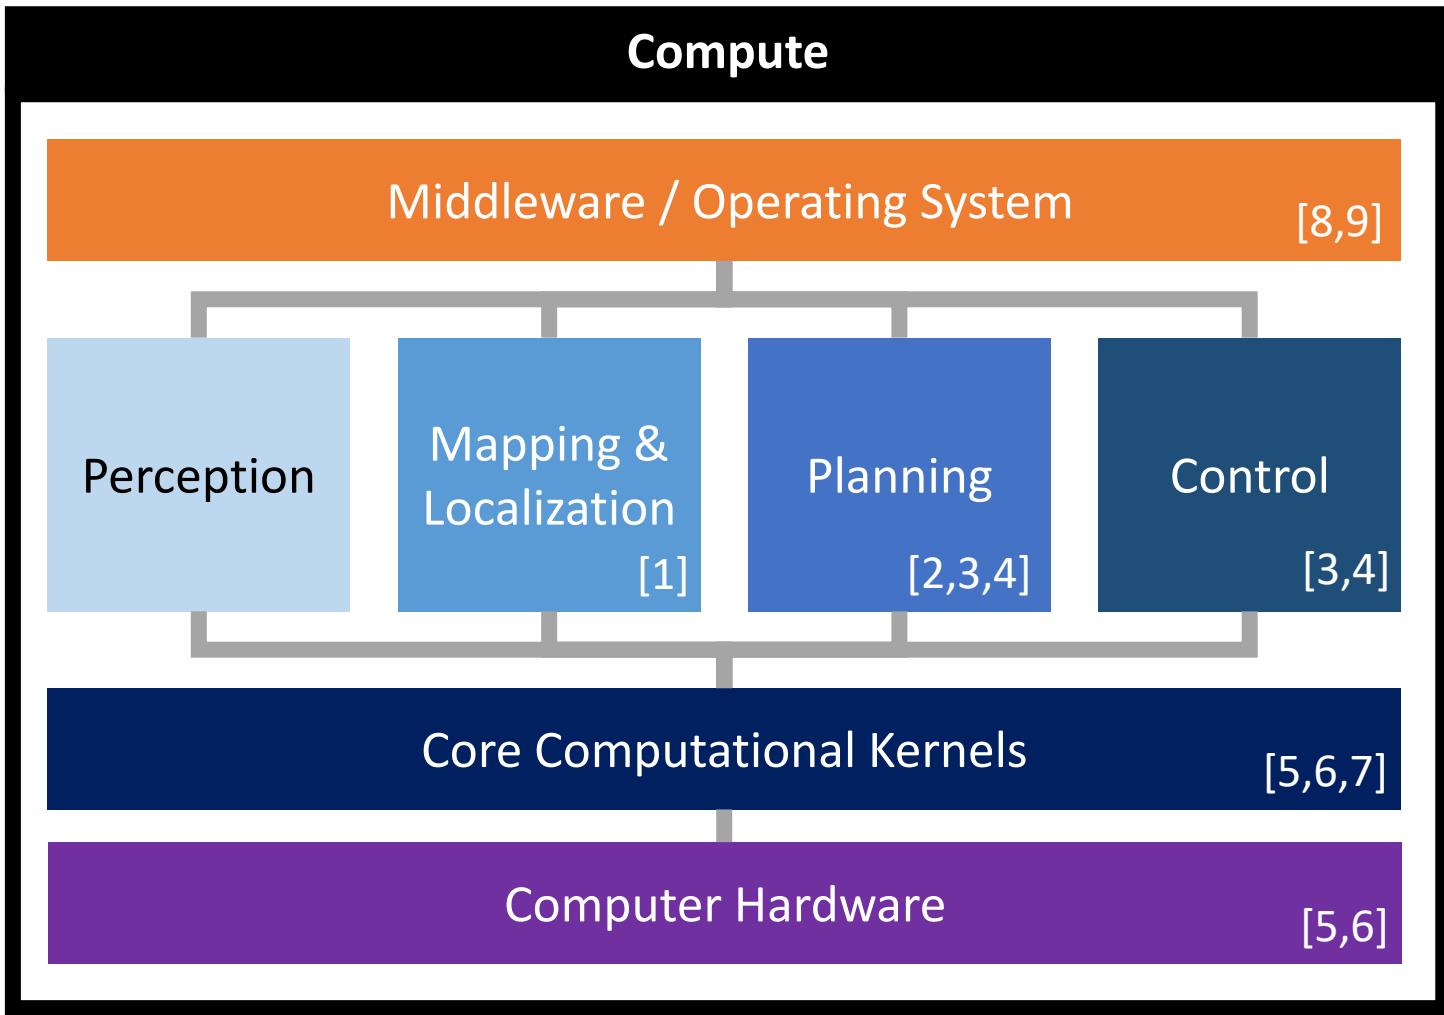
The Robotics Pipeline



The Robotics Pipeline



The Robotics Pipeline

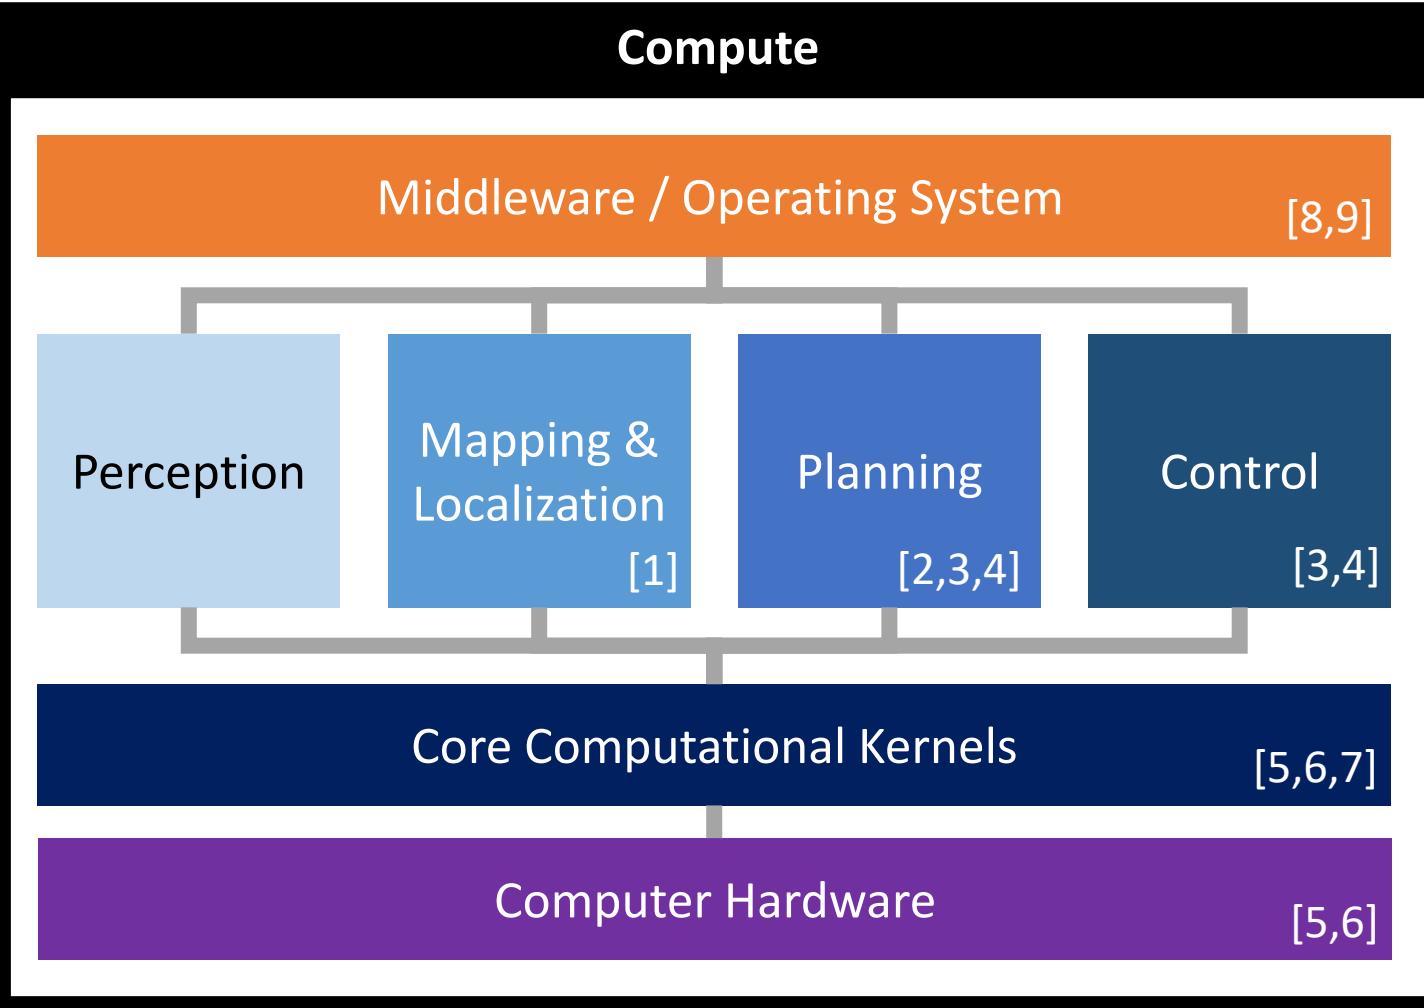


- [1] **B. Plancher**, C. D. Brumar, I. Brumar, L. Pentecost, S. Rama, D. Brooks. Application of Approximate Matrix Multiplication to Neural Networks and Distributed SLAM. HPEC 2019.
- [2] **B. Plancher**, Z. Manchester, S. Kuindersma. Constrained Unscented Dynamic Programming. IROS 2017.
- [3] **B. Plancher**, S. Kuindersma. A Performance Analysis of Parallel Differential Dynamic Programming on a GPU. WAFR 2018.
- [4] **B. Plancher**, S. Kuindersma. Realtime Model Predictive Control using Parallel DDP on a GPU. In Online Optimal Control Workshop at ICRA 2019.
- [5] **B. Plancher**, S. M. Neuman, T. Bourgeat, S. Kuindersma, S. Devadas, V. Janapa Reddi. Accelerating Robot Dynamics Gradients on a CPU, GPU, and FPGA. RA-L and ICRA 2021.
- [6] S. M. Neuman, **B. Plancher**, T. Bourgeat, T. Tambe, S. Devadas, V. Janapa Reddi. Robomorphic Computing: A Design Methodology for Domain-Specific Accelerators Parameterized by Robot Morphology. ASPLOS 2021.
- [7] [Under Review] **B. Plancher**, S. M. Neuman, R. Ghosal, S. Kuindersma, V. Janapa Reddi. GRiD: GPU-Accelerated Rigid Body Dynamics with Analytical Gradients. ICRA, 2022.
- [8] B. Boroujerdian, R. Ghosal, J. Cruz, **B. Plancher**, V. Janapa Reddi. RoboRun: A Robot Runtime to Exploit Spatial Heterogeneity. DAC 2021.
- [9] [Under Review] B. Boroujerdian, H. Genc, S. Krishnan, P. Bardienus, B. Duisterhof, **B. Plancher**, K. Mansoorshahi, M. Almeida, A. Faust, V. Janapa Reddi. The Role of Compute in Autonomous Aerial Vehicles. TOCS, 2022.
- [10] **B. Plancher**, V. Janapa Reddi. TinyMLedu: The Tiny Machine Learning Open Education Initiative. SIGCSE 2022.
- [11] V. Janapa Reddi, **B. Plancher**, S. Kennedy, L. Moroney, P. Warden, L. Suzuki, A. Agarwal, C. Banbury, M. Banzi, M. Bennett, B. Brown, S. Chitlangia, R. Ghosal, S. Grafman, R. Jaeger, S. Krishnan, M. Lam, D. Leiker, C. Mann, M. Mazumder, D. Pajak , D. Ramaprasad , J. E. Smith, M. Stewart, D. Tingley. Widening Access to Applied Machine Learning with TinyML. HDSR 2022.
- [12] S. Karaman, A. Anders, M. Boulet, J. Connor, K. Gregson, W. Guerra, O. Guldner, M. Mohamoud, **B. Plancher**, R. Shin, J. Vivilecchia. Project-based, collaborative, algorithmic robotics for high school students: Programming self-driving race cars at MIT. ISEC 2017.

The Robotics Pipeline

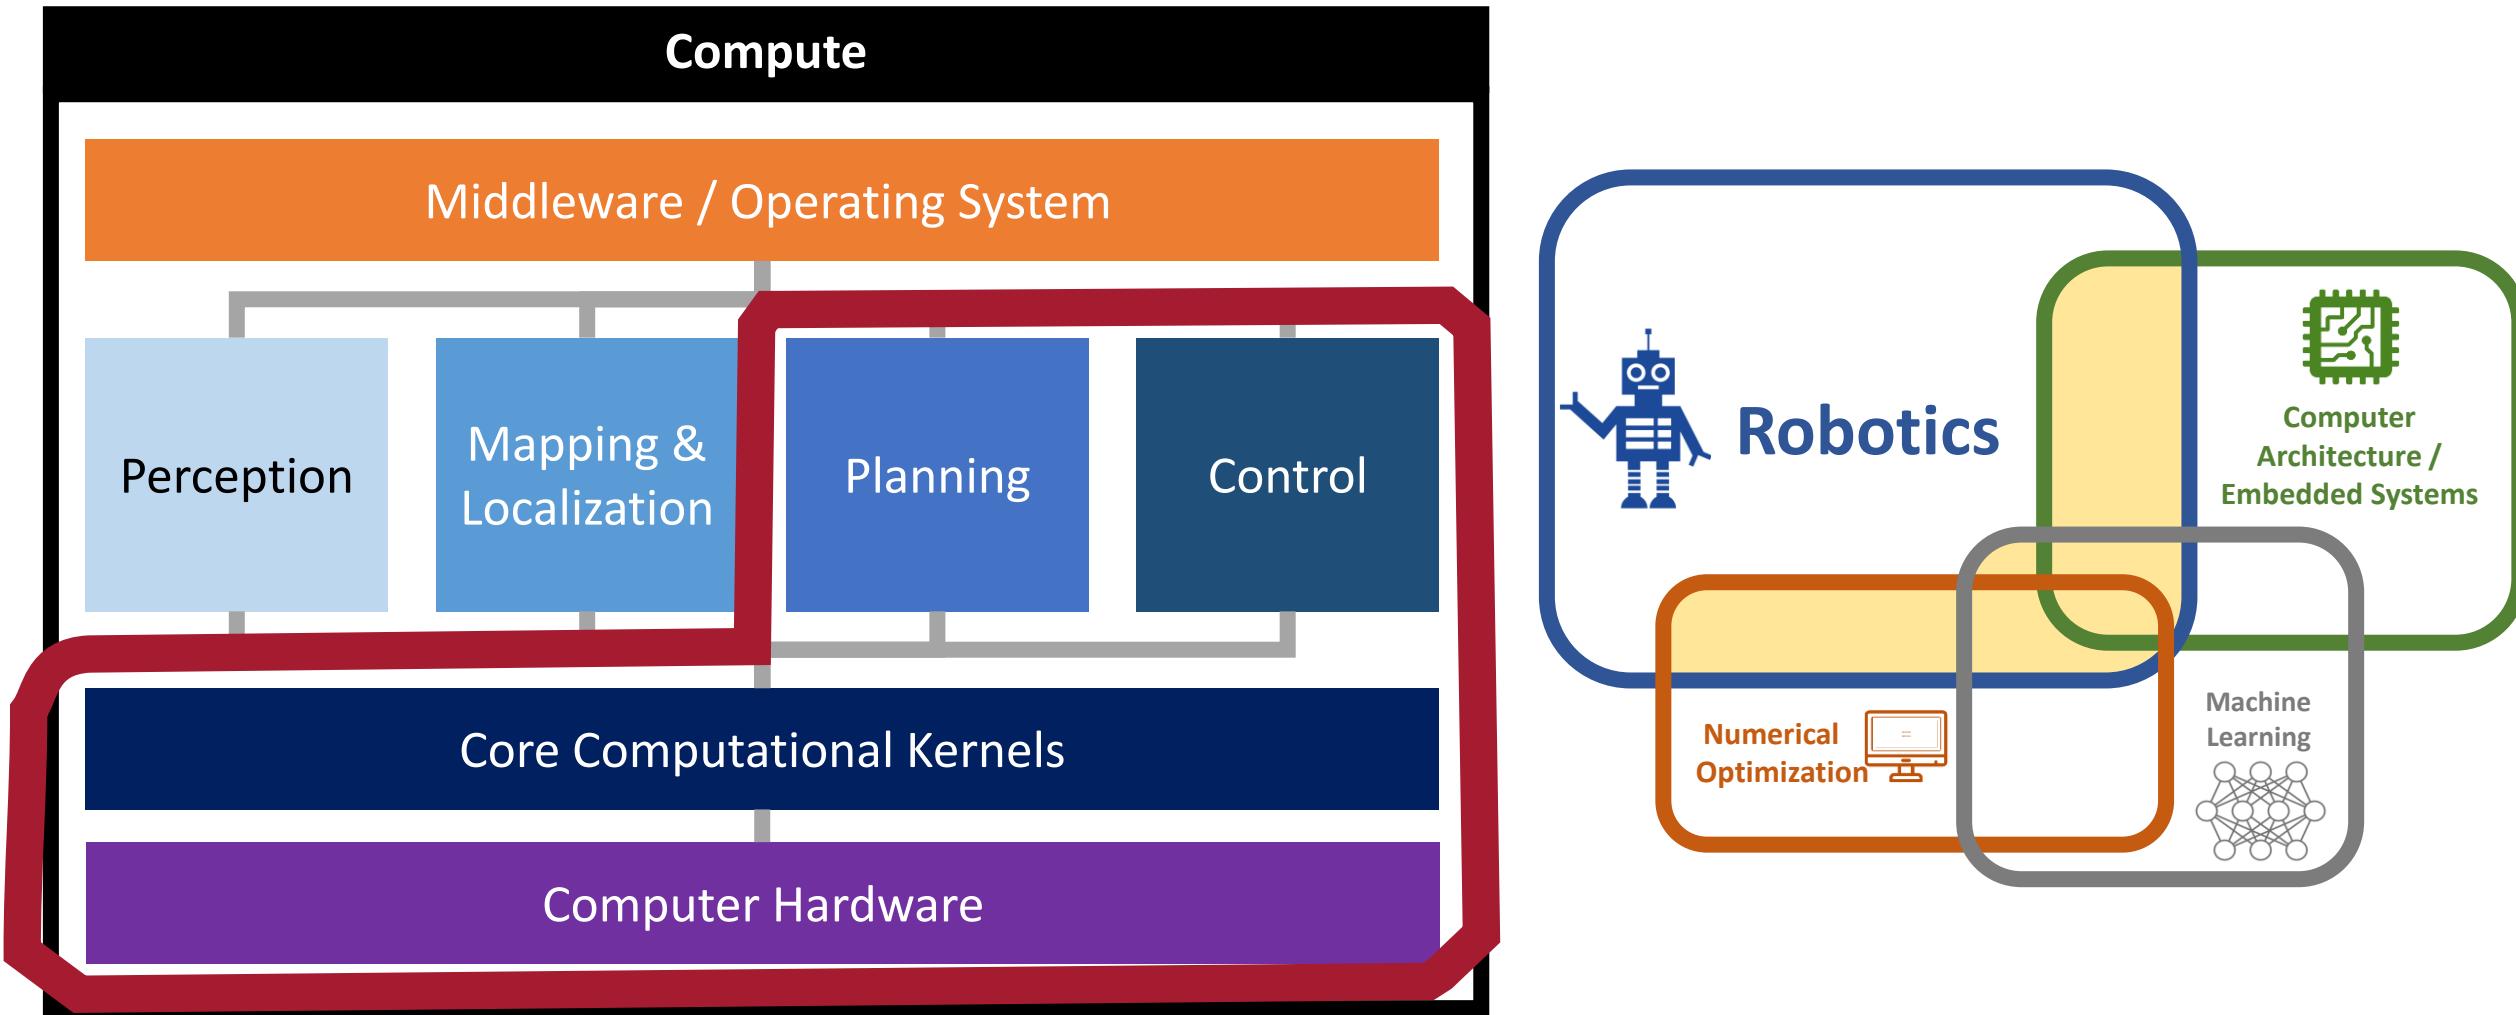


CS Education
[10,11,12]

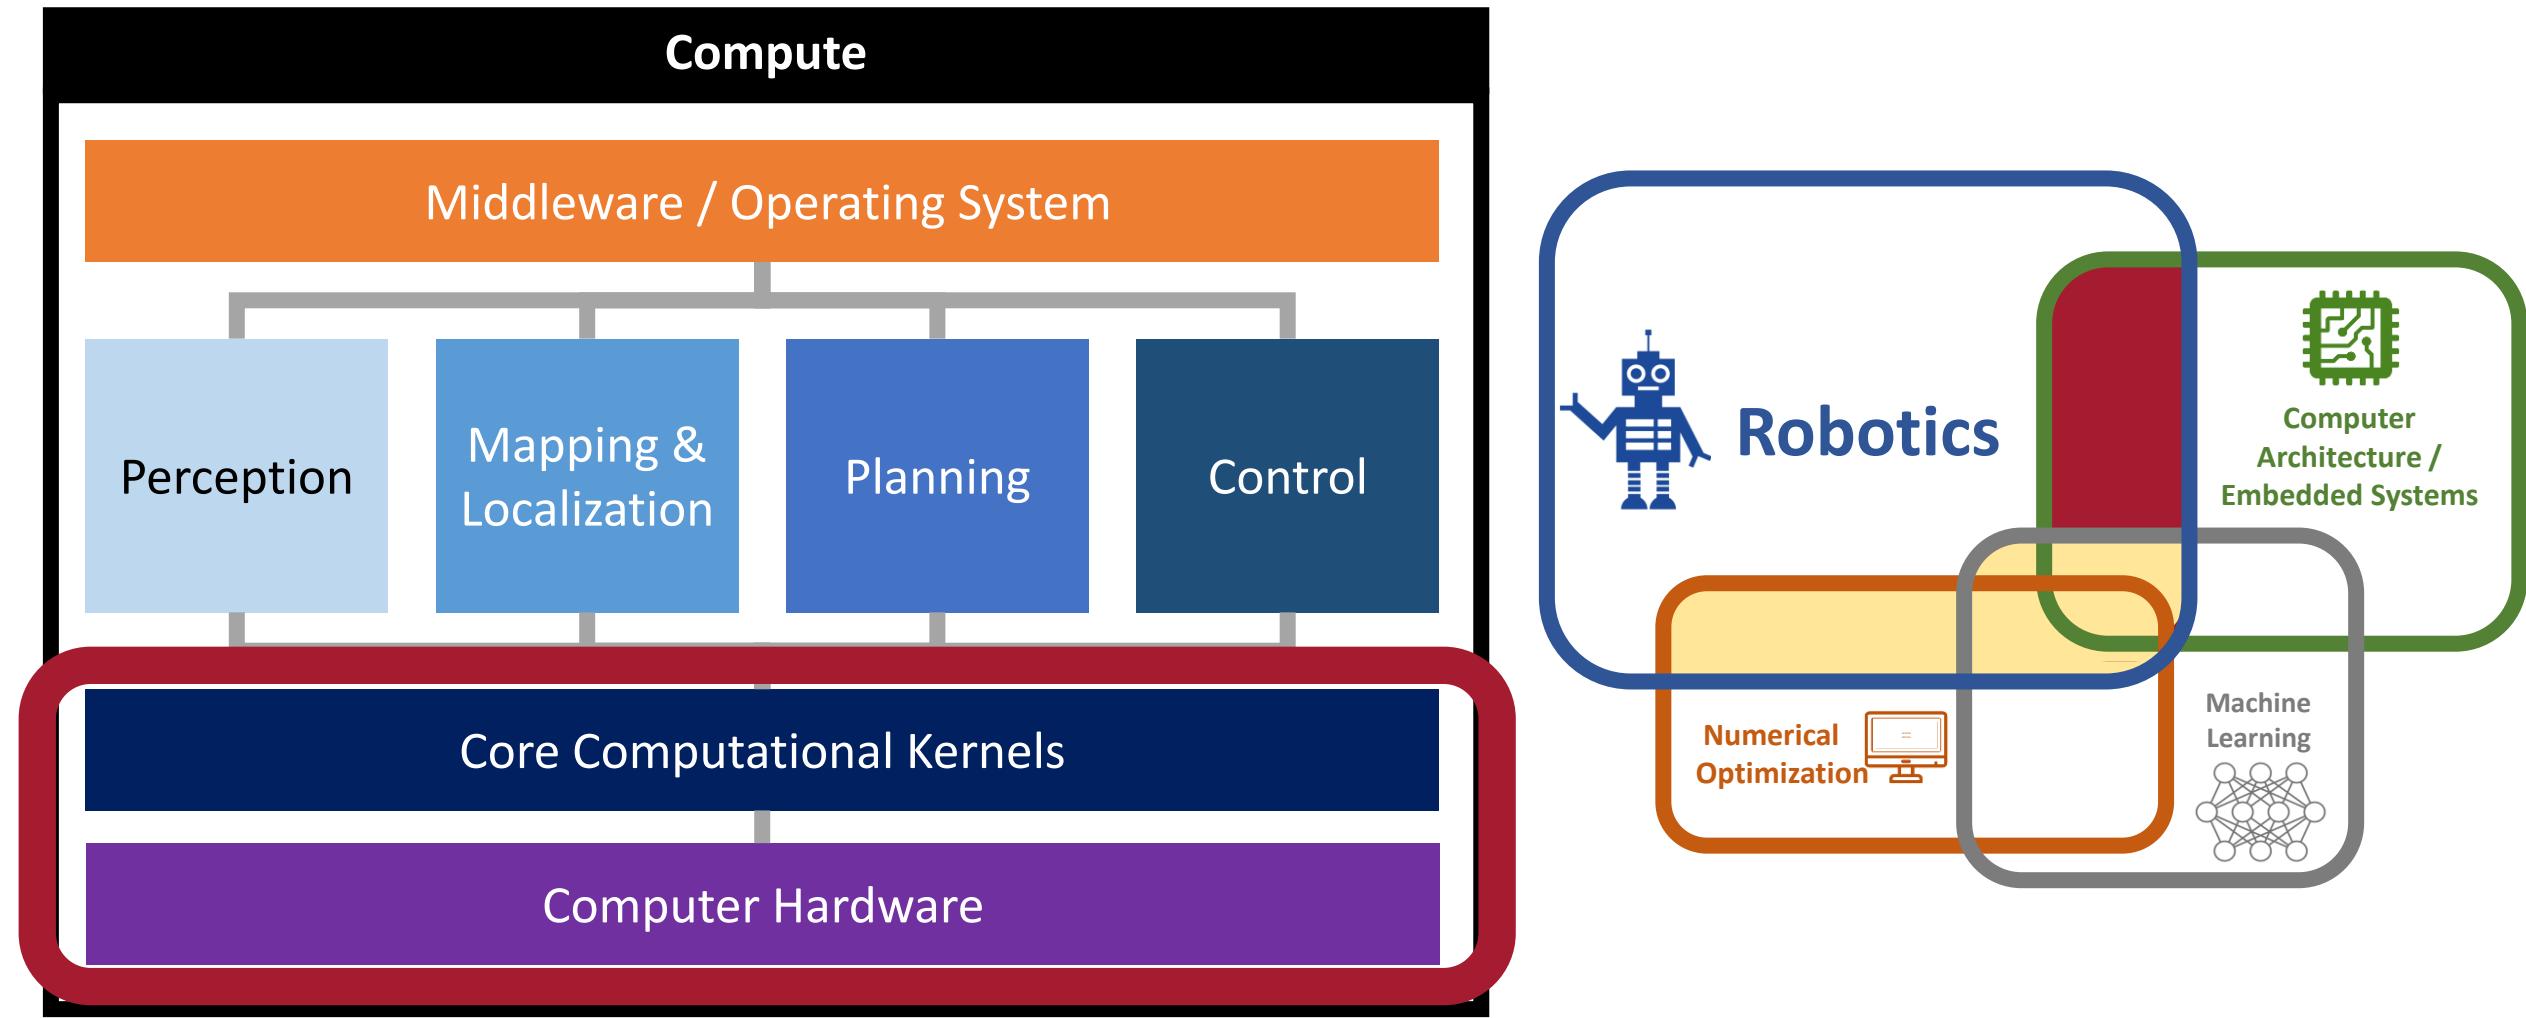


- [1] **B. Plancher**, C. D. Brumar, I. Brumar, L. Pentecost, S. Rama, D. Brooks. Application of Approximate Matrix Multiplication to Neural Networks and Distributed SLAM. HPEC 2019.
- [2] **B. Plancher**, Z. Manchester, S. Kuindersma. Constrained Unscented Dynamic Programming. IROS 2017.
- [3] **B. Plancher**, S. Kuindersma. A Performance Analysis of Parallel Differential Dynamic Programming on a GPU. WAFR 2018.
- [4] **B. Plancher**, S. Kuindersma. Realtime Model Predictive Control using Parallel DDP on a GPU. In Online Optimal Control Workshop at ICRA 2019.
- [5] **B. Plancher**, S. M. Neuman, T. Bourgeat, S. Kuindersma, S. Devadas, V. Janapa Reddi. Accelerating Robot Dynamics Gradients on a CPU, GPU, and FPGA. RA-L and ICRA 2021.
- [6] S. M. Neuman, **B. Plancher**, T. Bourgeat, T. Tambe, S. Devadas, V. Janapa Reddi. Robomorphic Computing: A Design Methodology for Domain-Specific Accelerators Parameterized by Robot Morphology. ASPLOS 2021.
- [7] [Under Review] **B. Plancher**, S. M. Neuman, R. Ghosal, S. Kuindersma, V. Janapa Reddi. GRiD: GPU-Accelerated Rigid Body Dynamics with Analytical Gradients. ICRA, 2022.
- [8] B. Boroujerdian, R. Ghosal, J. Cruz, **B. Plancher**, V. Janapa Reddi. RoboRun: A Robot Runtime to Exploit Spatial Heterogeneity. DAC 2021.
- [9] [Under Review] B. Boroujerdian, H. Genc, S. Krishnan, P. Bardienus, B. Duisterhof, **B. Plancher**, K. Mansoorshahi, M. Almeida, A. Faust, V. Janapa Reddi. The Role of Compute in Autonomous Aerial Vehicles. TOCS, 2022.
- [10] **B. Plancher**, V. Janapa Reddi. TinyMLedu: The Tiny Machine Learning Open Education Initiative. SIGCSE 2022.
- [11] V. Janapa Reddi, **B. Plancher**, S. Kennedy, L. Moroney, P. Warden, L. Suzuki, A. Agarwal, C. Banbury, M. Banzi, M. Bennett, B. Brown, S. Chitlangia, R. Ghosal, S. Grafman, R. Jaeger, S. Krishnan, M. Lam, D. Leiker, C. Mann, M. Mazumder, D. Pajak, D. Ramaprasad, J. E. Smith, M. Stewart, D. Tingley. Widening Access to Applied Machine Learning with TinyML. HDSR 2022.
- [12] S. Karaman, A. Anders, M. Boulet, J. Connor, K. Gregson, W. Guerra, O. Guldner, M. Mohamoud, **B. Plancher**, R. Shin, J. Vivilecchia. Project-based, collaborative, algorithmic robotics for high school students: Programming self-driving race cars at MIT. ISEC 2017.

I work at the intersection of robotics and adjacent fields



I work at the intersection of robotics and adjacent fields



Hardware Acceleration for Realtime Robotics



A Quick Overview of (And My Approach To) Robotics

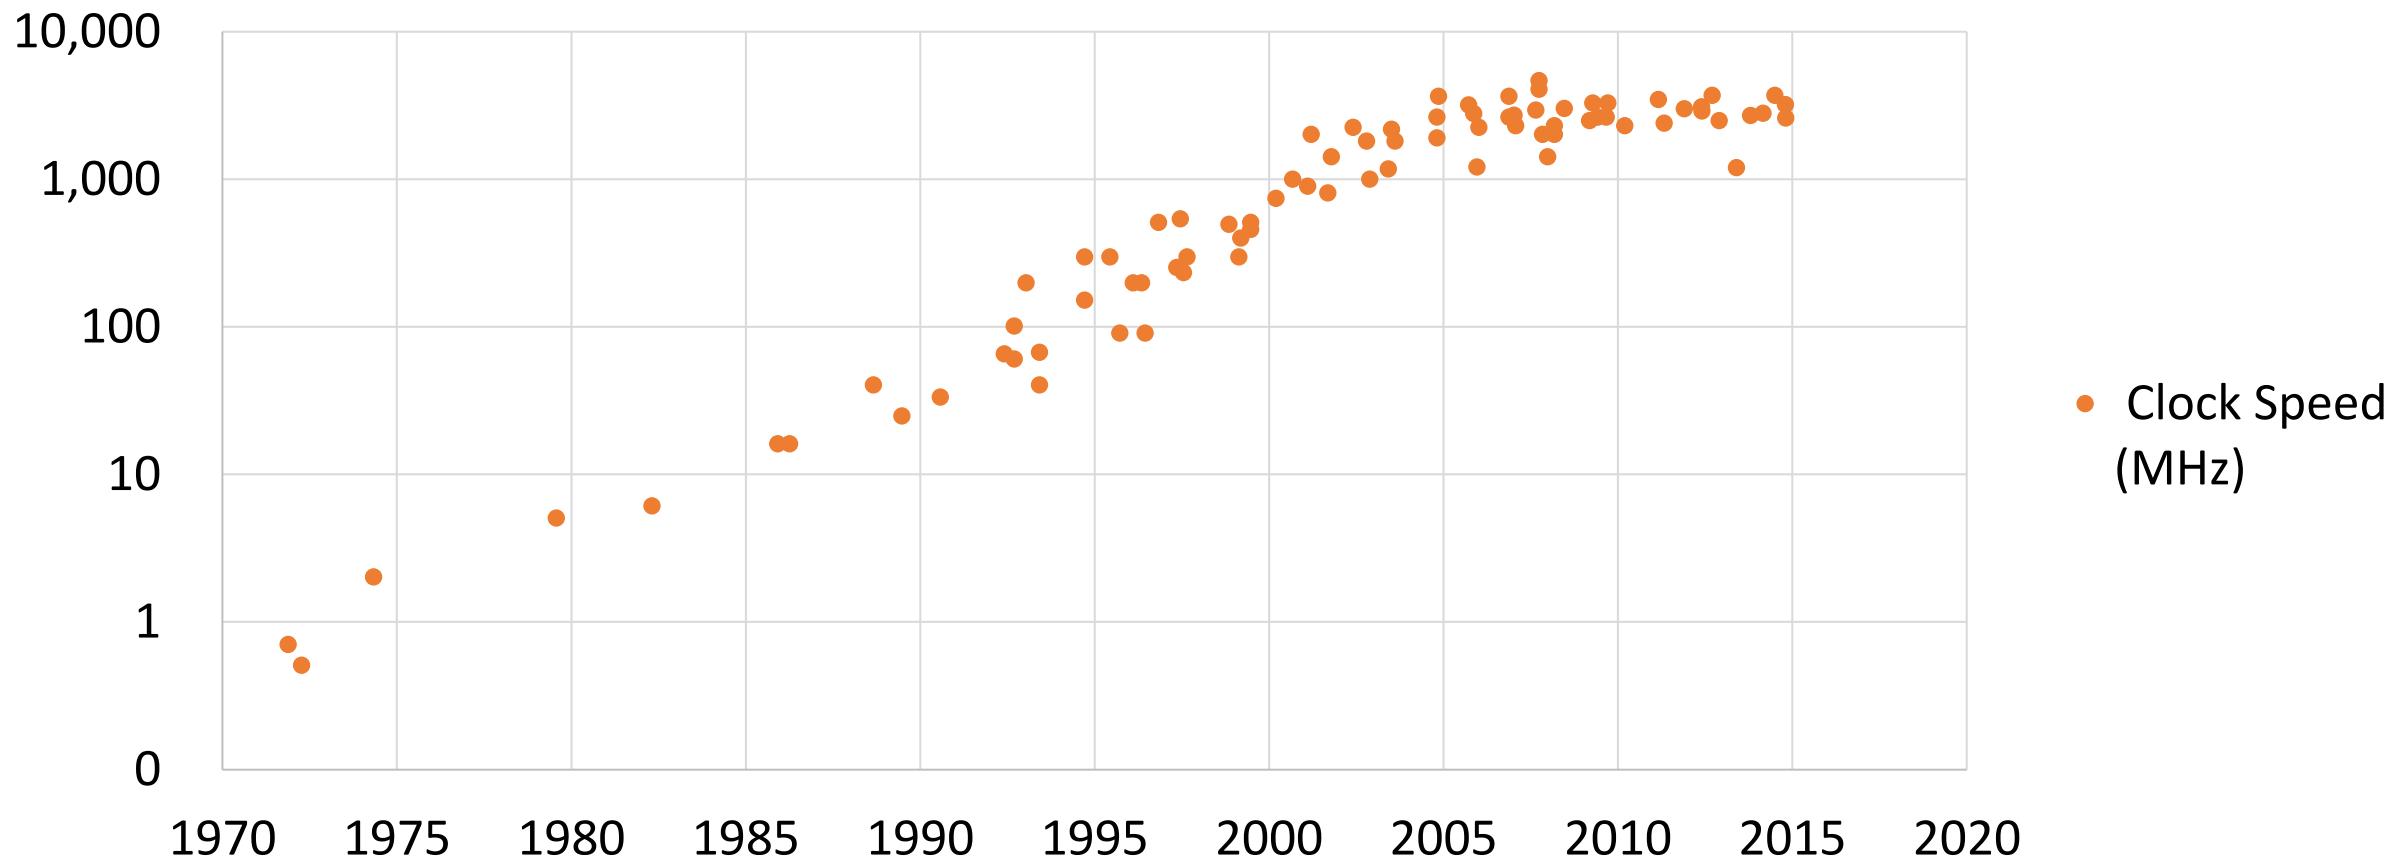
A Crash Course in (GPU) Architecture

A Case Study in Accelerating Rigid Body Dynamics

The Future of Accelerating Robotics

CPUs are not getting faster...

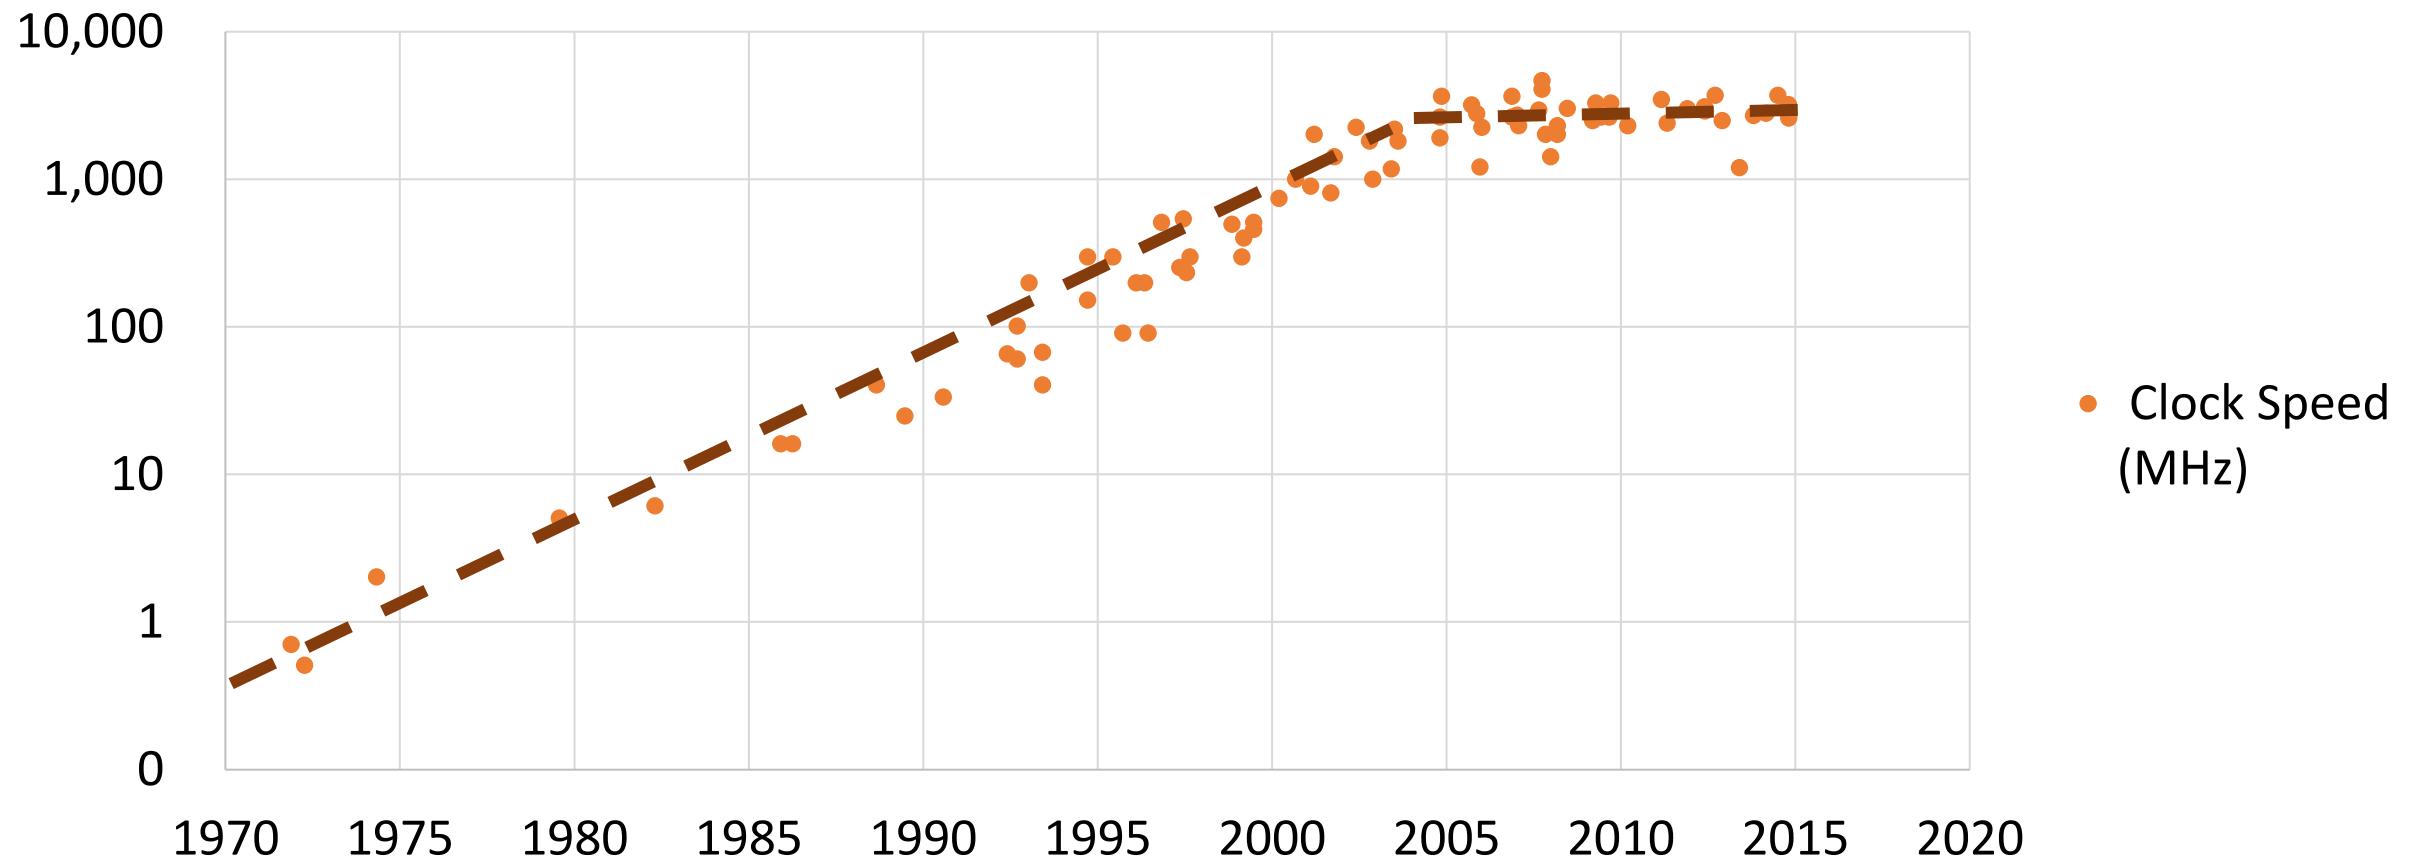
48 Years of Processor Trends



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2019 by K. Rupp <https://github.com/karlrupp/microprocessor-trend-data>

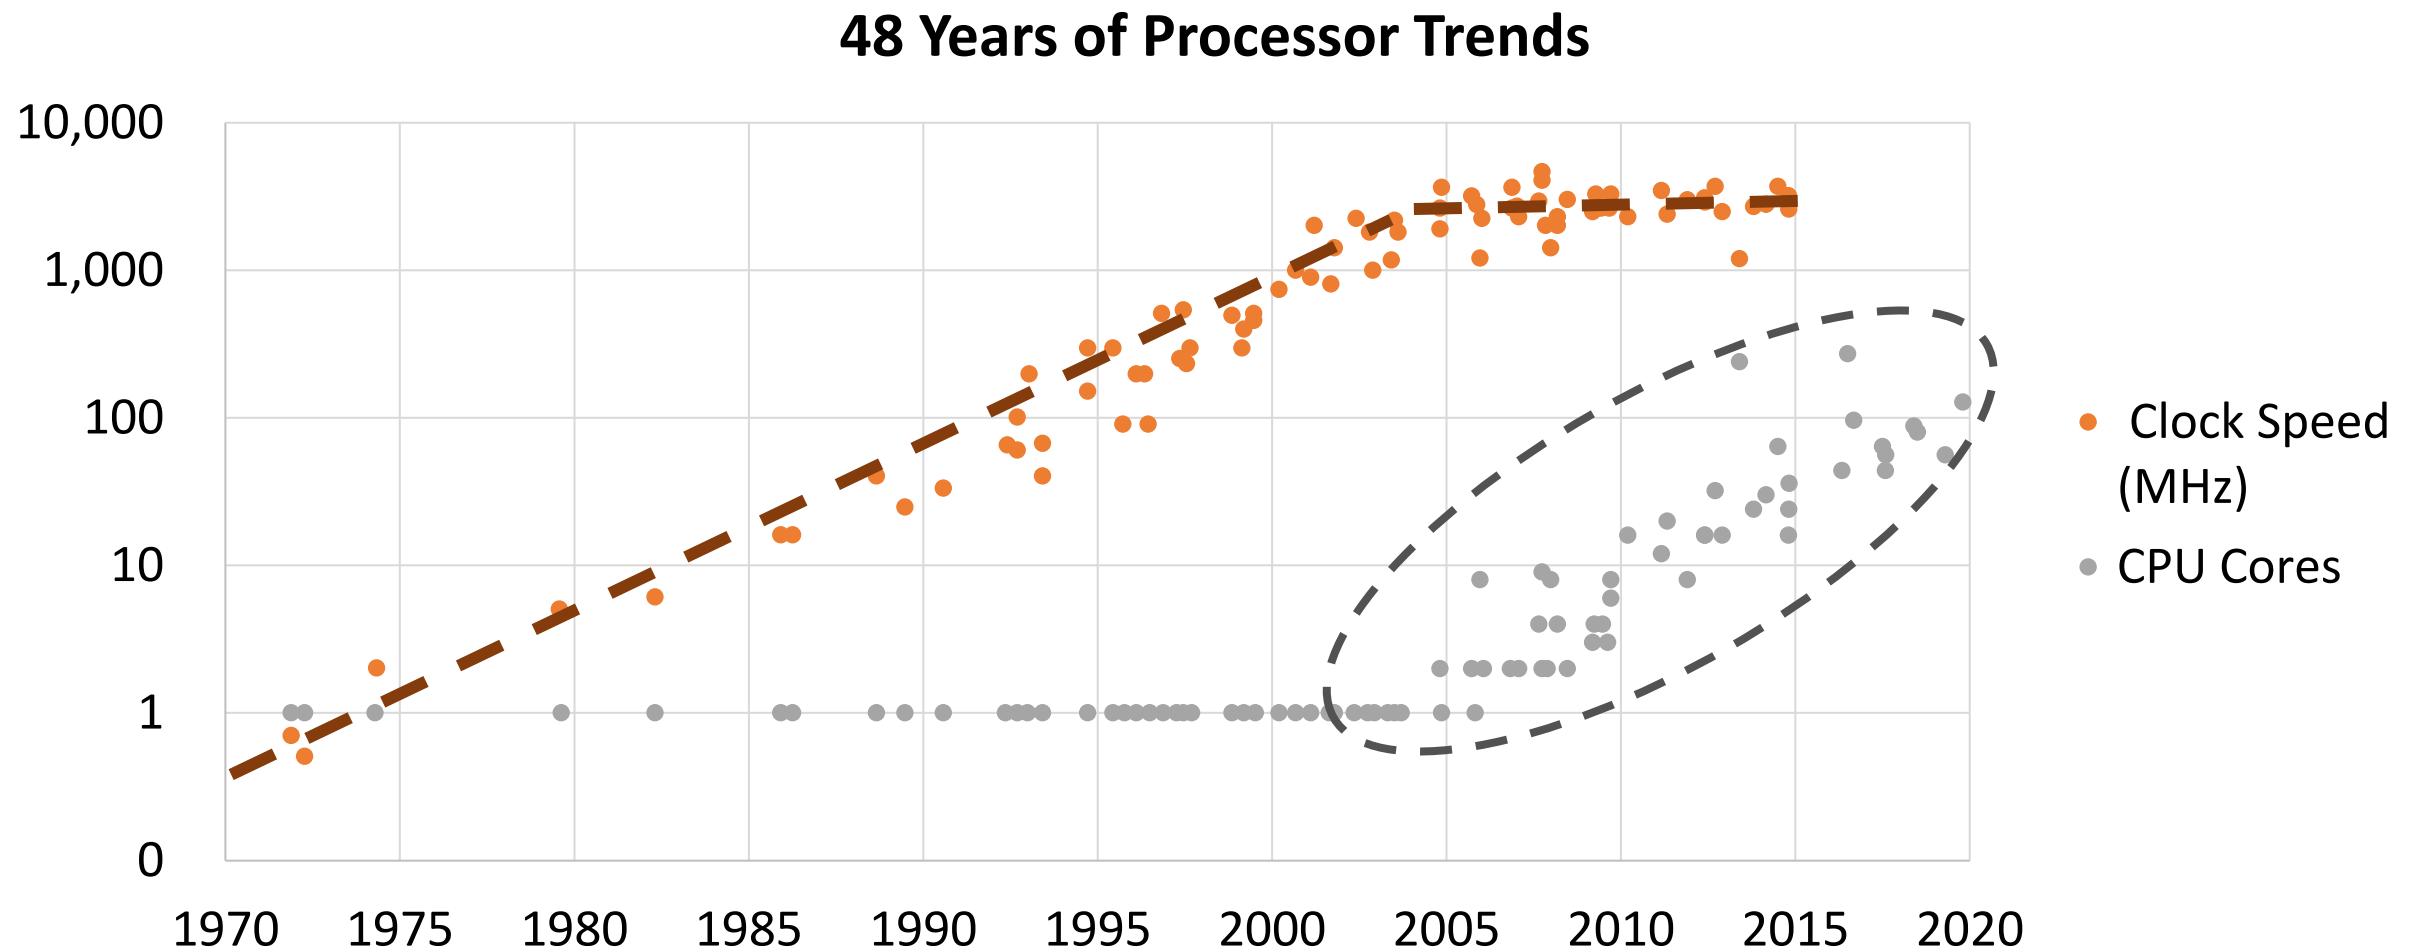
CPUs are not getting faster...

48 Years of Processor Trends



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2019 by K. Rupp <https://github.com/karlrupp/microprocessor-trend-data>

... instead, we need to focus on Parallelism



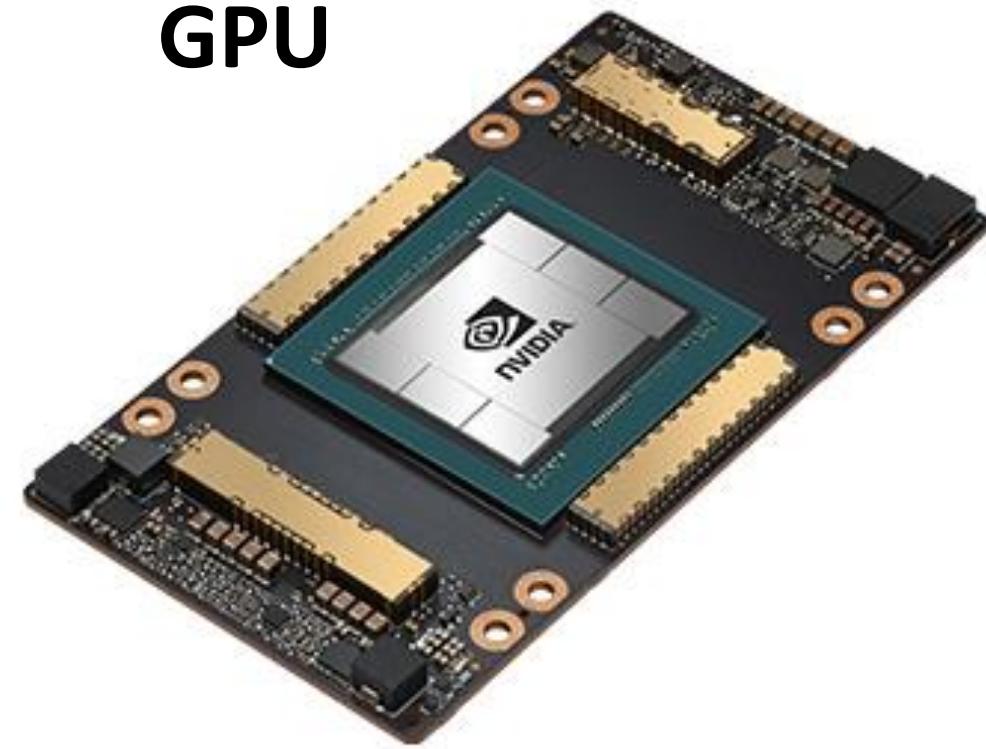
Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2019 by K. Rupp <https://github.com/karlrupp/microprocessor-trend-data>

... instead, we need to focus on **Parallelism**

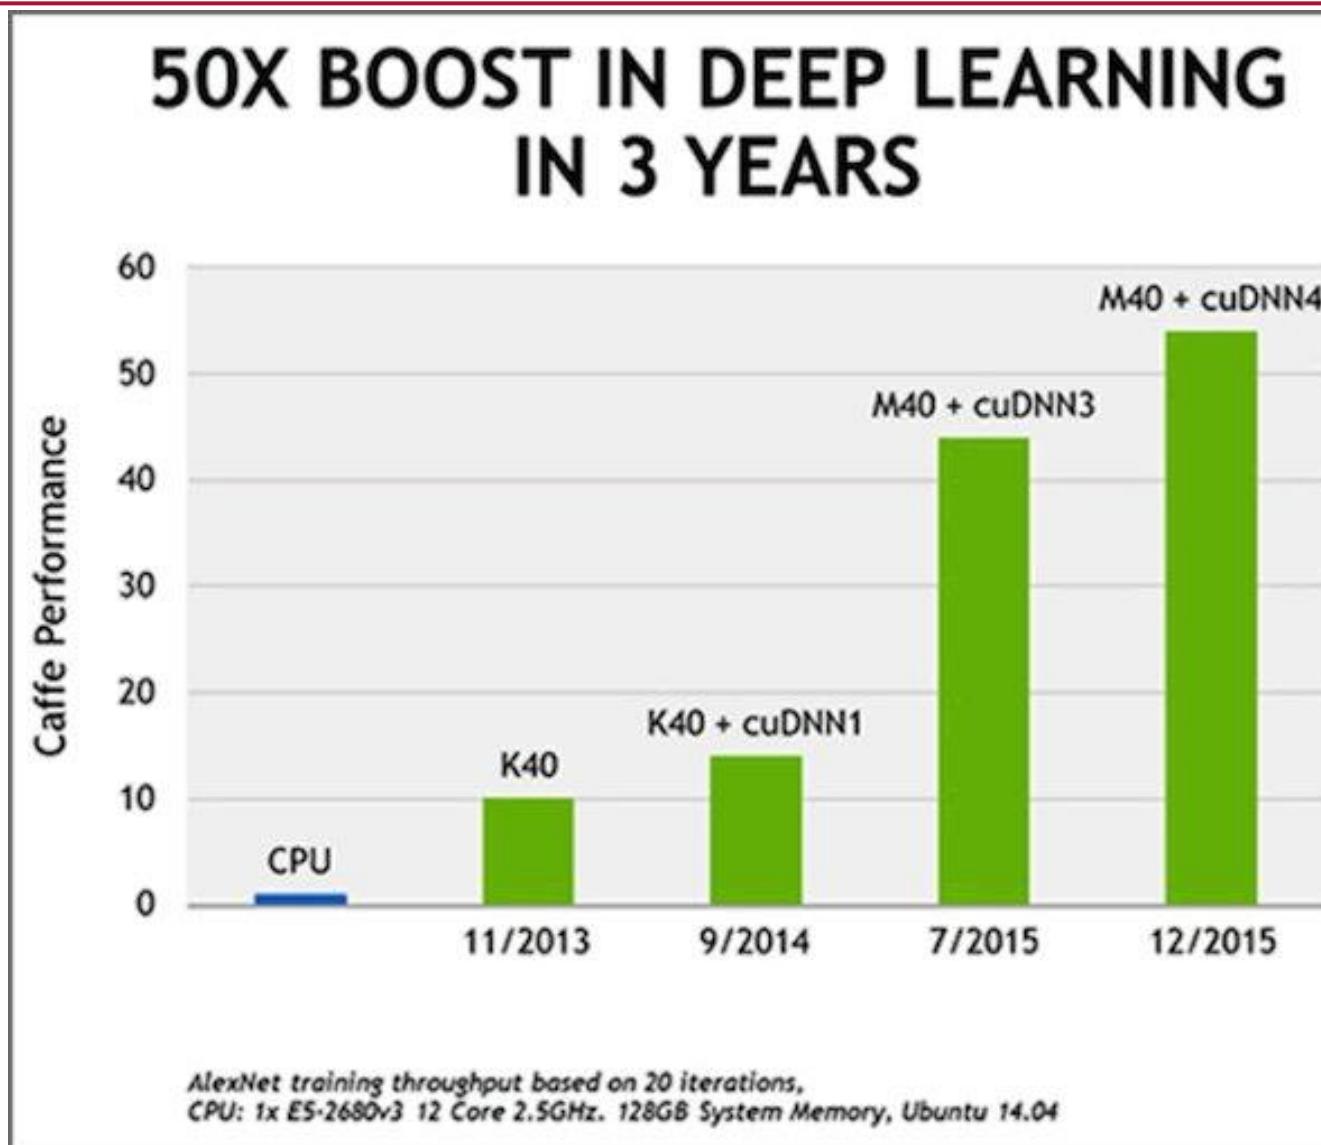
CPU



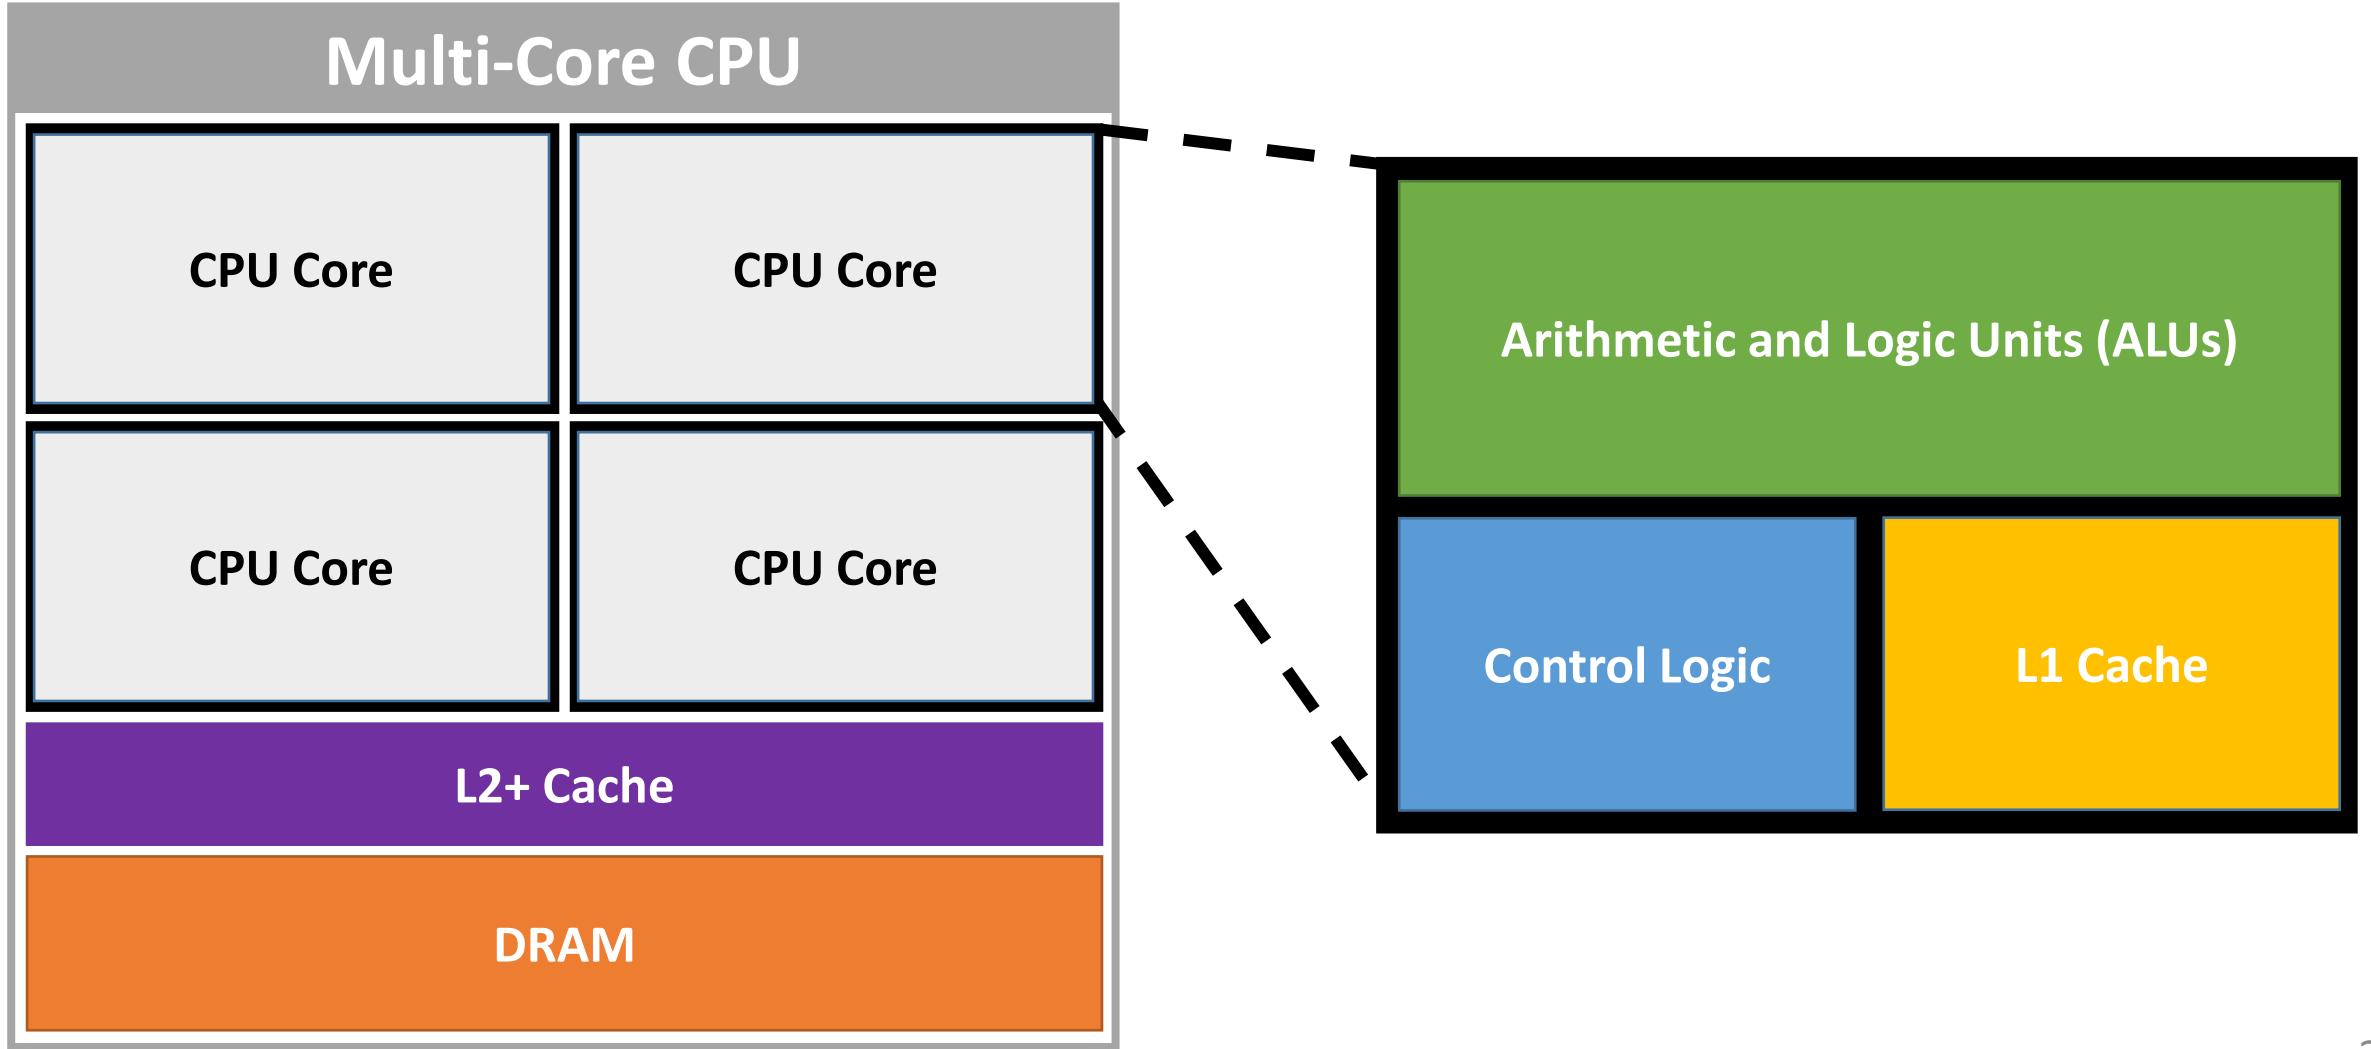
GPU



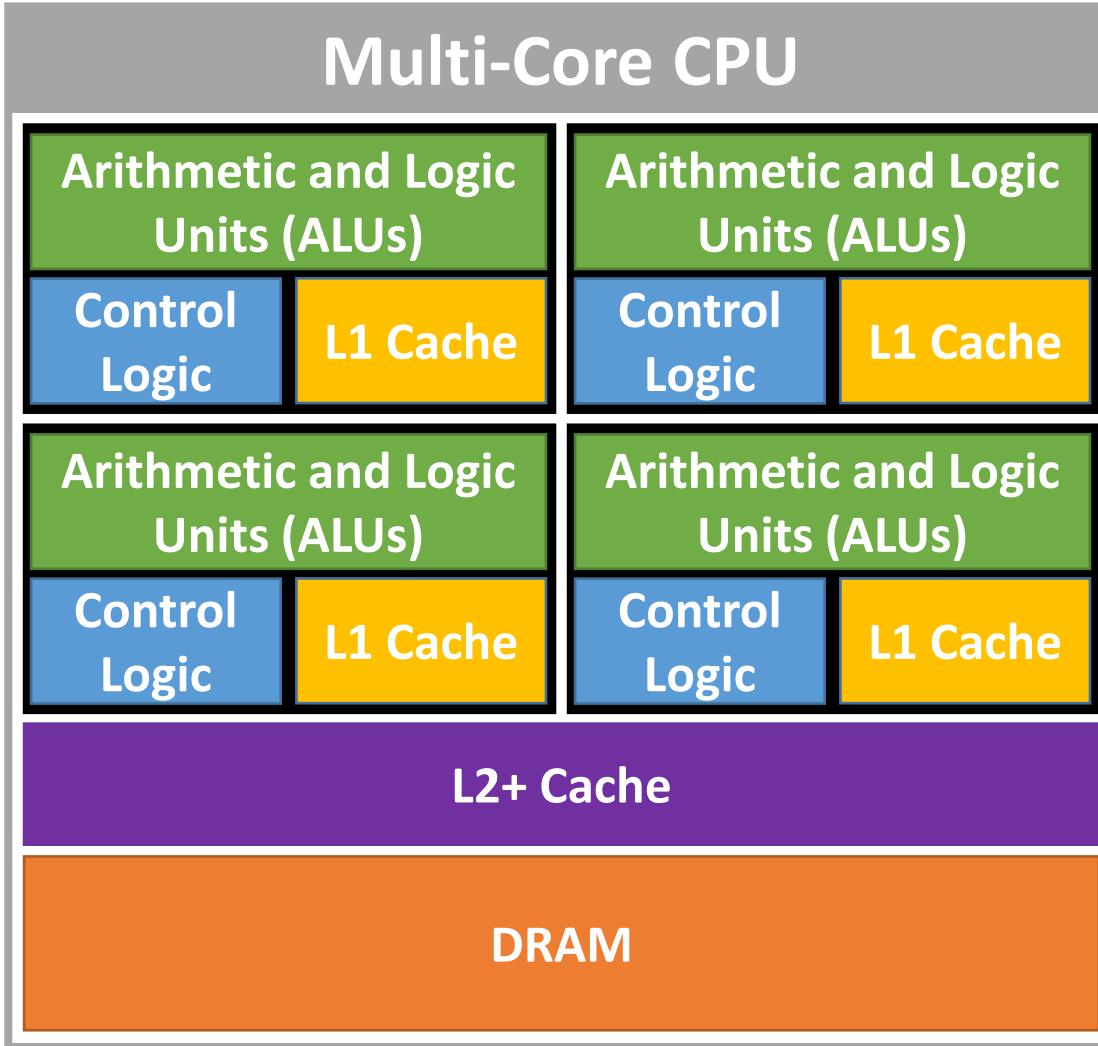
CPUs and GPUs have fundamentally different strengths and weaknesses



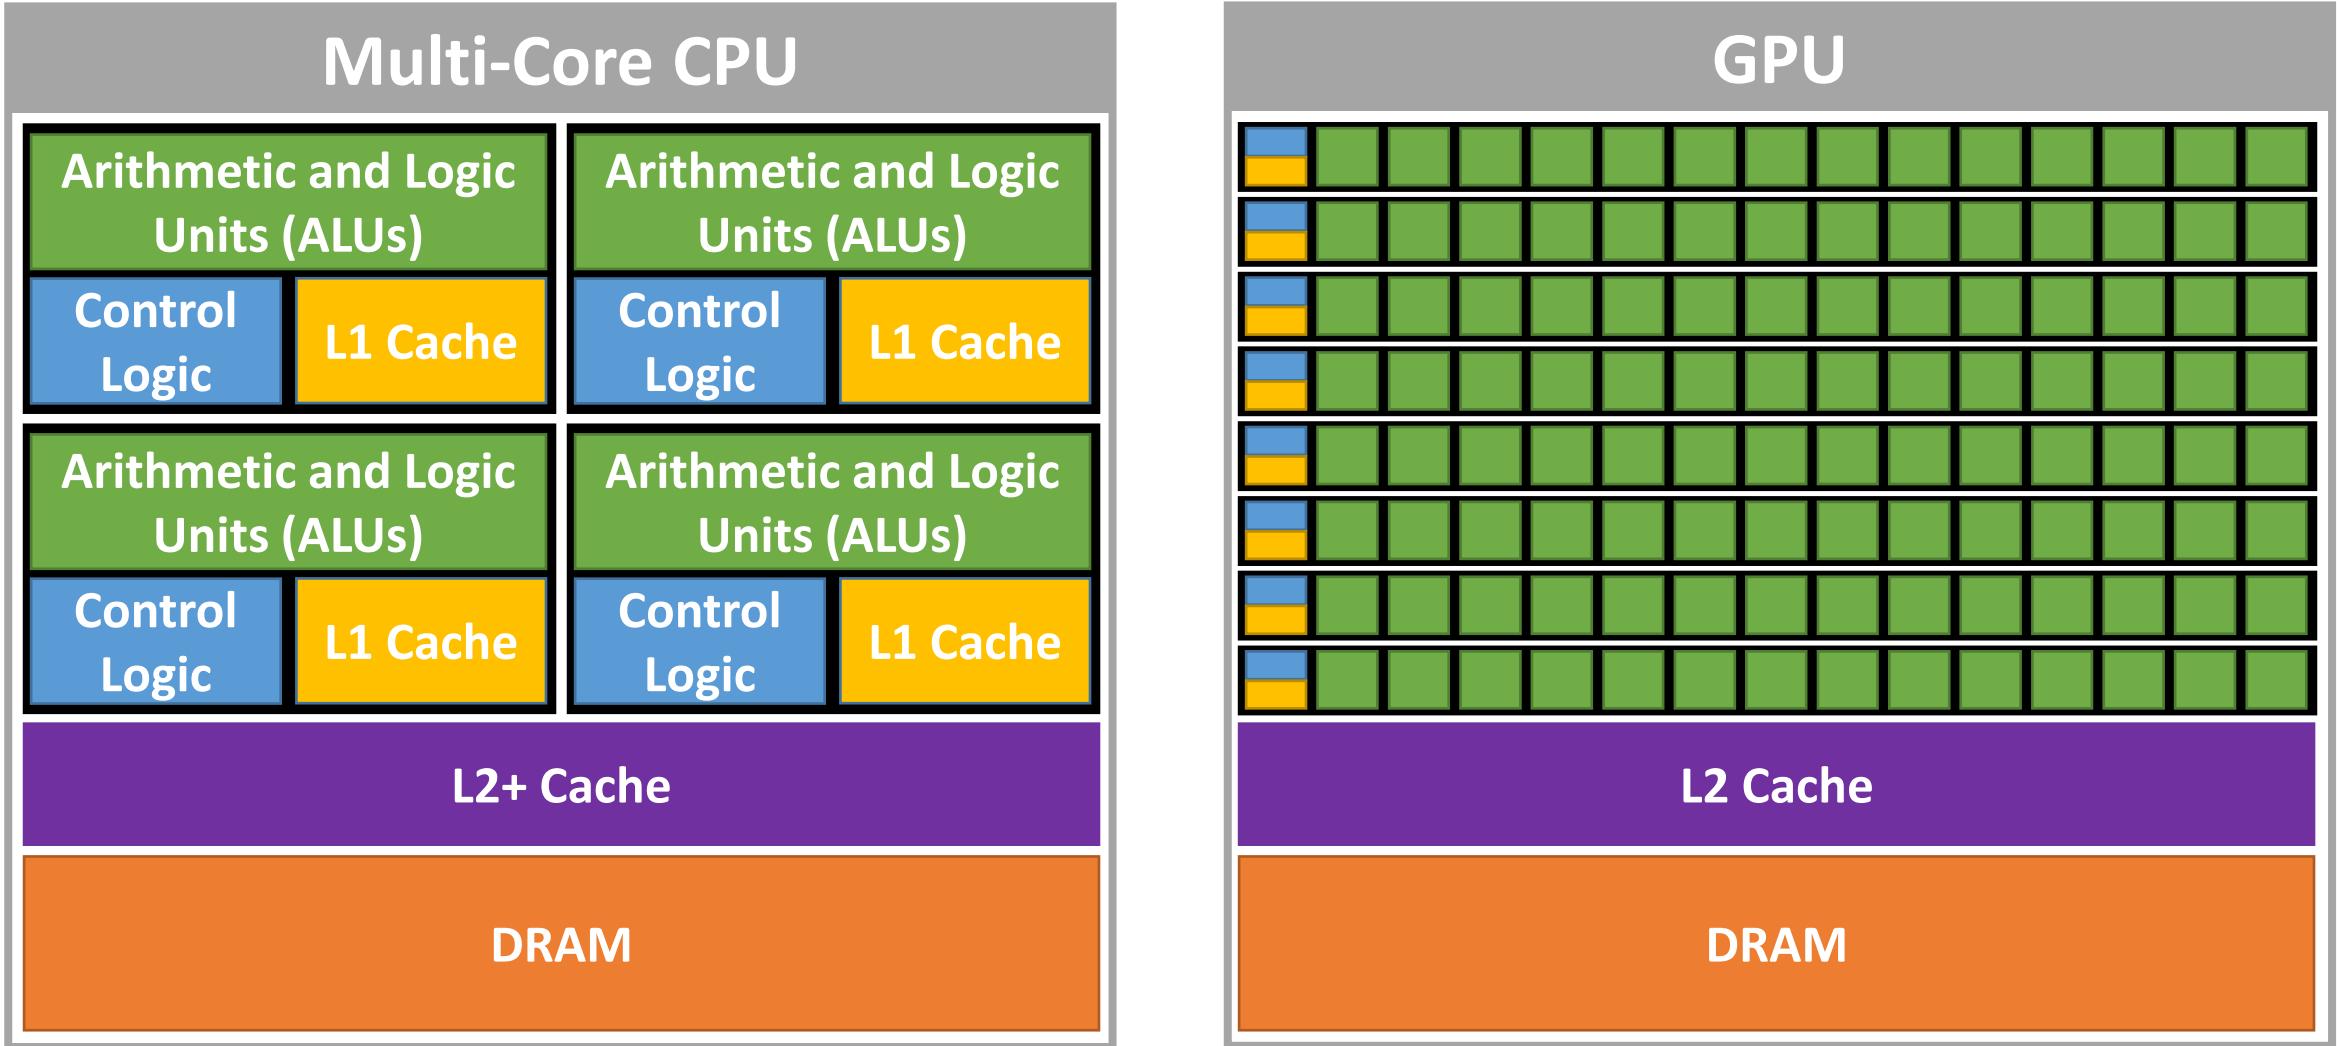
CPUs and GPUs have fundamentally different strengths and weaknesses



CPUs and GPUs have fundamentally different strengths and weaknesses



CPUs and GPUs have fundamentally different strengths and weaknesses



Hardware Acceleration for Realtime Robotics



A Quick Overview of (And My Approach To) Robotics

A Crash Course in (GPU) Architecture

A Case Study in Accelerating Rigid Body Dynamics

The Future of Accelerating Robotics

Hardware Acceleration for Realtime Robotics



A Quick Overview of (And My Approach To) Robotics

A Crash Course in (GPU) Architecture

A Case Study in Accelerating Rigid Body Dynamics

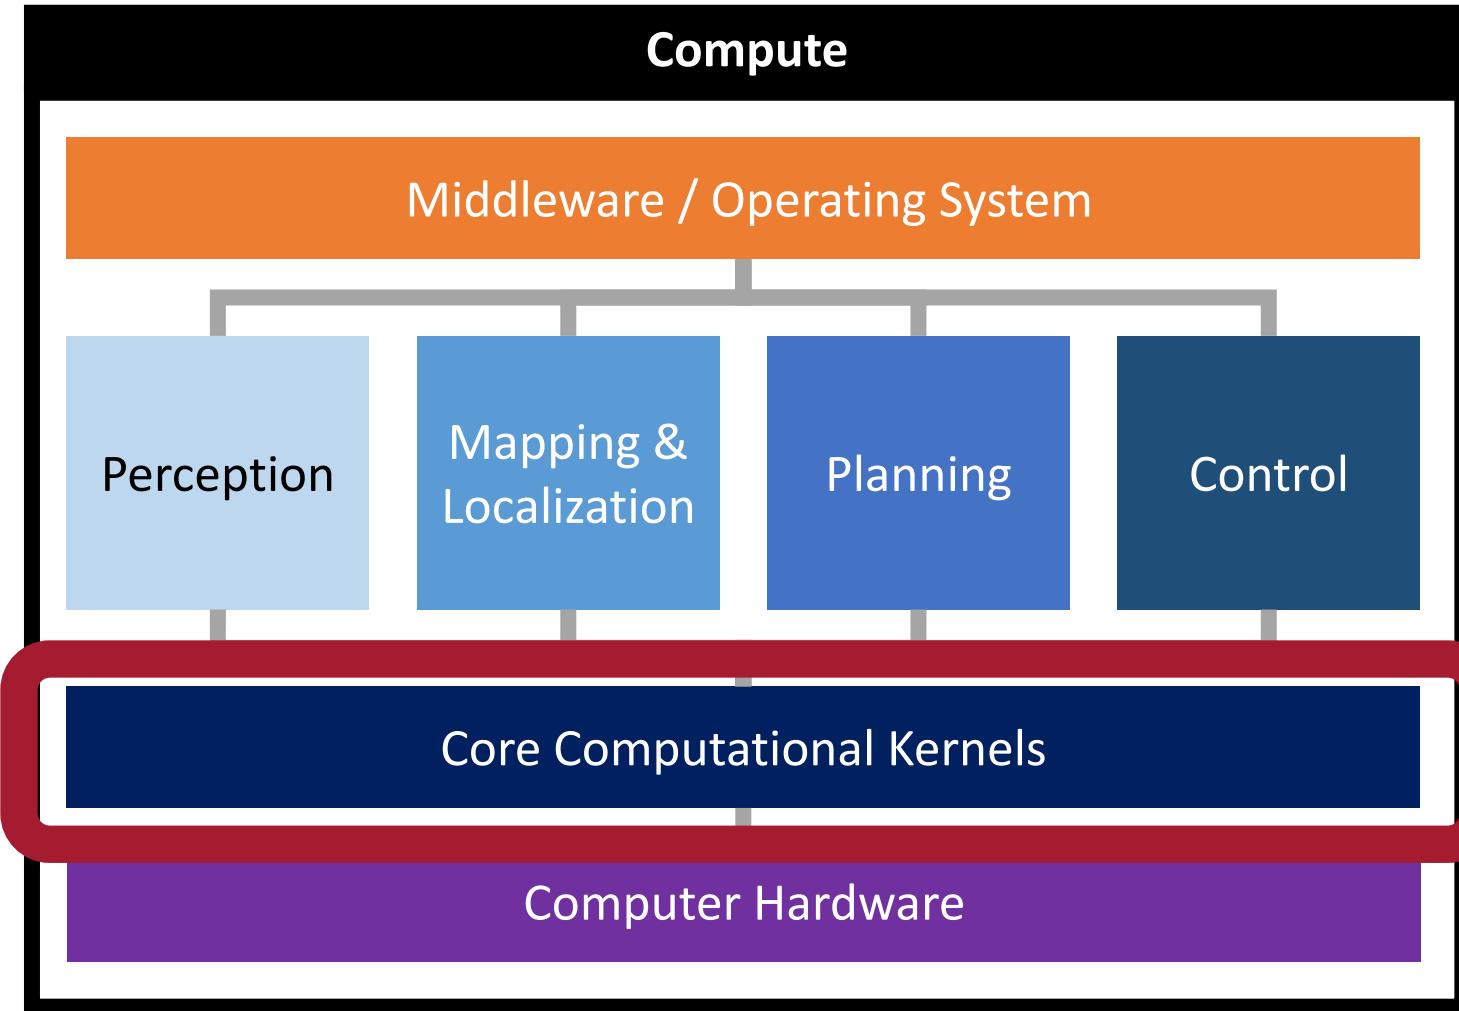
Why Accelerate Rigid Body Dynamics?

Accelerating Rigid Body Dynamics on the GPU

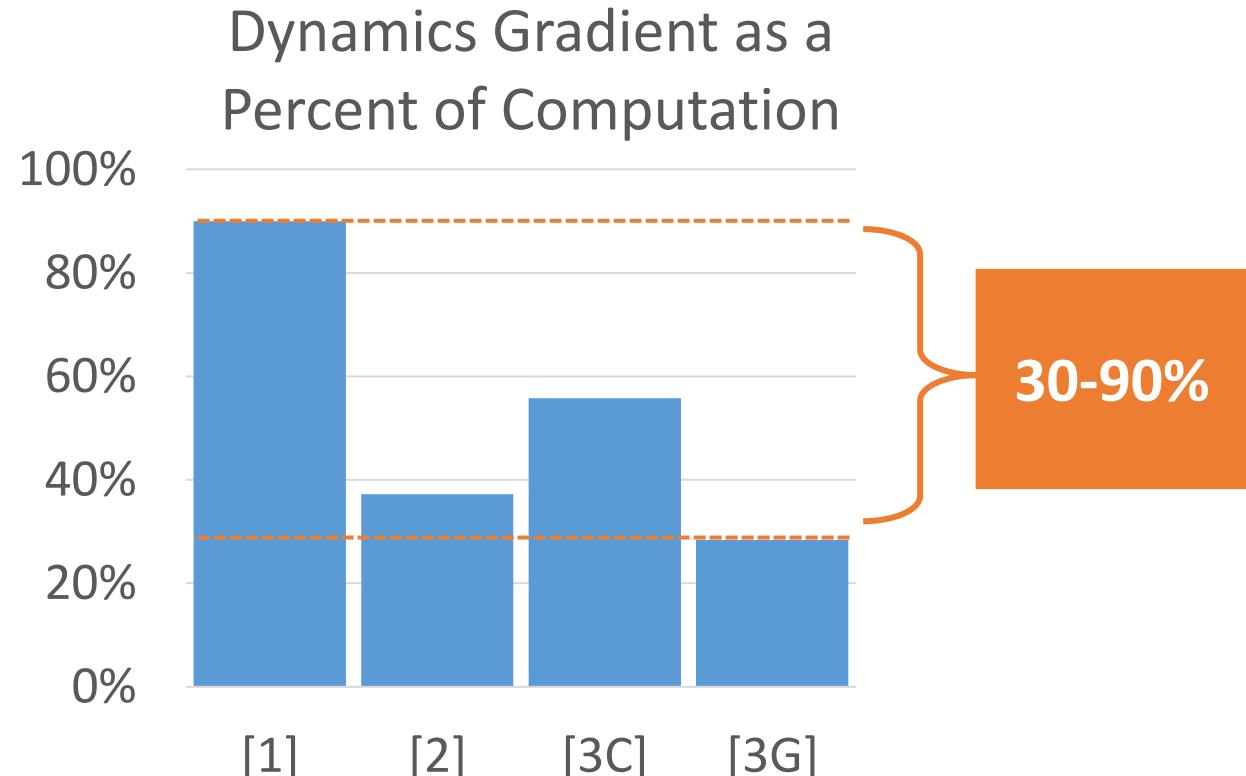
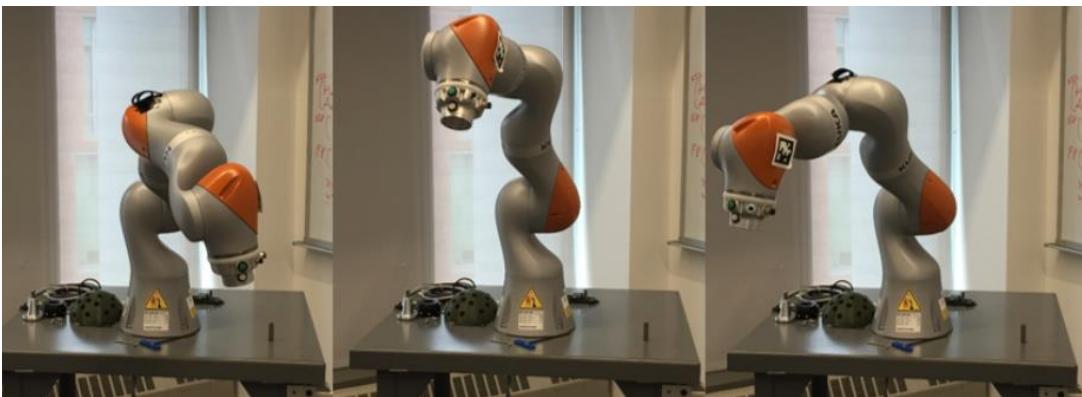
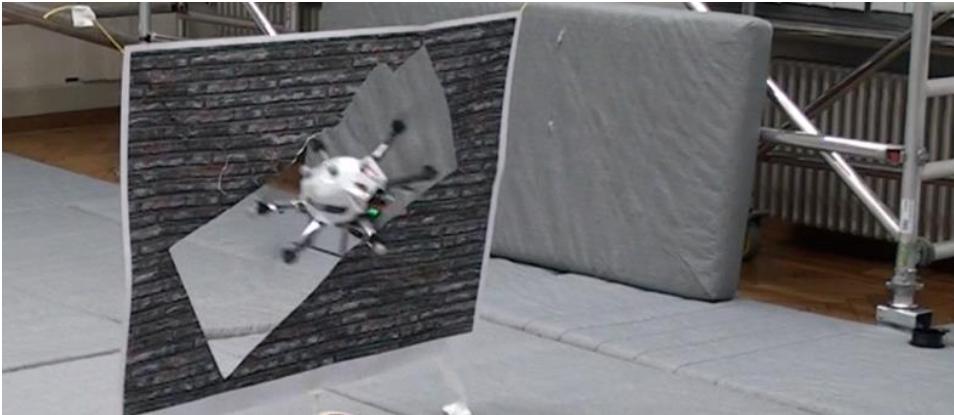
Moving Beyond the GPU

The Future of Accelerating Robotics

Why Rigid Body Dynamics



Rigid Body Dynamics Gradients are a bottleneck for planning and control



[1] J. Carpentier and N. Mansrud,
“Analytical Derivatives of Rigid Body
Dynamics Algorithms,” RSS 2018

[2] M. Neunert, et al., “Fast nonlinear Model
Predictive Control for unified trajectory
optimization and tracking,” ICRA 2016

[3] Best end-to-end [C]PU and [G]PU option from B. Plancher and S.
Kuindersma, “A Performance Analysis of Parallel Differential Dynamic
Programming,” WAFR 2018

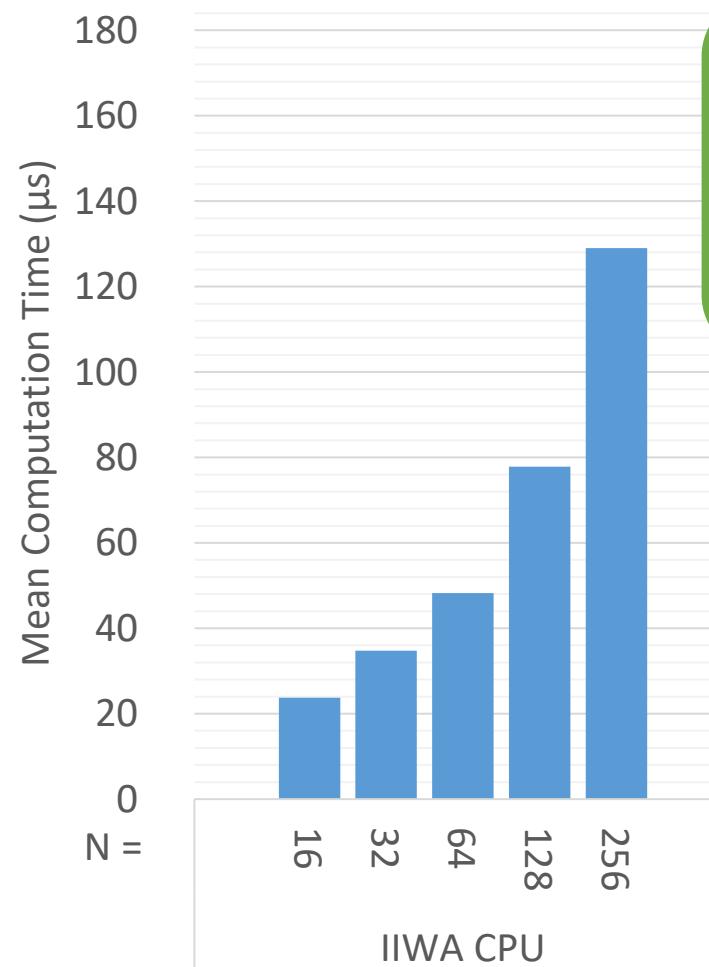
Why GPU Acceleration for Rigid Body Dynamics?

Planning and Control Algorithms often require many simultaneous computations of Rigid Body Dynamics



Why GPU Acceleration for Rigid Body Dynamics?

Forward Dynamics Gradient Multiple
Computation Latency



This can become very
time consuming
quickly! And is
naturally parallel!



Hardware Acceleration for Realtime Robotics



A Quick Overview of (And My Approach To) Robotics

A Crash Course in (GPU) Architecture

A Case Study in Accelerating Rigid Body Dynamics

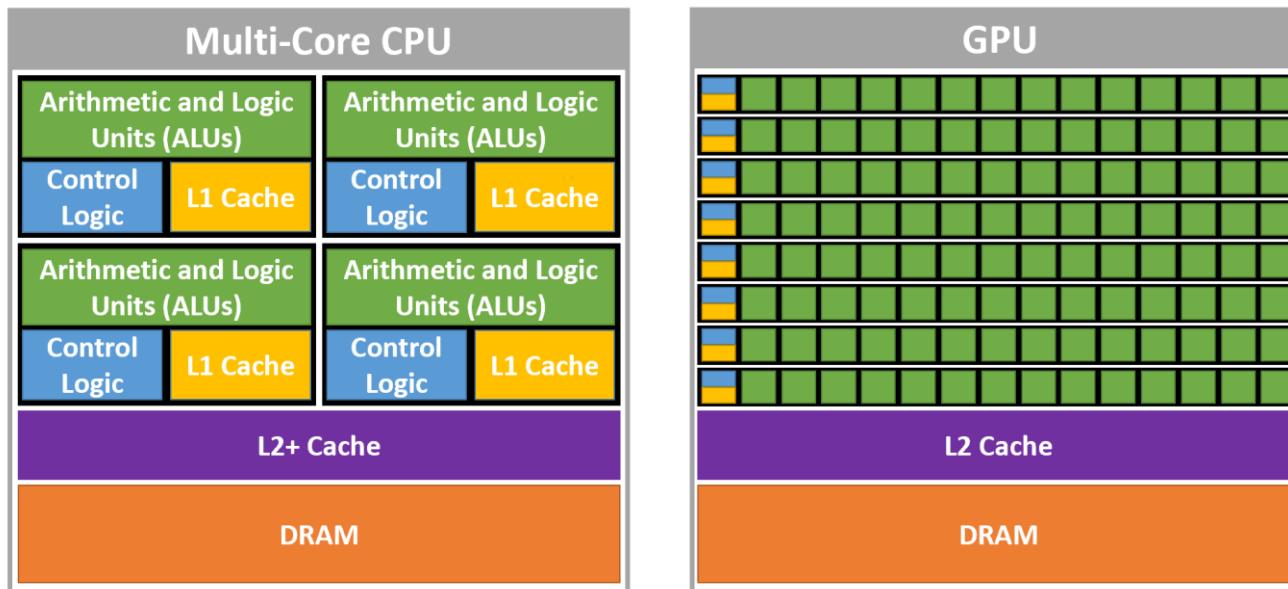
Why Accelerate Rigid Body Dynamics?

Accelerating Rigid Body Dynamics on the GPU

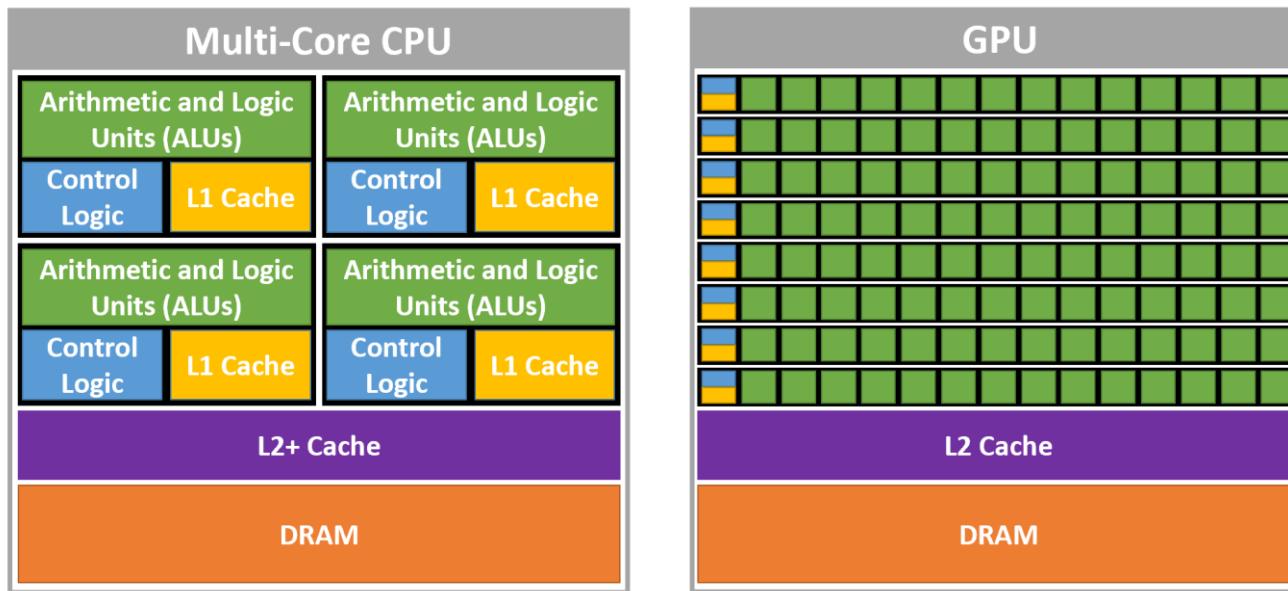
Moving Beyond the GPU

The Future of Accelerating Robotics

How do we Accelerate Rigid Body Dynamics on a GPU?



How do we Accelerate Rigid Body Dynamics on a GPU?

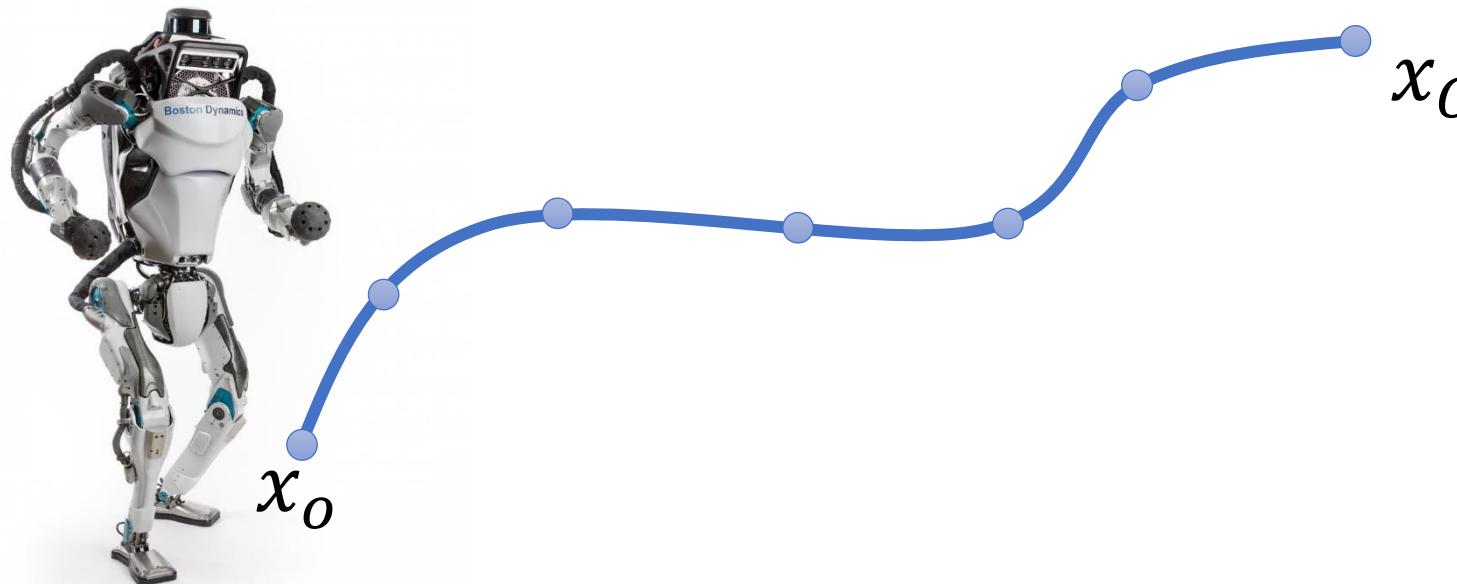


Hardware-Software Co-Design

High performance code needs to be **refactored** to take advantage of **different hardware** computational strengths and weaknesses

Algorithmic Features of the Gradient of Rigid Body Dynamics as a step of a planning and control algorithm

Algorithmic Feature	CPU	GPU
Parallelism	Bad	Good



Algorithmic Features of the Gradient of Rigid Body Dynamics as a step of a planning and control algorithm

Algorithmic Feature	CPU	GPU
Parallelism	Bad	Good
?		

What is Rigid Body Dynamics?

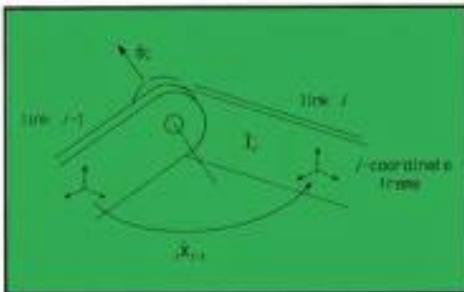
What is Rigid Body Dynamics?

Physics

What is Rigid Body Dynamics?

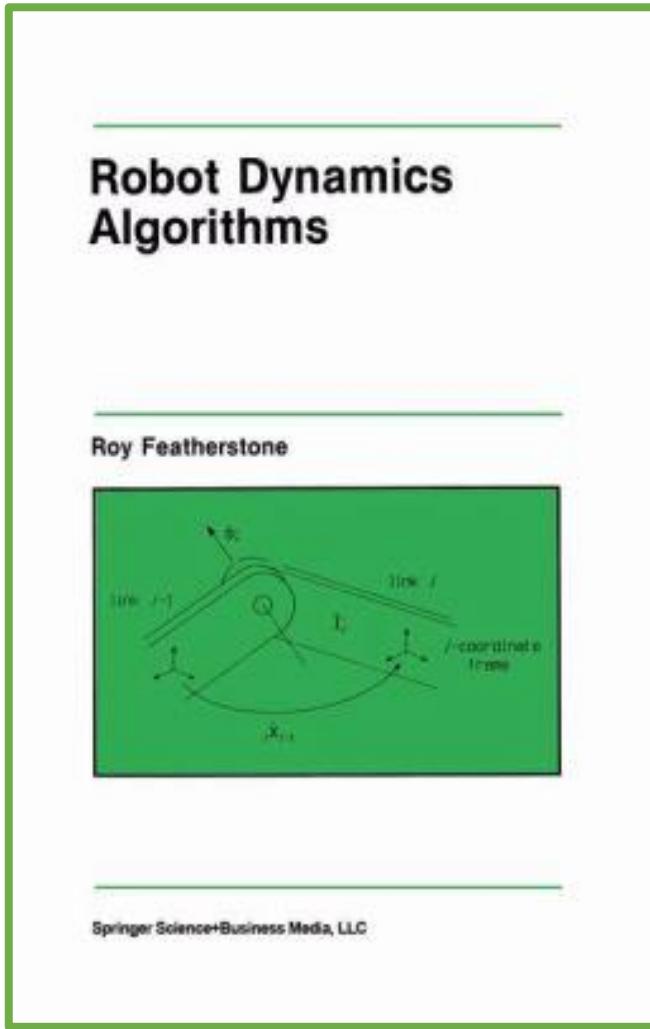
Robot Dynamics Algorithms

Roy Featherstone



Springer Science+Business Media, LLC

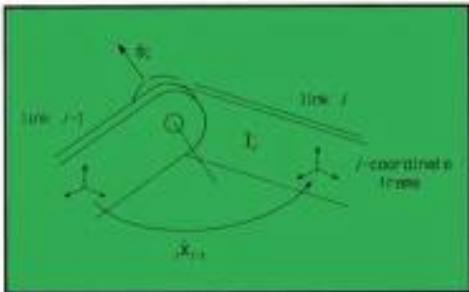
What is Rigid Body Dynamics?



What is Rigid Body Dynamics?

Robot Dynamics Algorithms

Roy Featherstone

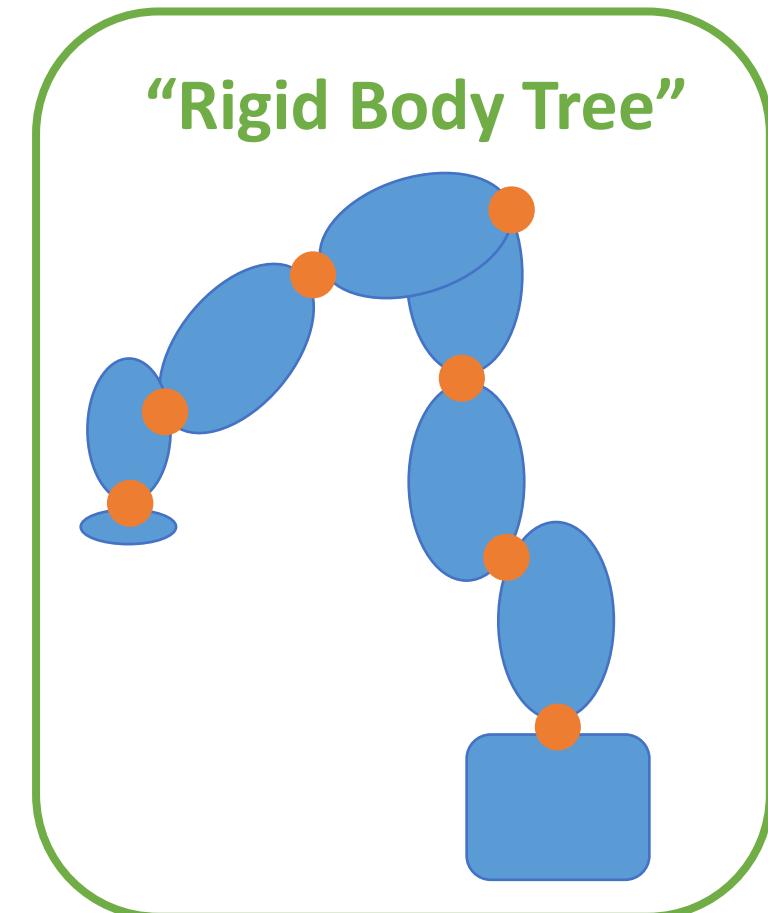
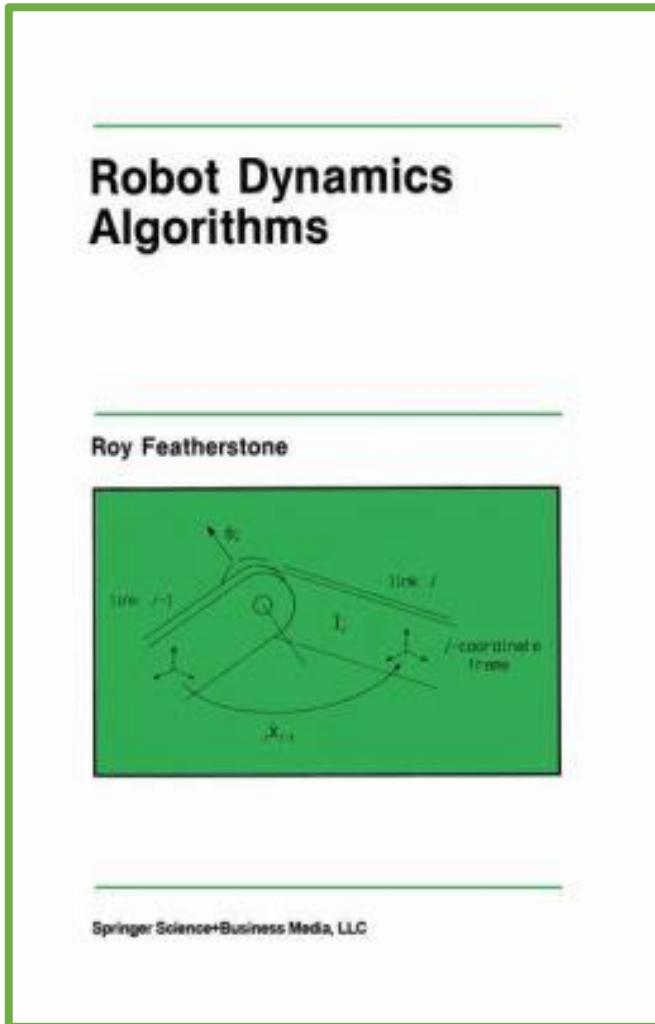


Springer Science+Business Media, LLC

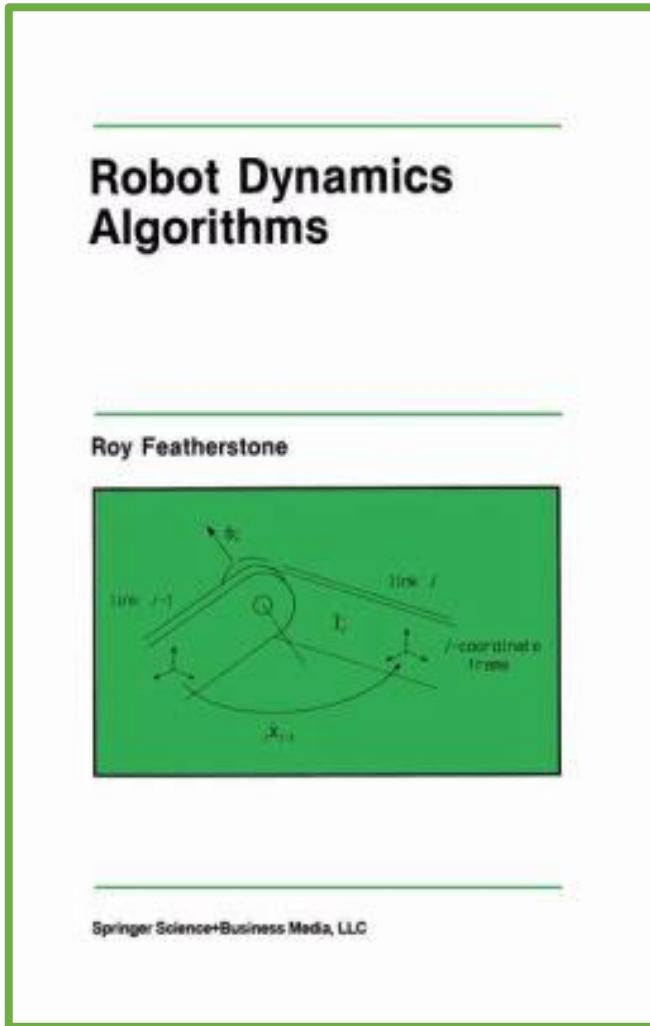


“Rigid Body Tree”

What is Rigid Body Dynamics?



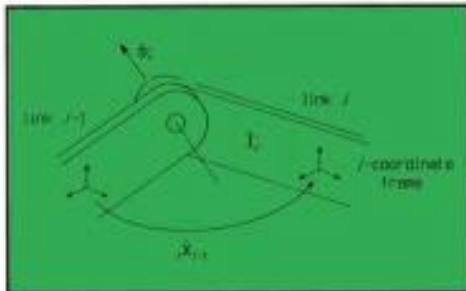
What is Rigid Body Dynamics?



What is Rigid Body Dynamics?

Robot Dynamics Algorithms

Roy Featherstone



Springer Science+Business Media, LLC



The study of the motion of bodies under the action of forces (**aka $F = ma$**)

How do Rigid Body Dynamics Algorithms Work?

How do Rigid Body Dynamics Algorithms Work?

Algorithm 1 RNEA($q, \dot{q}, \ddot{q}, f^{ext}$) $\rightarrow c$

2: **for** link $i = 1 : n$ **do**

Joints/Links are
numbered in
order

7: **for** link $i = n : 1$ **do**

The algorithms

up and down the tree

How do Rigid Body Dynamics Algorithms Work?

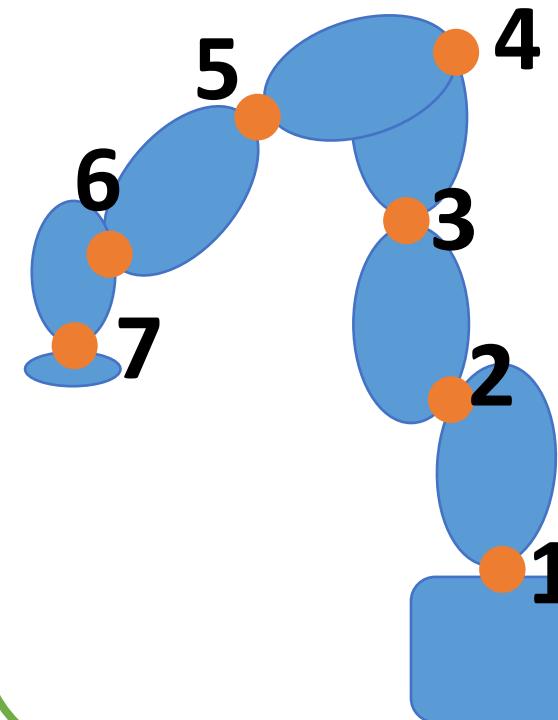
Algorithm 1 RNEA($q, \dot{q}, \ddot{q}, f^{ext}$) $\rightarrow c$

2: **for** link $i = 1 : n$ **do**

Joints/Links are numbered in order

7: **for** link $i = n : 1$ **do**

“Rigid Body Tree”



How do Rigid Body Dynamics Algorithms Work?

Algorithm 1 RNEA($q, \dot{q}, \ddot{q}, f^{ext} \rightarrow c$

2: **for** link $i = 1 : n$ **do**

4: $v_i =$

5: $a_i =$

6: $f_i =$

7: **for** link $i = n : 1$ **do**

8: c_i

9: f_{λ_i}

λ refers to the
parent link

The algorithms

pass the **forces (f)**,
velocities (v), and
accelerations (a)
up and down the tree

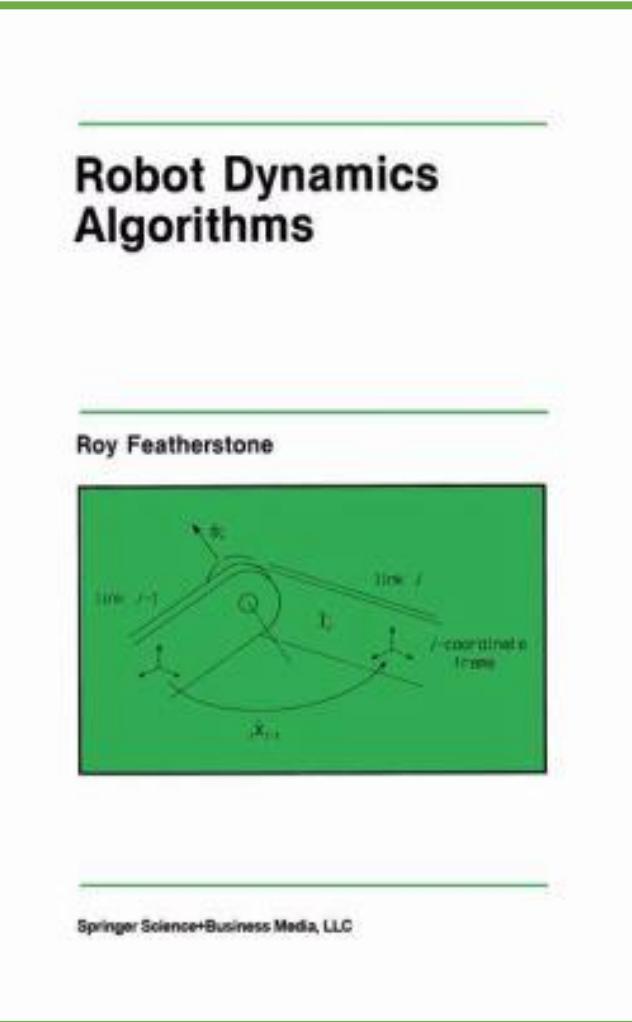
How do Rigid Body Dynamics Algorithms Work?

Algorithm 1 RNEA($q, \dot{q}, \ddot{q}, f^{ext} \rightarrow c$

```
1:  $v_0 = 0, a_0 = \text{gravity}$ 
2: for link  $i = 1 : n$  do
3:   Compute  ${}^i X_{\lambda_i}, S_i, I_i$ 
4:    $v_i = {}^i X_{\lambda_i} v_{\lambda_i} + S_i \dot{q}_i$ 
5:    $a_i = {}^i X_{\lambda_i} a_{\lambda_i} + S_i \ddot{q}_i + v_i \times S_i \dot{q}_i$ 
6:    $f_i = I_i a_i + v_i \times^* I_i v_i - f_i^{ext}$ 
7: for link  $i = n : 1$  do
8:    $c_i = S_i^T f_i$ 
9:    $f_{\lambda_i} += {}^i X_{\lambda_i}^T f_i$ 
```

The algorithms use **transformations (X)** and **inertias (I)** along with **joint direction vectors (S)** and **cross products (\times)** to pass the **forces (f)**, **velocities (v)**, and **accelerations (a)** up and down the tree

Rigid Body Dynamics Algorithms



- The study of the motion of bodies under the action of forces ($F = ma$)
- It **loops up and down** the rigid body tree
- It uses a **variety of different operations** in each step

What is the Gradient of Rigid Body Dynamics?

Algorithm 2 $\nabla\text{RNEA}(\dot{q}, v, a, f, X, S, I)$

1: **for** link $i = 1 : N$ **do**

$$2: \quad \frac{\partial v_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial v_{\lambda_i}}{\partial u} + \begin{cases} ({}^i X_{\lambda_i} v_{\lambda_i}) \times S_i & u = q \\ S_i & u \equiv \dot{q} \end{cases}$$

$$3: \quad \frac{\partial a_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial a_{\lambda_i}}{\partial u} + \frac{\partial v_{\lambda_i}}{\partial u} \times S_i \dot{q}_i + \begin{cases} ({}^i X_{\lambda_i} a_{\lambda_i}) \times S_i \\ v_i \times S_i \end{cases}$$

$$4: \quad \frac{\partial f_i}{\partial u} = I_i \frac{\partial a_i}{\partial u} + \frac{\partial v_i}{\partial u} \times^* I_i v_i + v_i \times^* I_i \frac{\partial v_i}{\partial u}$$

5: **for** link $i = N : 1$ **do**

$$6: \quad \frac{\partial c_i}{\partial u} = S_i^T \frac{\partial f_i}{\partial u}$$

$$7: \quad \frac{\partial f_{\lambda_i}}{\partial u} += {}^i X_{\lambda_i}^T \frac{\partial f_i}{\partial u} + {}^i X_{\lambda_i}^T (S_i \times^* f_i)$$

The multi-dimensional derivative

Basically the same algorithm (just a lot of *dus*)

Algorithmic Features of the Gradient of Rigid Body Dynamics as a step of a planning and control algorithm

Algorithmic Feature	CPU	GPU
Parallelism	Bad	Good
Serial Dependencies		
Irregular Operations	Good	

Each of the cores is designed to be good at general purpose computations

The diagram is a table comparing algorithmic features across two computing platforms: CPU and GPU. The rows represent different features: Parallelism, Serial Dependencies, and Irregular Operations. The columns represent the platforms: CPU and GPU. The entries indicate whether each feature is 'Bad' or 'Good' for that platform. A blue arrow points from the 'Good' entry for Irregular Operations on the GPU column to a callout box. The callout box contains the text: 'Each of the cores is designed to be good at general purpose computations'. The table has dashed horizontal lines between the rows and solid vertical lines between the columns.

Algorithmic Features of the Gradient of Rigid Body Dynamics as a step of a planning and control algorithm

Algorithmic Feature	CPU	GPU
Parallelism	Bad	Good
Serial Dependencies	Good	Bad
Irregular Operations		

All cores are optimized for regular and parallel operations

Algorithmic refactoring is needed to effectively target GPUs

Algorithm 2 $\nabla\text{RNEA}(\dot{q}, v, a, f, X, S, I) \rightarrow \partial c / \partial u$

1: **for** link $i = 1 : N$ **do**

2: $\frac{\partial v_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial v_{\lambda_i}}{\partial u} + \begin{cases} ({}^i X_{\lambda_i} v_{\lambda_i}) \times S_i & u \equiv q \\ S_i & u \equiv \dot{q} \end{cases}$

3: $\frac{\partial a_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial a_{\lambda_i}}{\partial u} + \frac{\partial v_{\lambda_i}}{\partial u} \times S_i \dot{q}_i + \begin{cases} ({}^i X_{\lambda_i} a_{\lambda_i}) \times S_i \\ v_i \times S_i \end{cases}$

4: $\frac{\partial f_i}{\partial u} = I_i \frac{\partial a_i}{\partial u} + \frac{\partial v_i}{\partial u} \times^* I_i v_i + v_i \times^* I_i \frac{\partial v_i}{\partial u}$

5: **for** link $i = N : 1$ **do**

6: $\frac{\partial c_i}{\partial u} = S_i^T \frac{\partial f_i}{\partial u}$

7: $\frac{\partial f_{\lambda_i}}{\partial u} += {}^i X_{\lambda_i}^T \frac{\partial f_i}{\partial u} + {}^i X_{\lambda_i}^T (S_i \times^* f_i)$

Algorithmic refactoring is needed to effectively target GPUs

Generate temporary values in parallel to reduce the computations done in serial loops

Algorithm 2 $\nabla\text{RNEA}(\dot{q}, v, a, f, X, S, I) \rightarrow \partial c / \partial u$

```

1: for link  $i = 1 : N$  do
2:    $\frac{\partial v_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial v_{\lambda_i}}{\partial u} + \begin{cases} ({}^i X_{\lambda_i} v_{\lambda_i}) \times S_i & u \equiv q \\ S_i & u \equiv \dot{q} \end{cases}$ 
3:    $\frac{\partial a_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial a_{\lambda_i}}{\partial u} + \frac{\partial v_{\lambda_i}}{\partial u} \times S_i \dot{q}_i + \begin{cases} ({}^i X_{\lambda_i} a_{\lambda_i}) \times S_i & \\ v_i \times S_i & \end{cases}$ 
4:    $\frac{\partial f_i}{\partial u} = I_i \frac{\partial a_i}{\partial u} + \frac{\partial v_i}{\partial u} \times^* I_i v_i + v_i \times^* I_i \frac{\partial v_i}{\partial u}$ 
5: for link  $i = N : 1$  do
6:    $\frac{\partial c_i}{\partial u} = S_i^T \frac{\partial f_i}{\partial u}$ 
7:    $\frac{\partial f_{\lambda_i}}{\partial u} += {}^i X_{\lambda_i}^T \frac{\partial f_i}{\partial u} + {}^i X_{\lambda_i}^T (S_i \times^* f_i)$ 

```



Algorithm 4 $\nabla\text{RNEA-GPU}(\dot{q}, v, a, f, X, S, I) \rightarrow \partial c / \partial u$

```

1: for link  $i = 1 : n$  in parallel do
2:    $\alpha_i = {}^i X_{\lambda_i} v_{\lambda_i}$   $\beta_i = {}^i X_{\lambda_i} a_{\lambda_i}$   $\gamma_i = I_i v_i$ 
3:    $\alpha_i = \alpha_i \times S_i$   $\beta_i = \beta_i \times S_i$   $\delta_i = v_i \times S_i$ 
4:    $\zeta_i = f_i \times S_i$   $\eta_i = v_i \times^*$ 
5:    $\zeta_i = -{}^i X_{\lambda_i}^T \zeta_i$   $\eta_i = \eta_i I_i$ 
6:   for link  $i = 1 : n$  do
7:      $\frac{\partial v_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial v_{\lambda_i}}{\partial u} + \begin{cases} \alpha_i & u \equiv q \\ S_i & u \equiv \dot{q} \end{cases}$ 
8:      $\mu_i = \frac{\partial v_i}{\partial u} \times^*$   $\rho_i = \frac{\partial v_{\lambda_i}}{\partial u} \times S_i \dot{q}_i + \begin{cases} \beta_i & \\ \delta_i & \end{cases}$ 
9:   for link  $i = 1 : n$  do
10:     $\frac{\partial a_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial a_{\lambda_i}}{\partial u} + \rho_i$ 
11:   for link  $i = 1 : n$  in parallel do
12:      $\frac{\partial f_i}{\partial u} = I_i \frac{\partial a_i}{\partial u} + \mu_i \gamma_i + \eta_i \frac{\partial v_i}{\partial u}$ 
13:   for link  $i = n : 1$  do
14:      $\frac{\partial f_{\lambda_i}}{\partial u} += {}^i X_{\lambda_i}^T \frac{\partial f_i}{\partial u} + \zeta_i$ 
15:   for link  $i = n : 1$  in parallel do
16:      $\frac{\partial c_i}{\partial u} = S_i^T \frac{\partial f_i}{\partial u}$ 

```

Algorithmic refactoring is needed to effectively target GPUs

Re-ordering computations to unify operations

Algorithm 2 $\nabla\text{RNEA}(\dot{q}, v, a, f, X, S, I) \rightarrow \partial c / \partial u$

```

1: for link  $i = 1 : N$  do
2:    $\frac{\partial v_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial v_{\lambda_i}}{\partial u} + \begin{cases} ({}^i X_{\lambda_i} v_{\lambda_i}) \times S_i & u \equiv q \\ S_i & u \equiv \dot{q} \end{cases}$ 
3:    $\frac{\partial a_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial a_{\lambda_i}}{\partial u} + \frac{\partial v_{\lambda_i}}{\partial u} \times S_i \dot{q}_i + \begin{cases} ({}^i X_{\lambda_i} a_{\lambda_i}) \times S_i & u \equiv q \\ v_i \times S_i & u \equiv \dot{q} \end{cases}$ 
4:    $\frac{\partial f_i}{\partial u} = I_i \frac{\partial a_i}{\partial u} + \frac{\partial v_i}{\partial u} \times^* I_i v_i + v_i \times^* I_i \frac{\partial v_i}{\partial u}$ 
5: for link  $i = N : 1$  do
6:    $\frac{\partial c_i}{\partial u} = S_i^T \frac{\partial f_i}{\partial u}$ 
7:    $\frac{\partial f_{\lambda_i}}{\partial u} += {}^i X_{\lambda_i}^T \frac{\partial f_i}{\partial u} + {}^i X_{\lambda_i}^T (S_i \times^* f_i)$ 

```



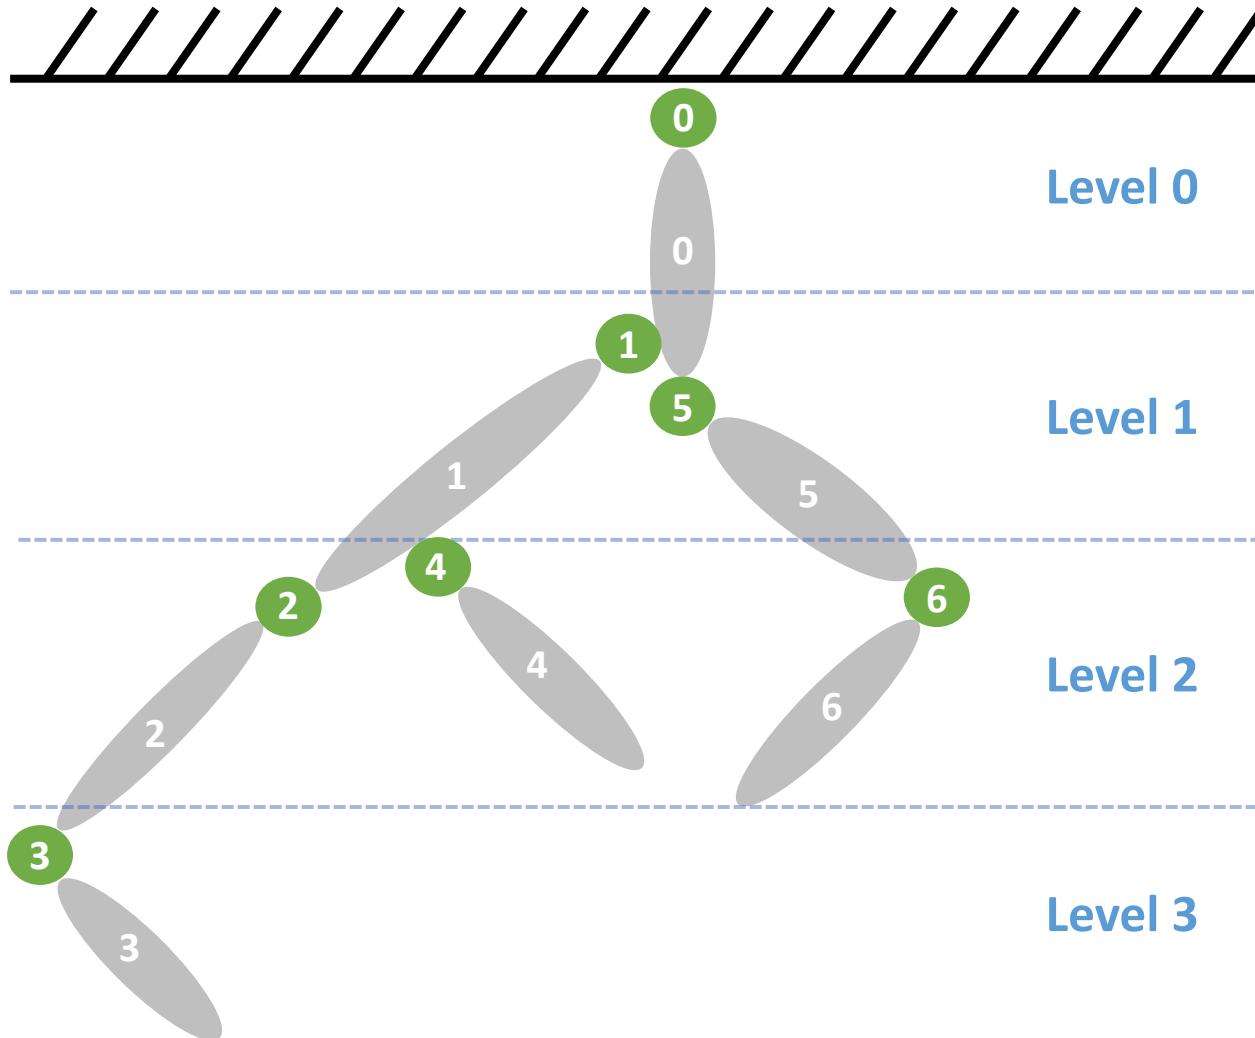
Algorithm 4 $\nabla\text{RNEA-GPU}(\dot{q}, v, a, f, X, S, I) \rightarrow \partial c / \partial u$

```

1: for link  $i = 1 : n$  in parallel do
2:    $\alpha_i = {}^i X_{\lambda_i} v_{\lambda_i}$   $\beta_i = {}^i X_{\lambda_i} a_{\lambda_i}$   $\gamma_i = I_i v_i$ 
3:    $\alpha_i = \alpha_i \times S_i$   $\beta_i = \beta_i \times S_i$   $\delta_i = v_i \times S_i$ 
4:    $\zeta_i = f_i \times S_i$   $\eta_i = v_i \times^*$ 
5:    $\zeta_i = -{}^i X_{\lambda_i}^T \zeta_i$   $\eta_i = \eta_i I_i$ 
6: for link  $i = 1 : n$  do
7:    $\frac{\partial v_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial v_{\lambda_i}}{\partial u} + \begin{cases} \alpha_i & u \equiv q \\ S_i & u \equiv \dot{q} \end{cases}$ 
8:    $\mu_i = \frac{\partial v_i}{\partial u} \times^*$   $\rho_i = \frac{\partial v_{\lambda_i}}{\partial u} \times S_i \dot{q}_i + \begin{cases} \beta_i & u \equiv q \\ \delta_i & u \equiv \dot{q} \end{cases}$ 
9: for link  $i = 1 : n$  do
10:   $\frac{\partial a_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial a_{\lambda_i}}{\partial u} + \rho_i$ 
11: for link  $i = 1 : n$  in parallel do
12:   $\frac{\partial f_i}{\partial u} = I_i \frac{\partial a_i}{\partial u} + \mu_i \gamma_i + \eta_i \frac{\partial v_i}{\partial u}$ 
13: for link  $i = n : 1$  do
14:   $\frac{\partial f_{\lambda_i}}{\partial u} += {}^i X_{\lambda_i}^T \frac{\partial f_i}{\partial u} + \zeta_i$ 
15: for link  $i = n : 1$  in parallel do
16:   $\frac{\partial c_i}{\partial u} = S_i^T \frac{\partial f_i}{\partial u}$ 

```

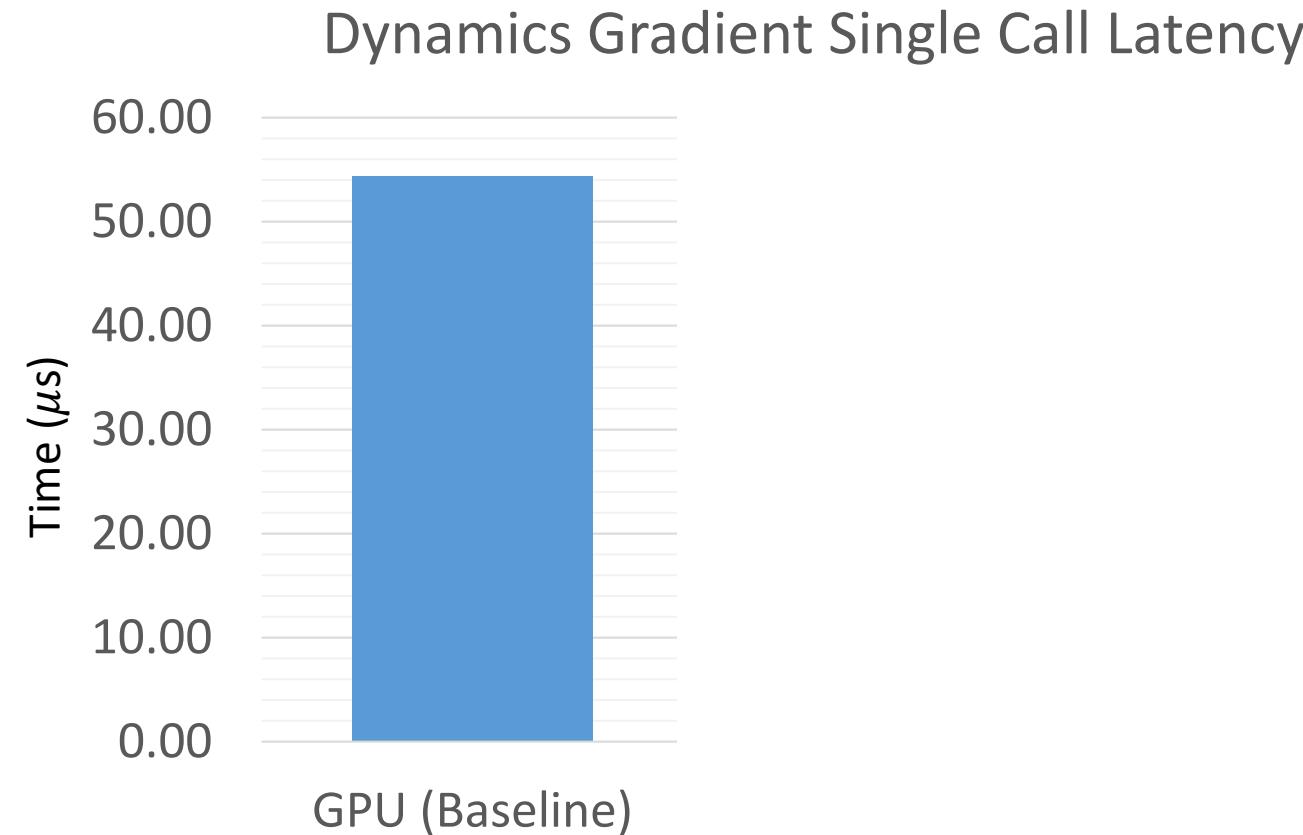
Algorithmic refactoring is needed to effectively target GPUs



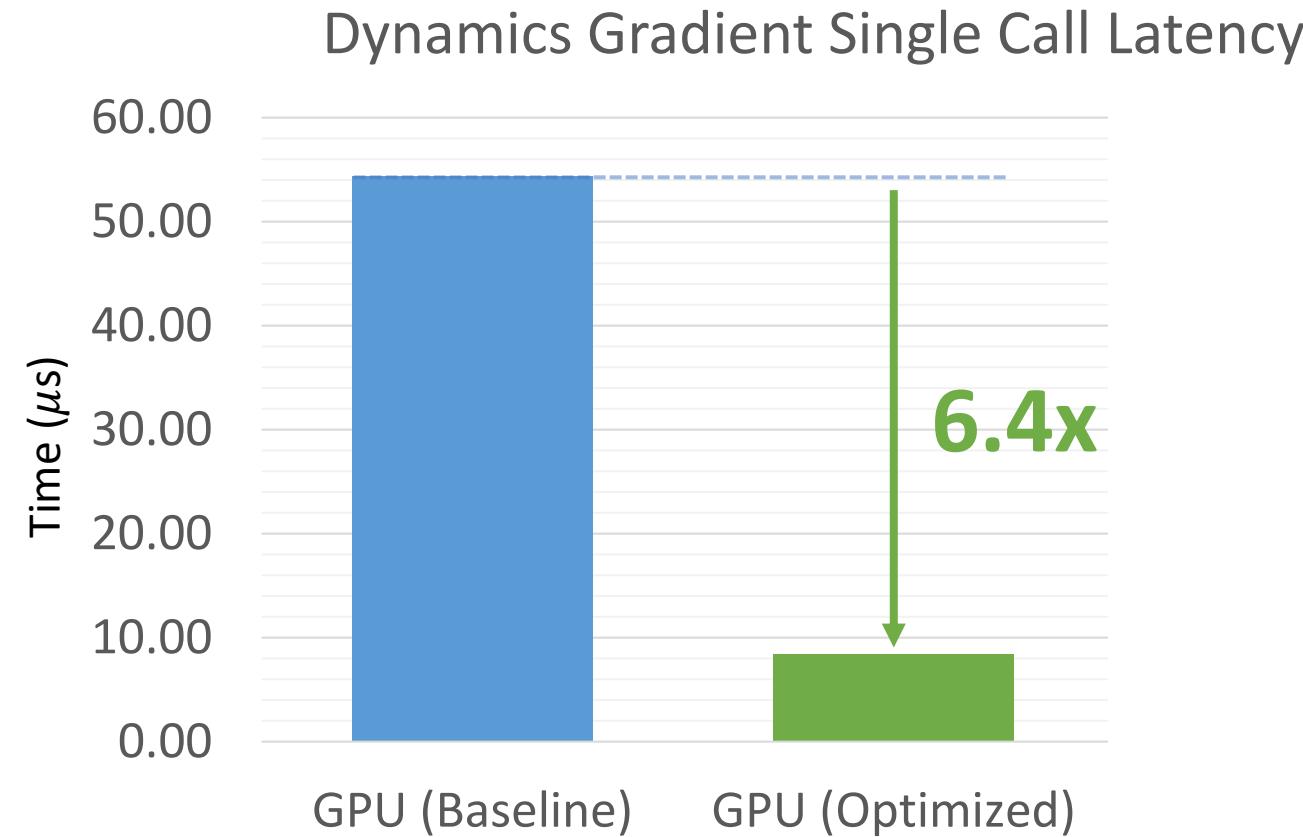
Algorithm 4 $\nabla_{\text{RNEA-GPU}}(\dot{q}, v, a, f, X, S, I) \rightarrow \partial c / \partial u$

```
1: for link  $i = 1 : n$  in parallel do
2:    $\alpha_i = {}^i X_{\lambda_i} v_{\lambda_i}$   $\beta_i = {}^i X_{\lambda_i} a_{\lambda_i}$   $\gamma_i = I_i v_i$ 
3:    $\alpha_i = \alpha_i \times S_i$   $\beta_i = \beta_i \times S_i$   $\delta_i = v_i \times S_i$ 
    $\zeta_i = f_i \times S_i$   $\eta_i = v_i \times {}^*$ 
4:    $\zeta_i = -{}^i X_{\lambda_i}^T \zeta_i$   $\eta_i = n_i I_i$ 
5: for link  $i = 1 : n$  do In parallel within each level
6:    $\frac{\partial v_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial v_{\lambda_i}}{\partial u} + \begin{cases} \alpha_i & u = q \\ S_i & u \equiv \dot{q} \end{cases}$ 
7: for link  $i = 1 : n$  in parallel do
8:    $\mu_i = \frac{\partial v_i}{\partial u} \times {}^*$   $\rho_i = \frac{\partial v_{\lambda_i}}{\partial u} \times S_i \dot{q}_i + \begin{cases} \beta_i & \\ \delta_i & \end{cases}$ 
In parallel within each level
9: for link  $i = 1 : n$  do
10:   $\frac{\partial a_i}{\partial u} = {}^i X_{\lambda_i} \frac{\partial a_{\lambda_i}}{\partial u} + \rho_i$ 
11: for link  $i = 1 : n$  in parallel do
12:   $\frac{\partial f_i}{\partial u} = I_i \frac{\partial a_i}{\partial u} + \mu_i \gamma_i + n_i \frac{\partial v_i}{\partial u}$ 
In parallel within each level
13: for link  $i = n : 1$  do
14:   $\frac{\partial f_{\lambda_i}}{\partial u} += {}^i X_{\lambda_i}^T \frac{\partial f_i}{\partial u} + \zeta_i$ 
15: for link  $i = n : 1$  in parallel do
16:   $\frac{\partial c_i}{\partial u} = S_i^T \frac{\partial f_i}{\partial u}$ 
```

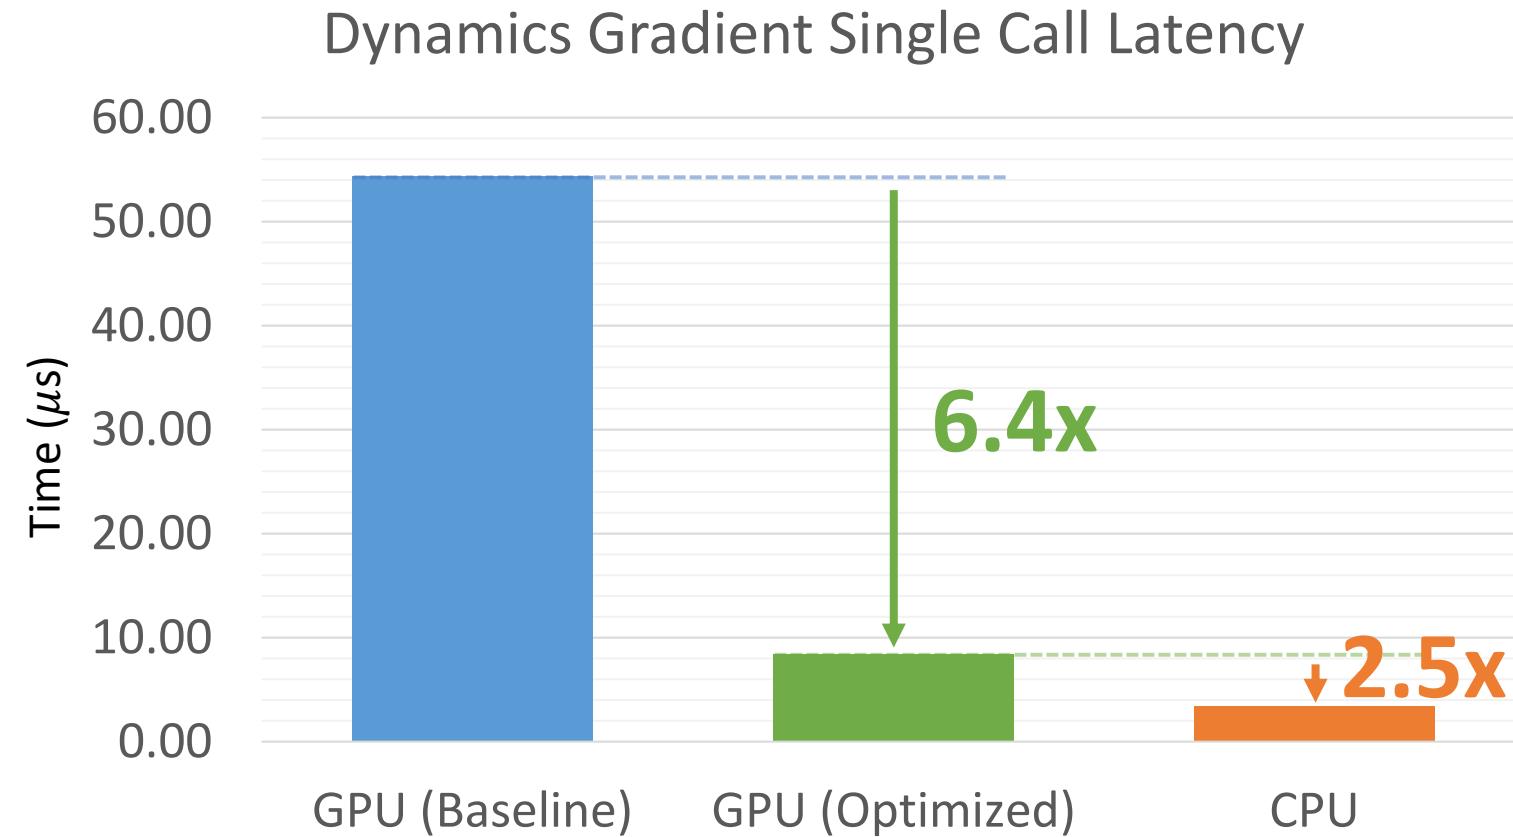
Refactoring improved single computation latency



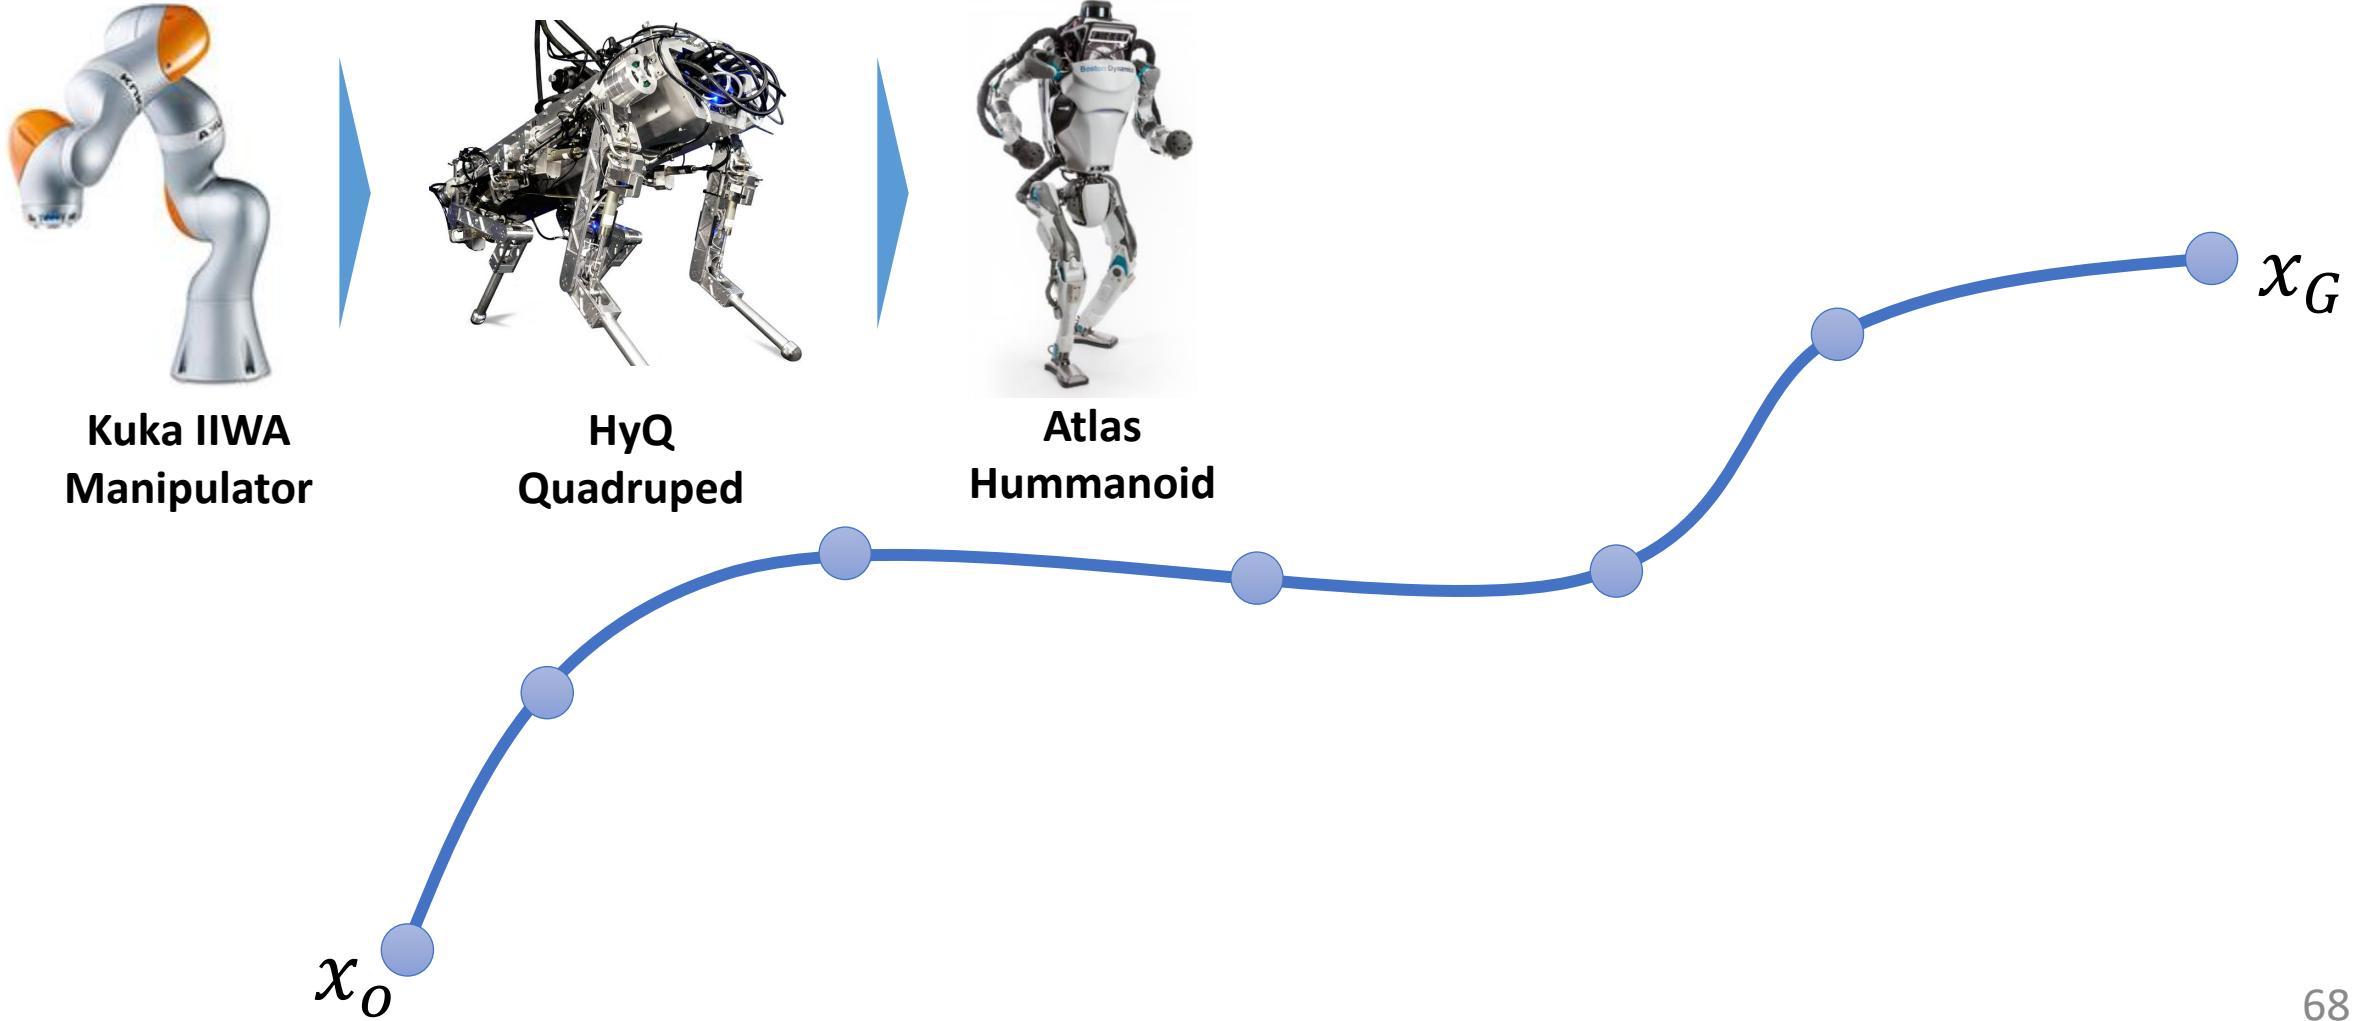
Refactoring improved single computation latency



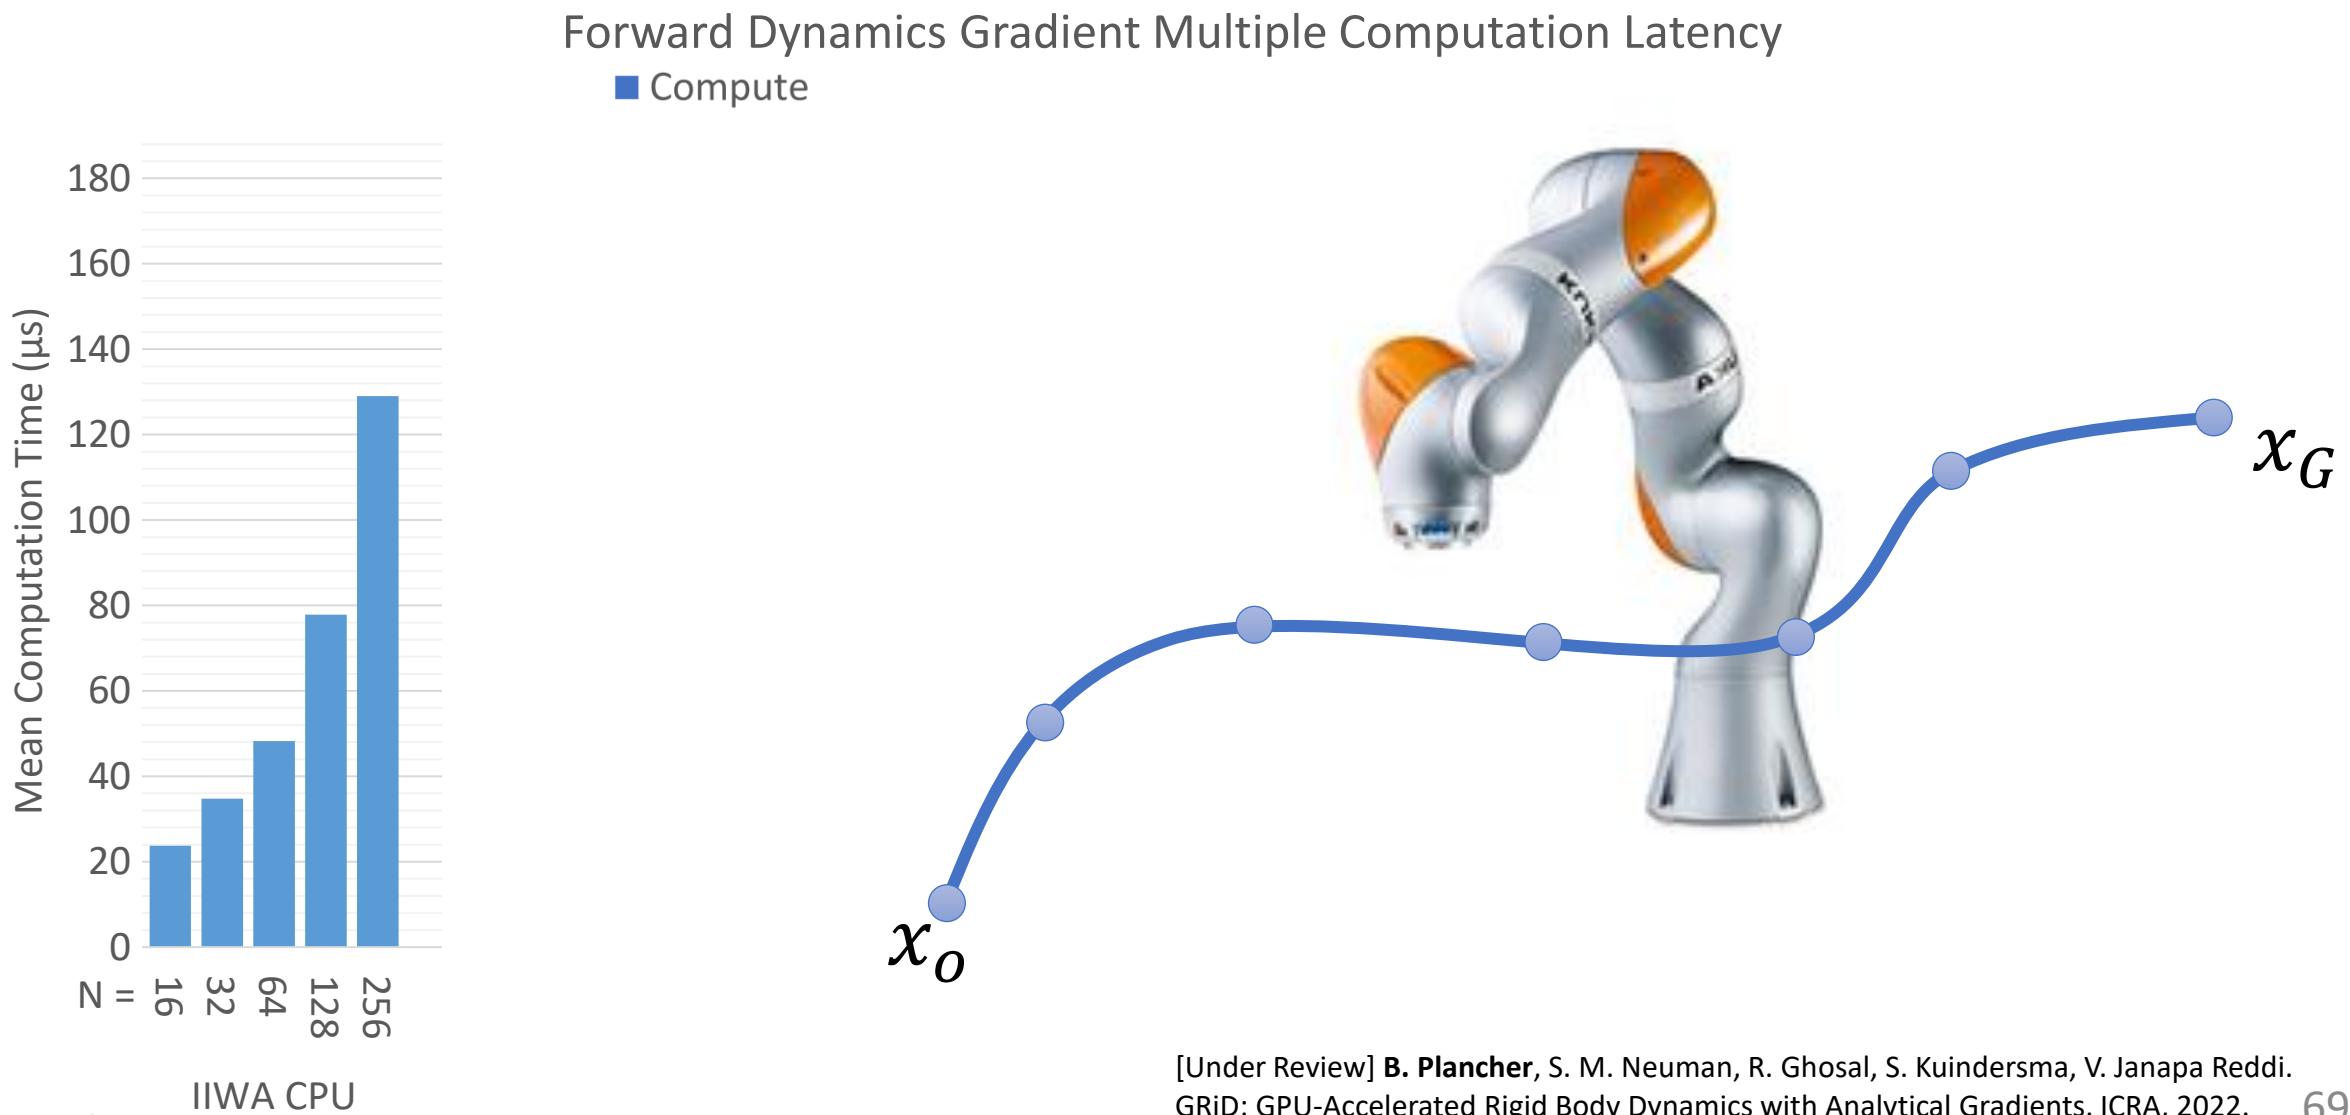
Refactoring improved single computation latency



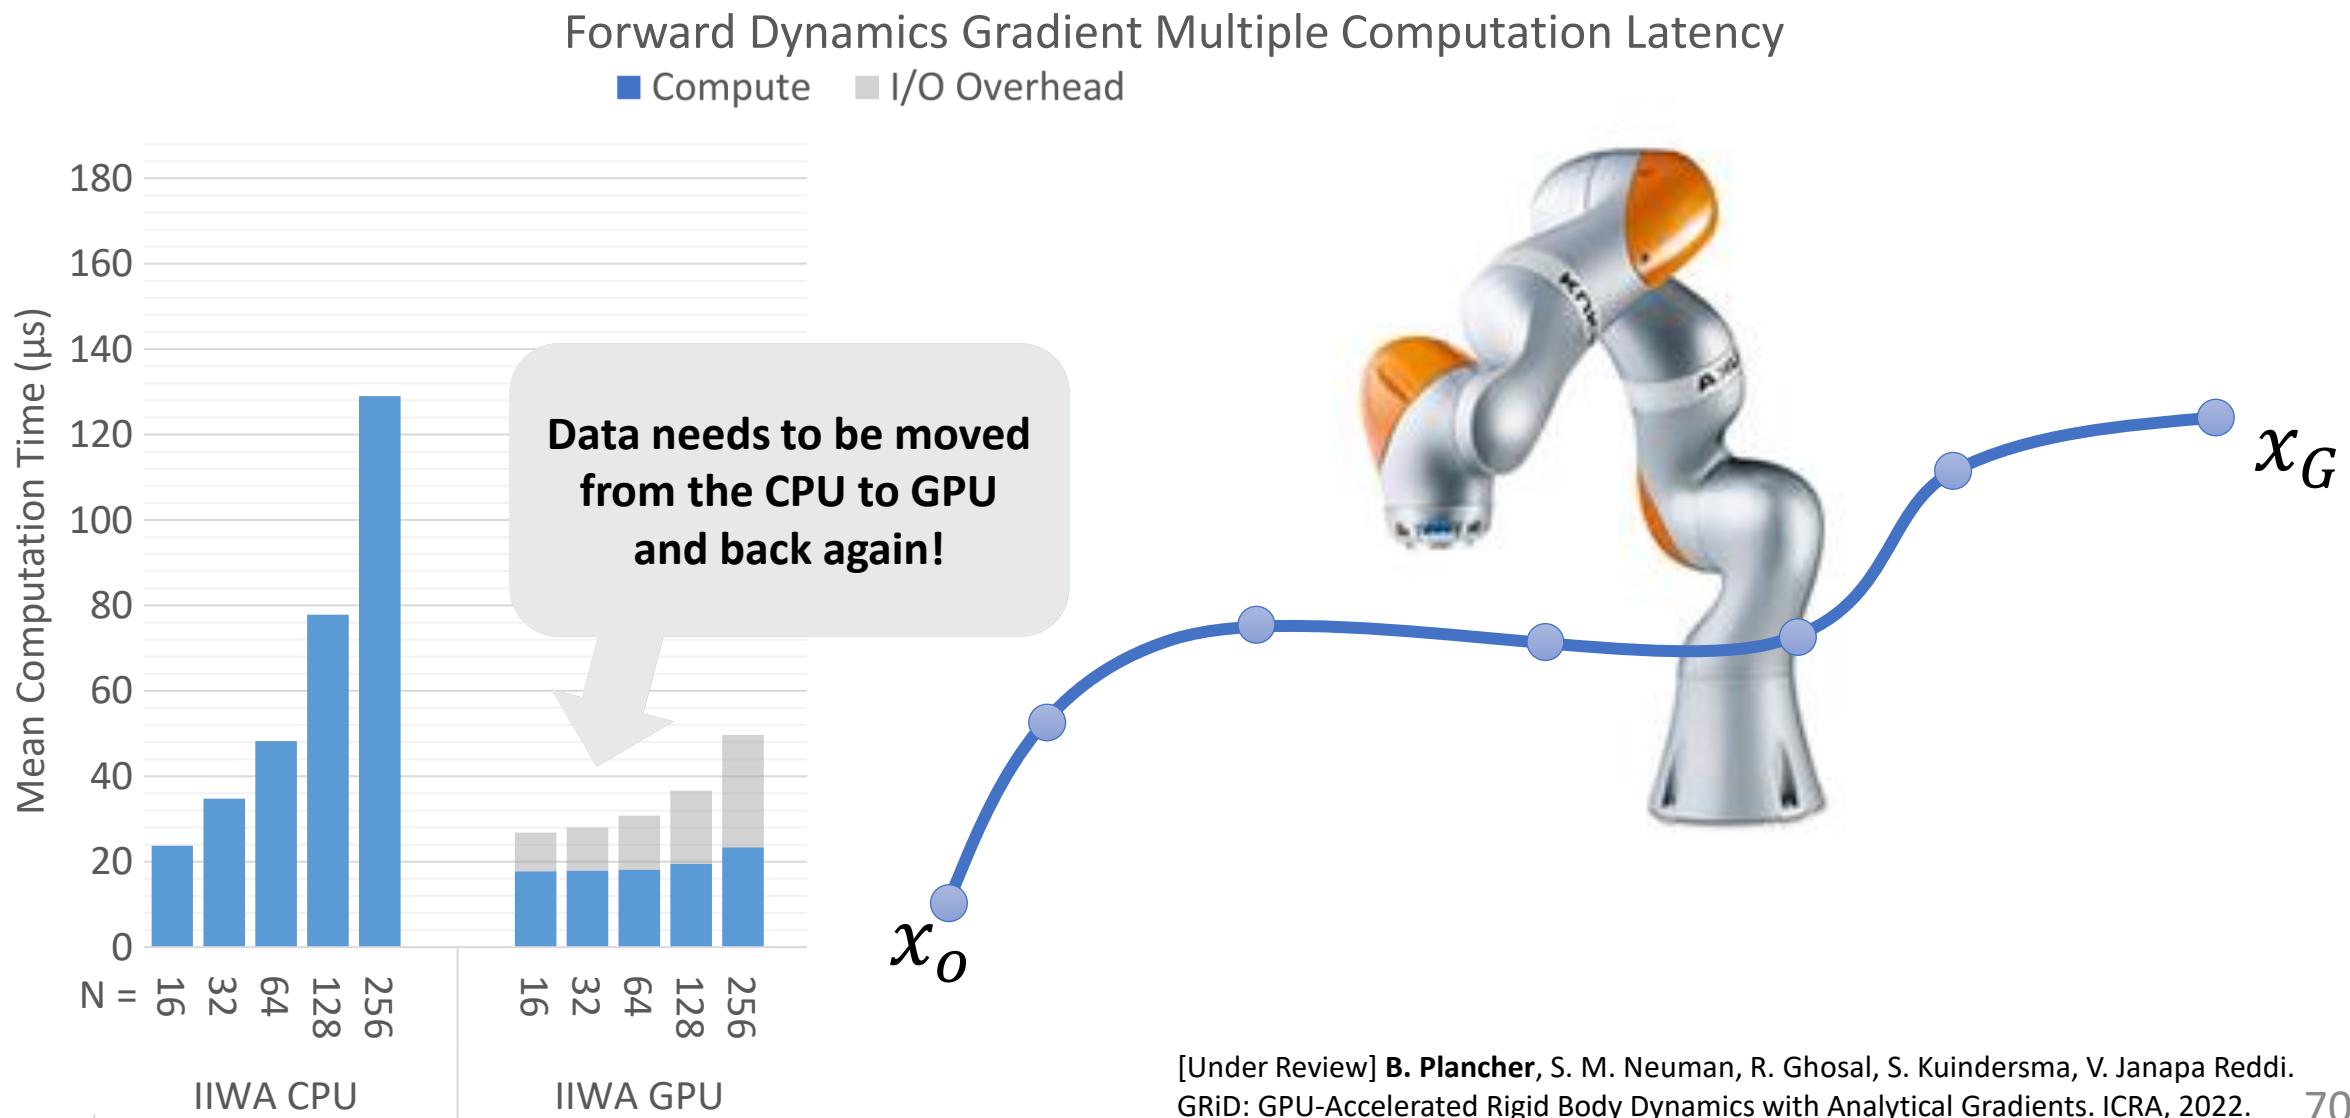
What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



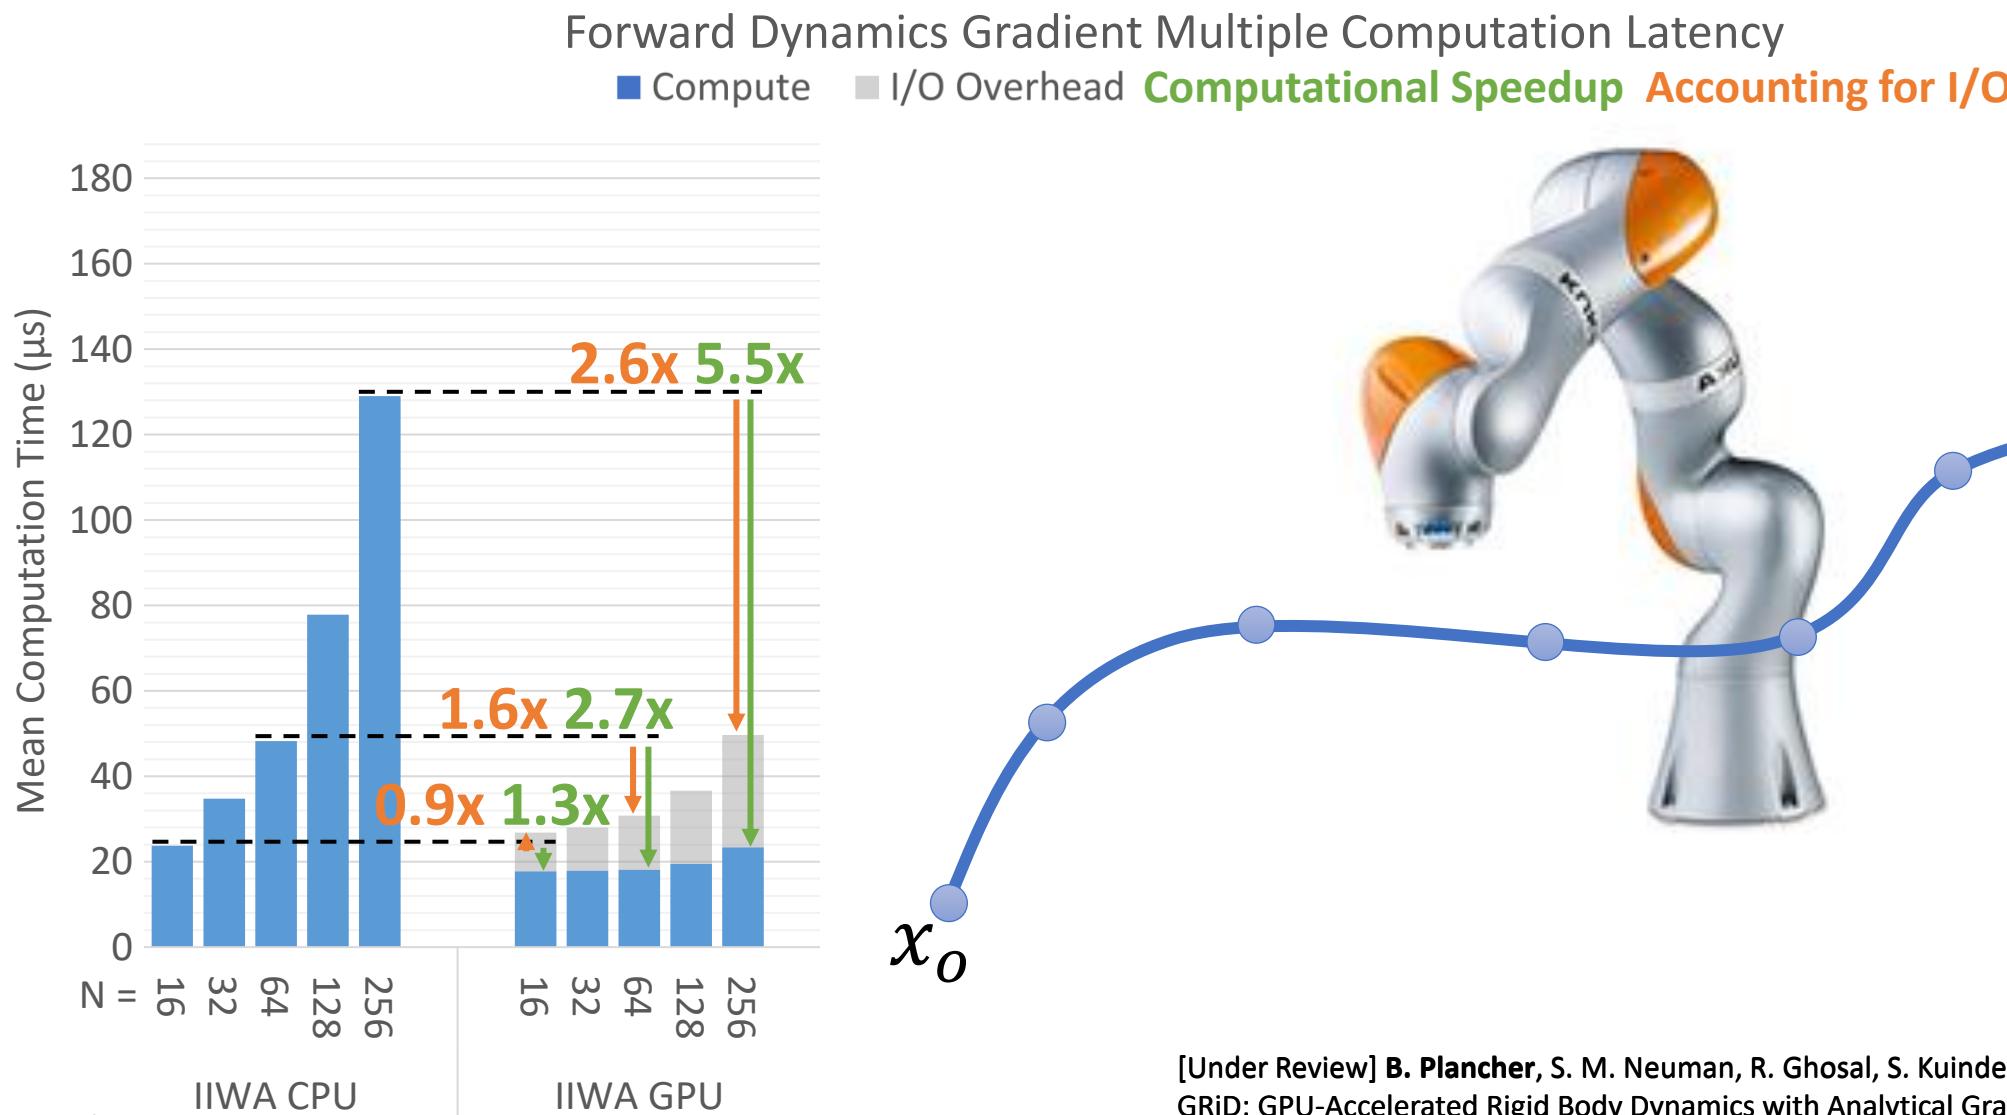
What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



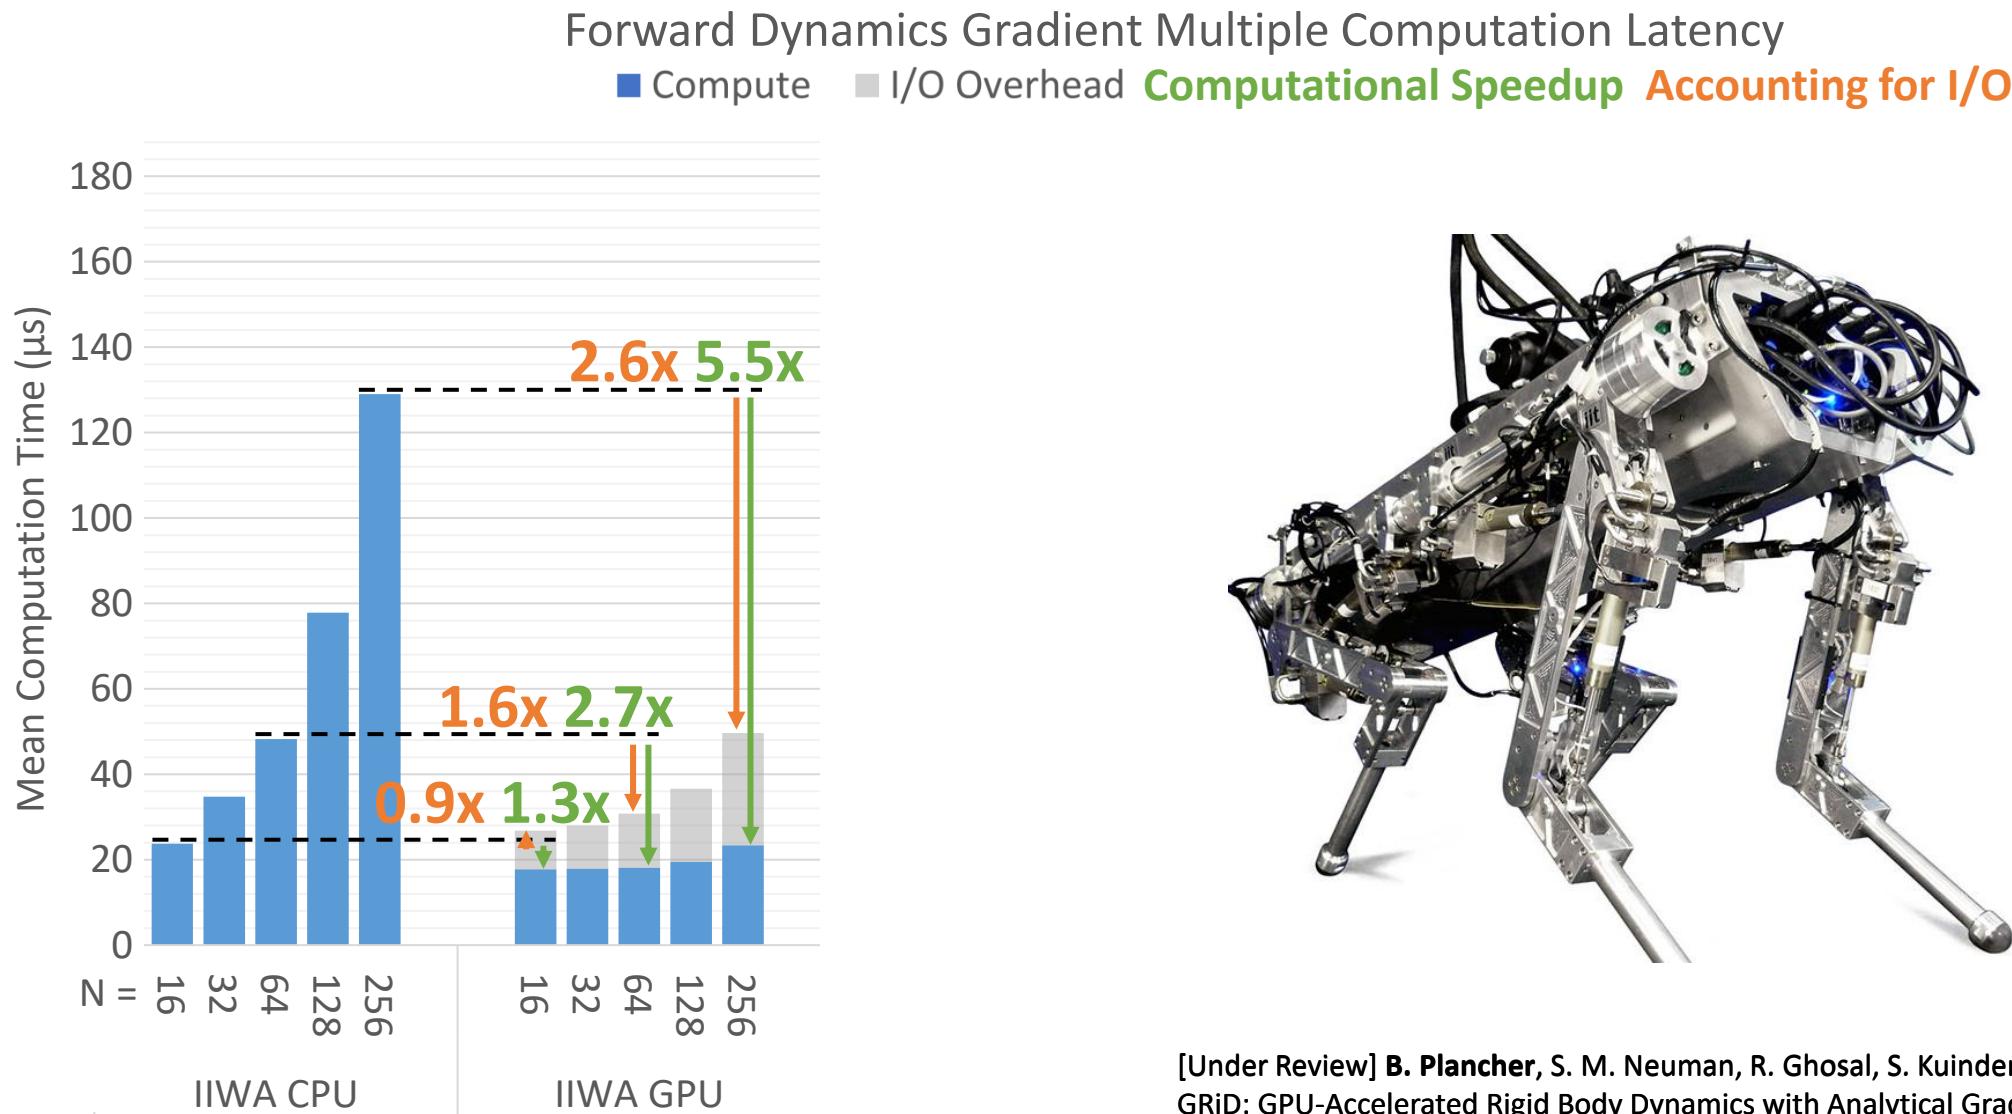
What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



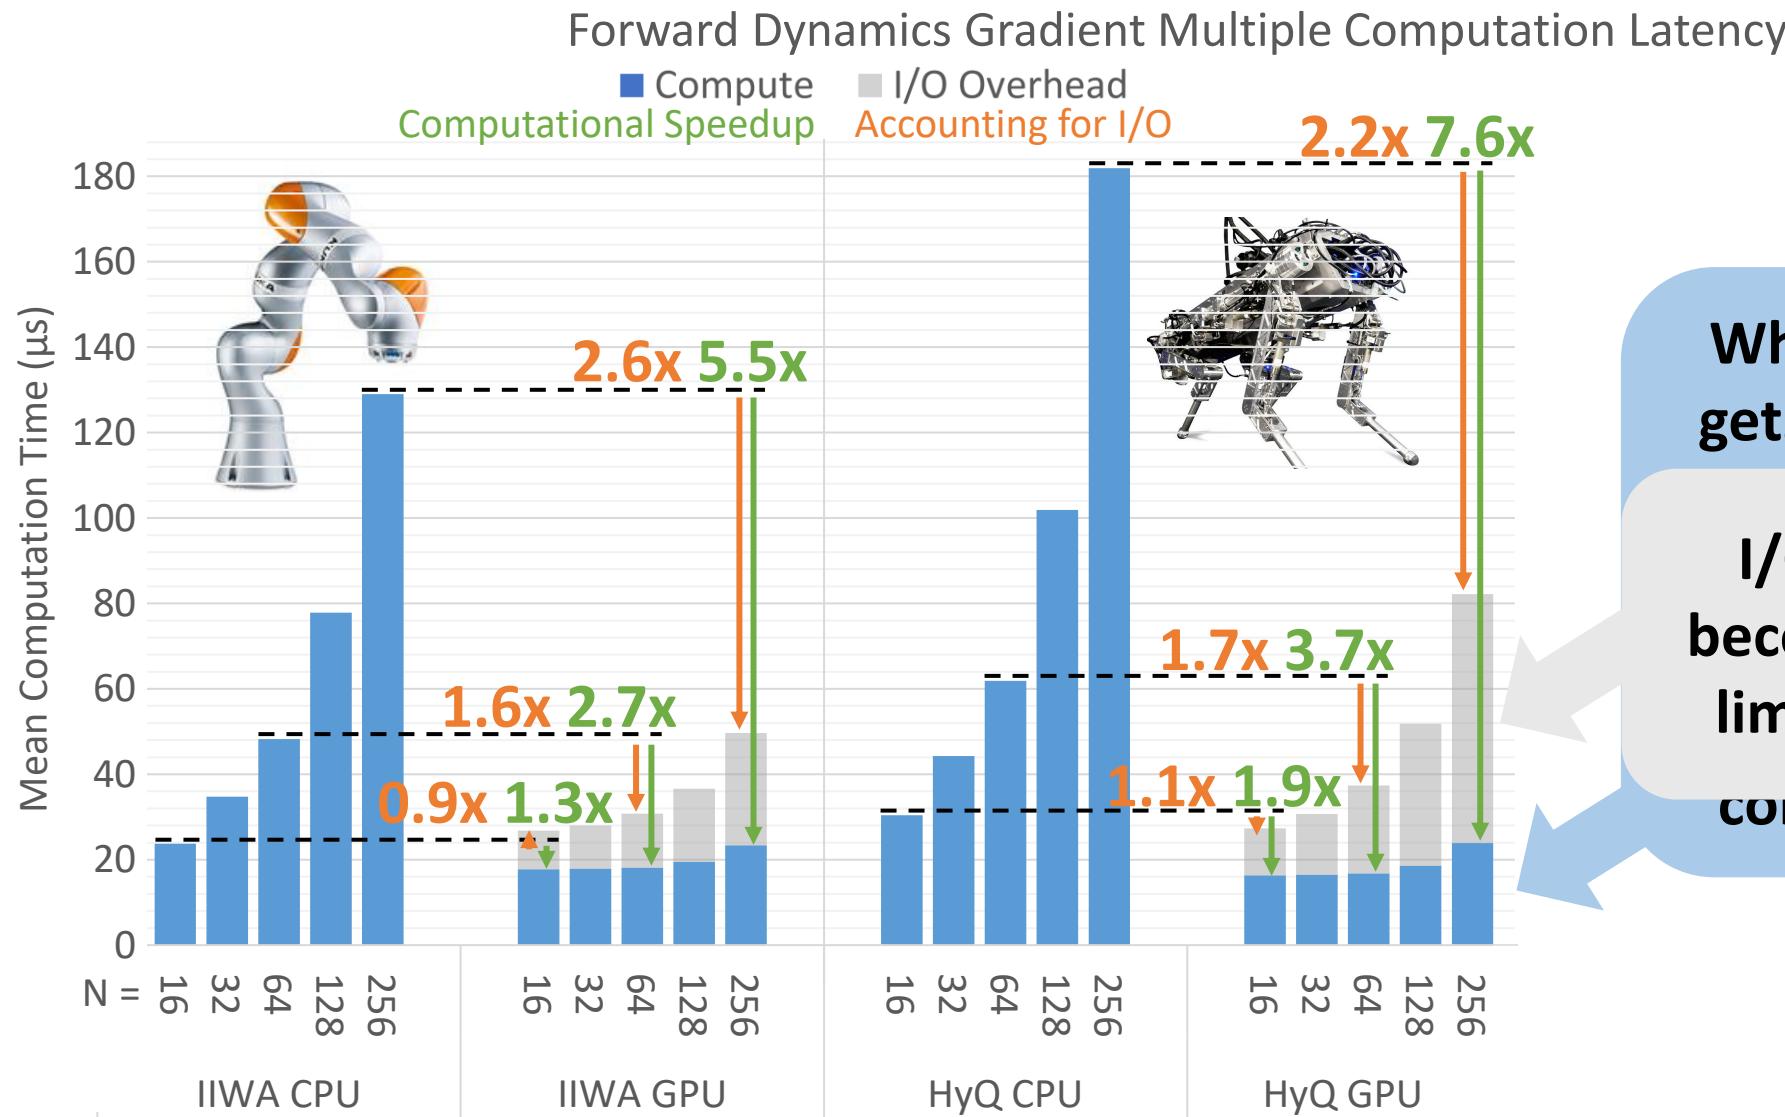
What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



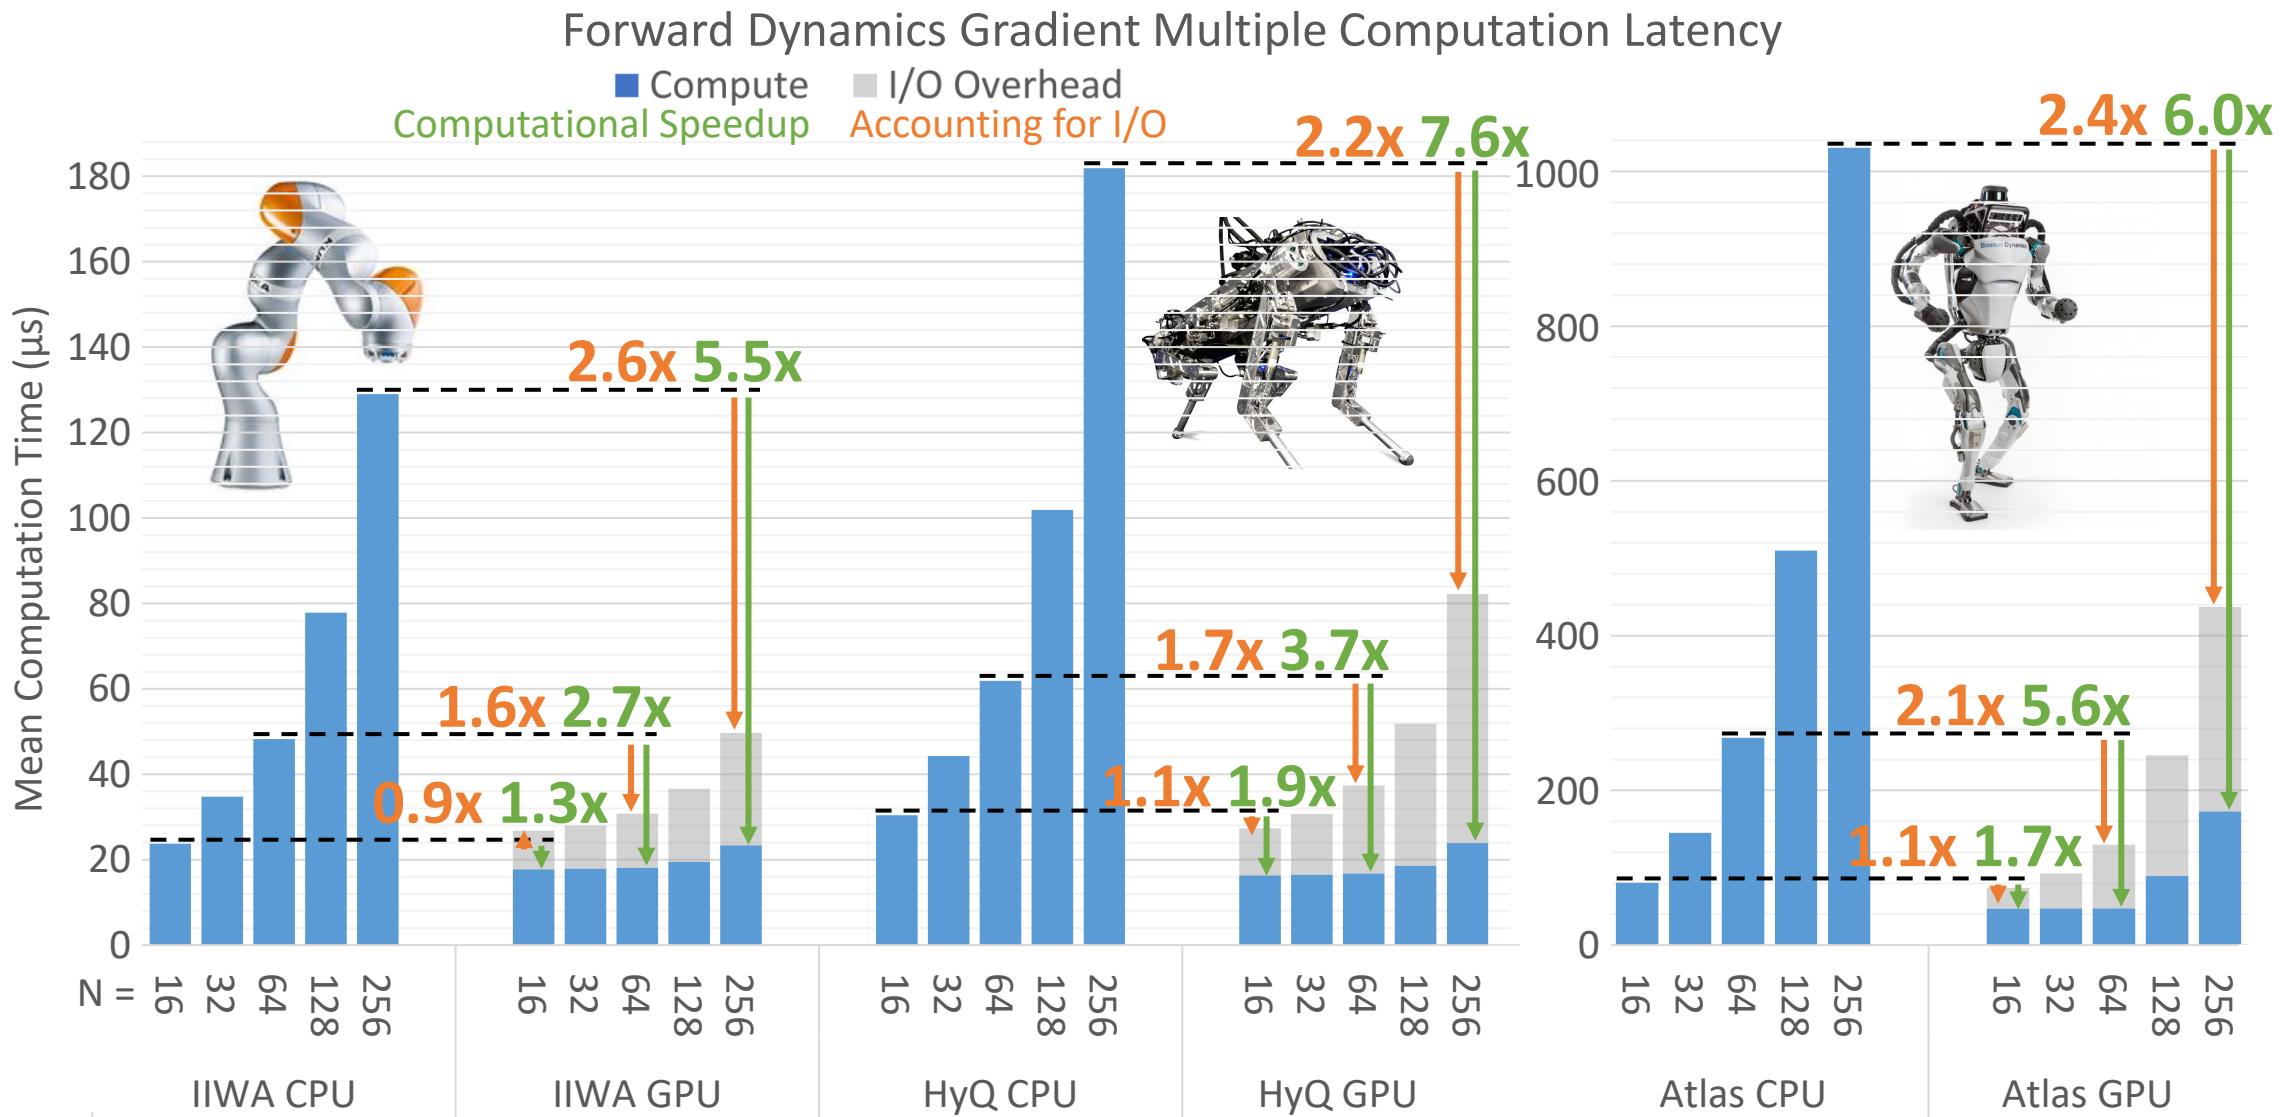
What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



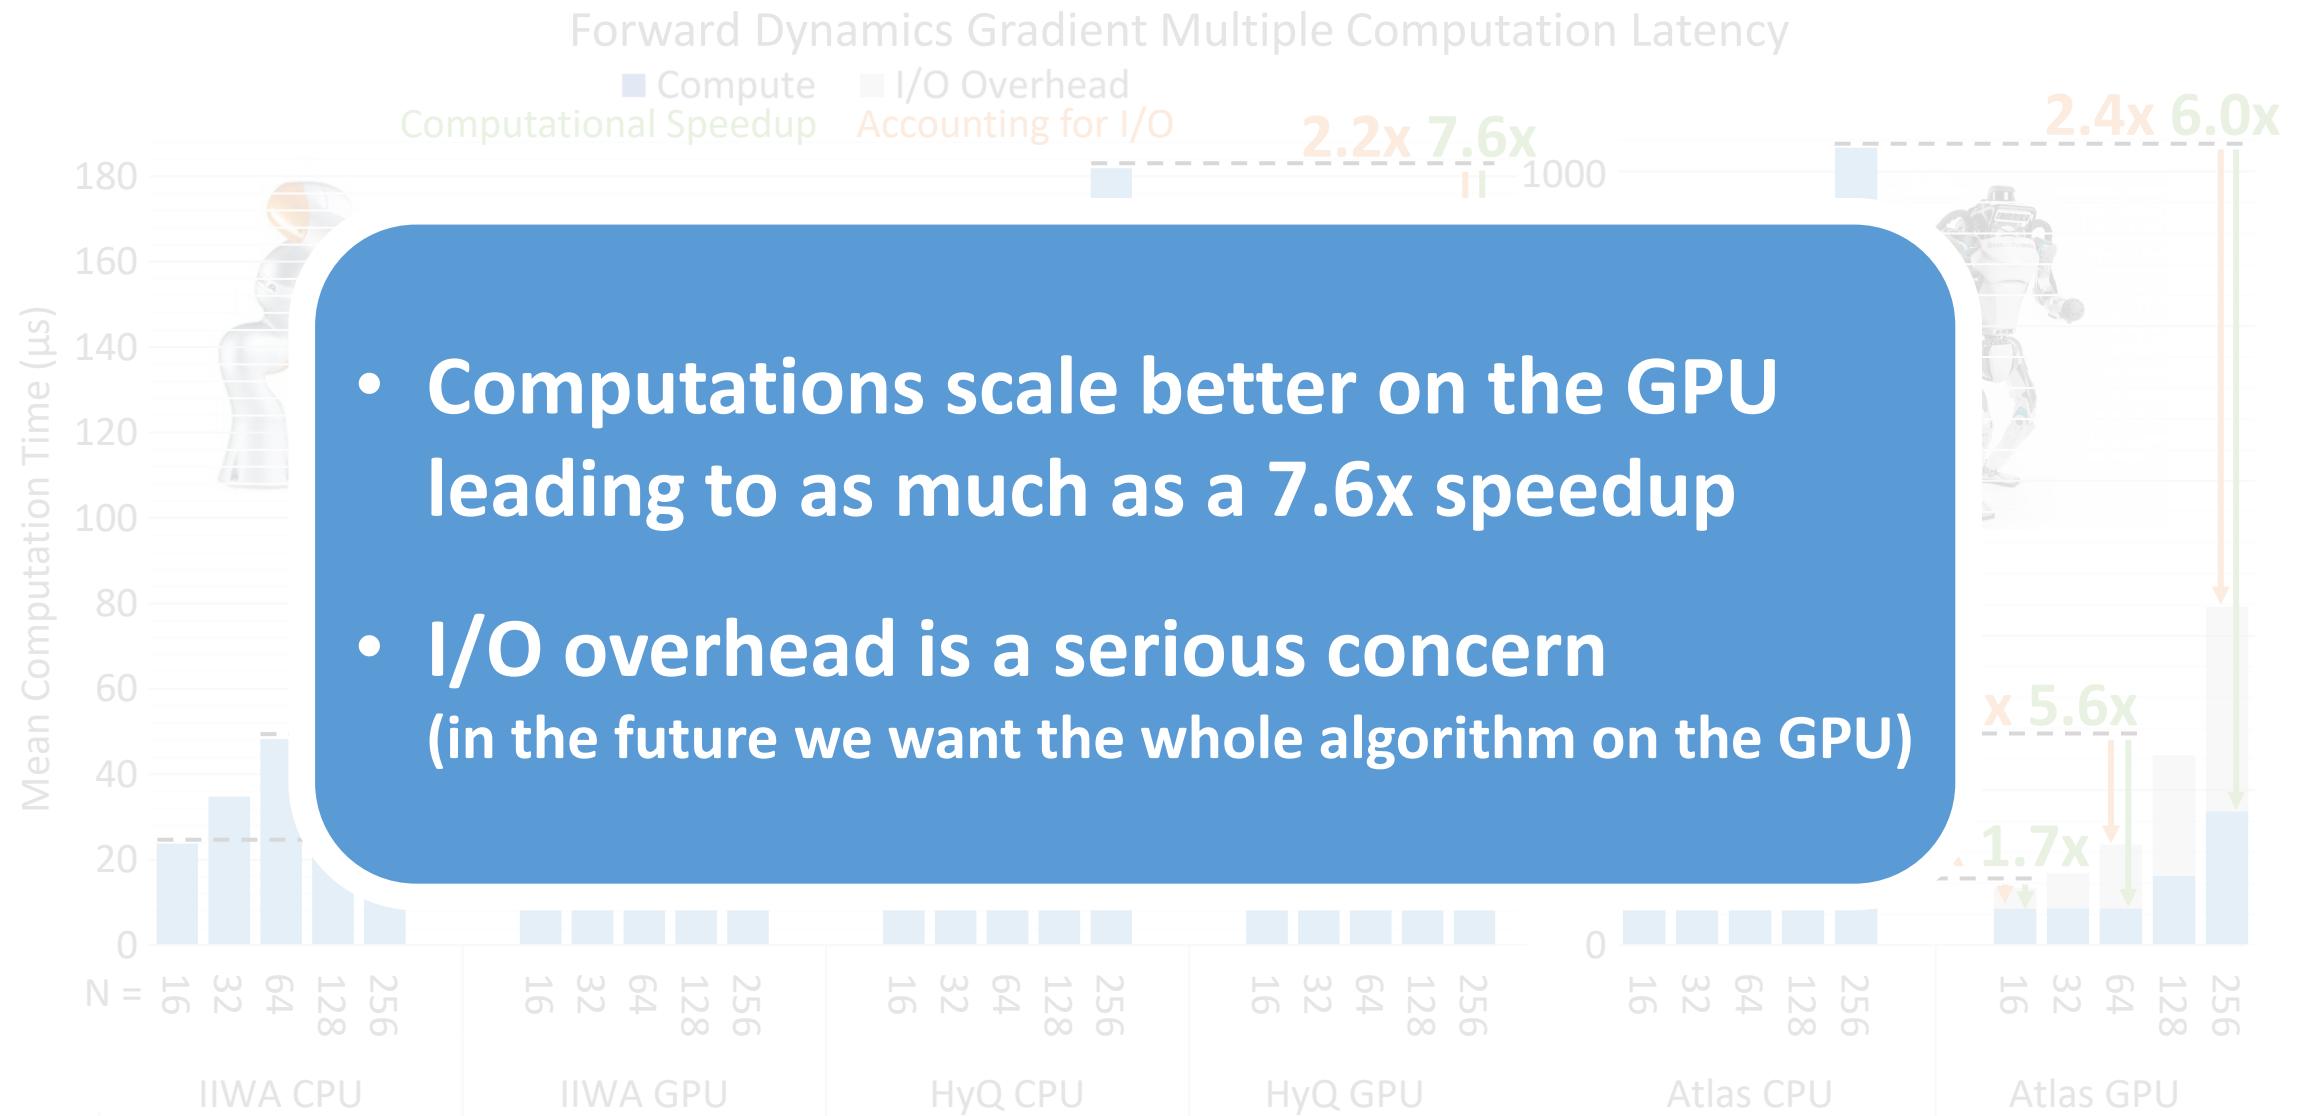
What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



What happens when we Scale Robot Complexity and consider Multiple Simultaneous Computations?



Hardware Acceleration for Realtime Robotics



A Quick Overview of (And My Approach To) Robotics

A Crash Course in (GPU) Architecture

A Case Study in Accelerating Rigid Body Dynamics

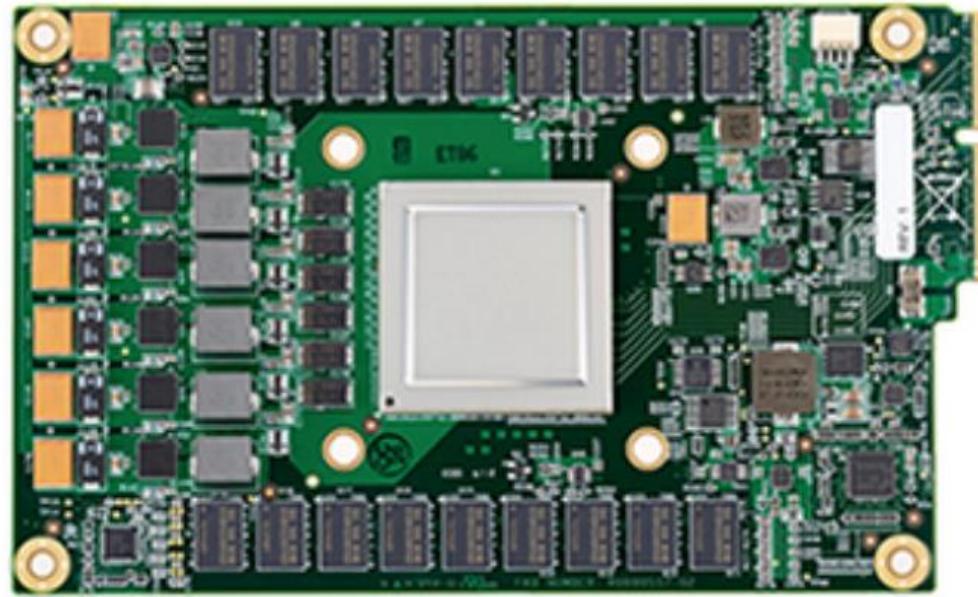
Why Accelerate Rigid Body Dynamics?

Accelerating Rigid Body Dynamics on the GPU

Moving Beyond the GPU

The Future of Accelerating Robotics

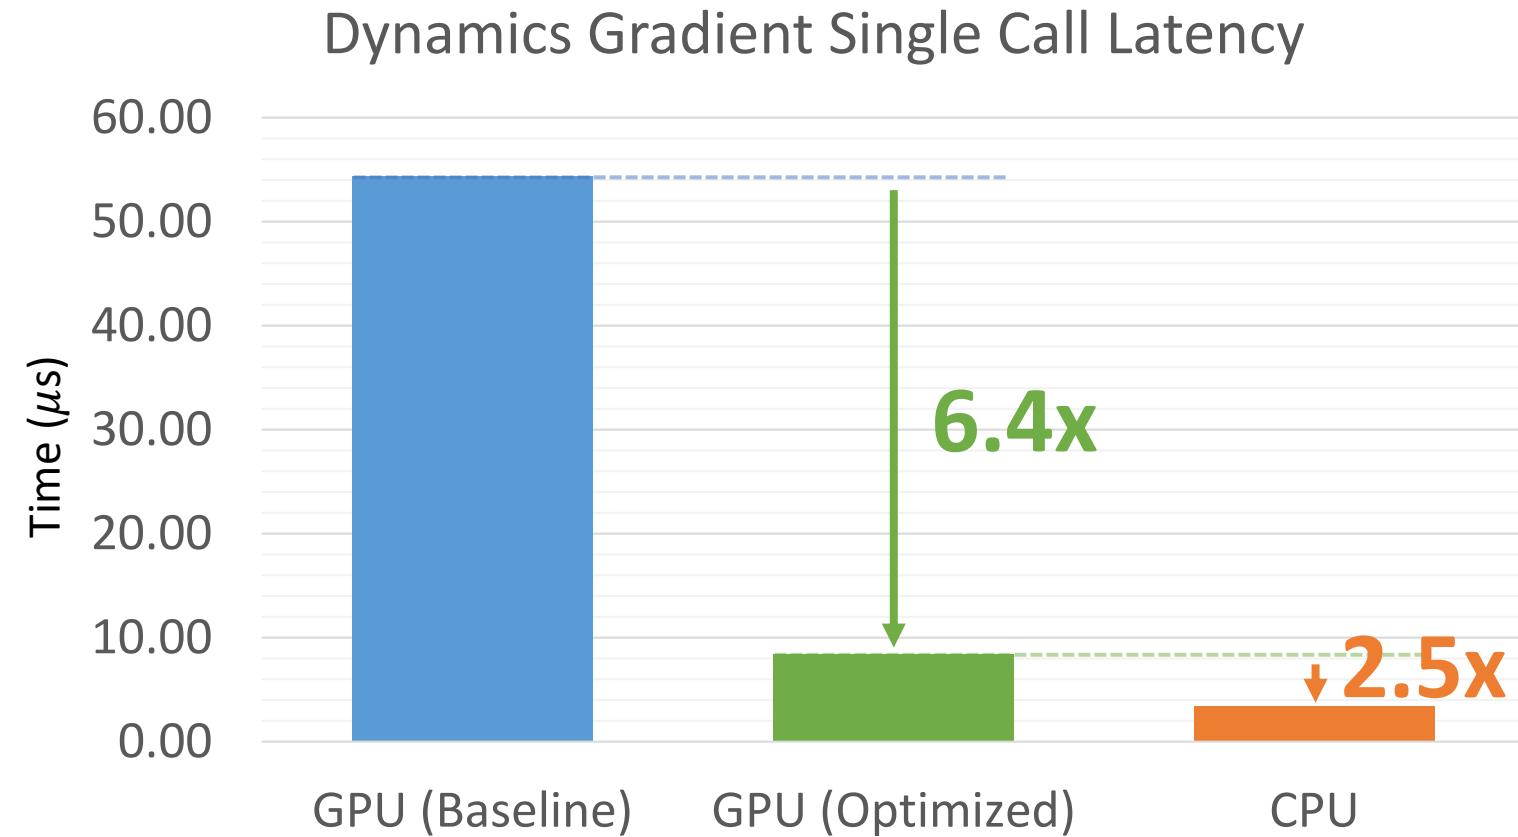
Custom Hardware Can Provide More Acceleration



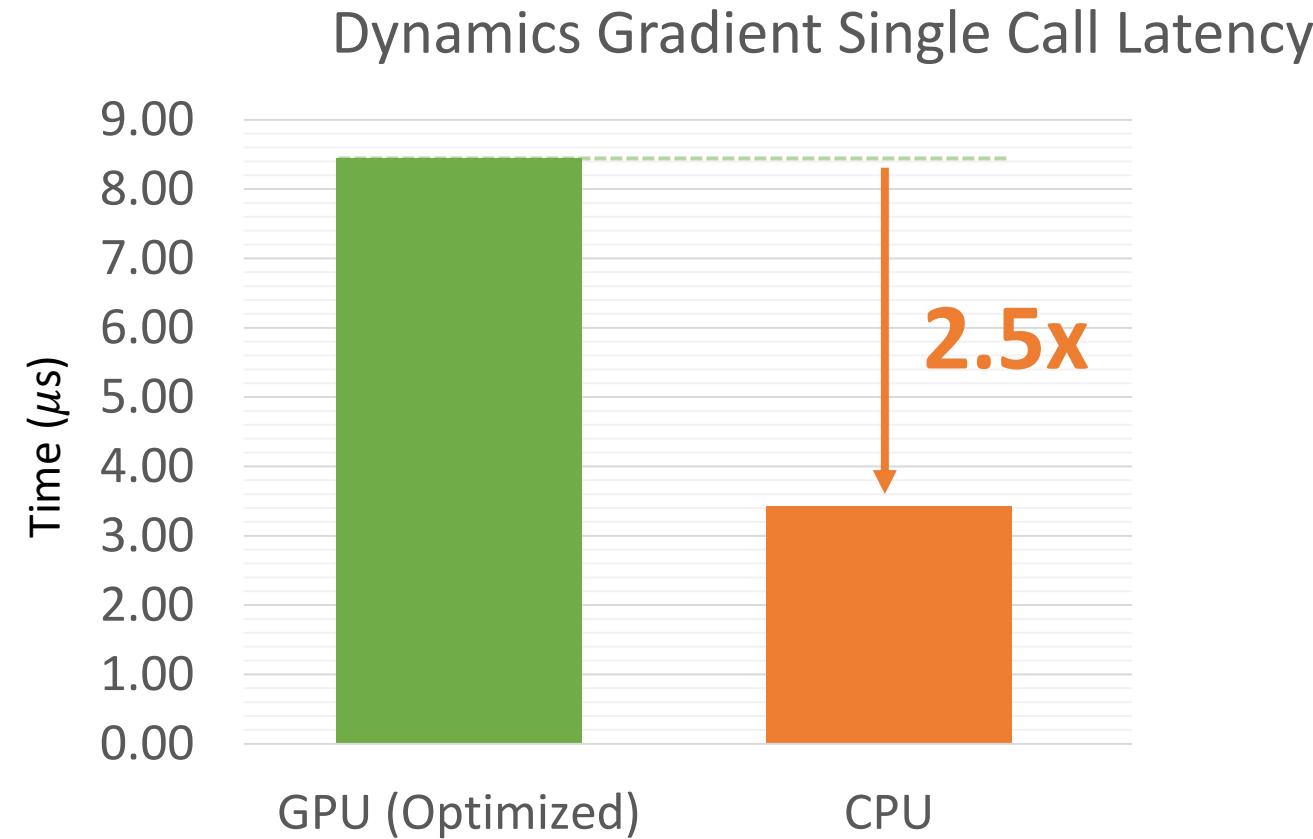
Google's first Tensor Processing Unit (TPU) on a printed circuit board (left); TPUs deployed in a Google datacenter
(right)

<https://cloud.google.com/blog/products/ai-machine-learning/an-in-depth-look-at-googles-first-tensor-processing-unit-tpu>

Custom Hardware Can Provide More Acceleration



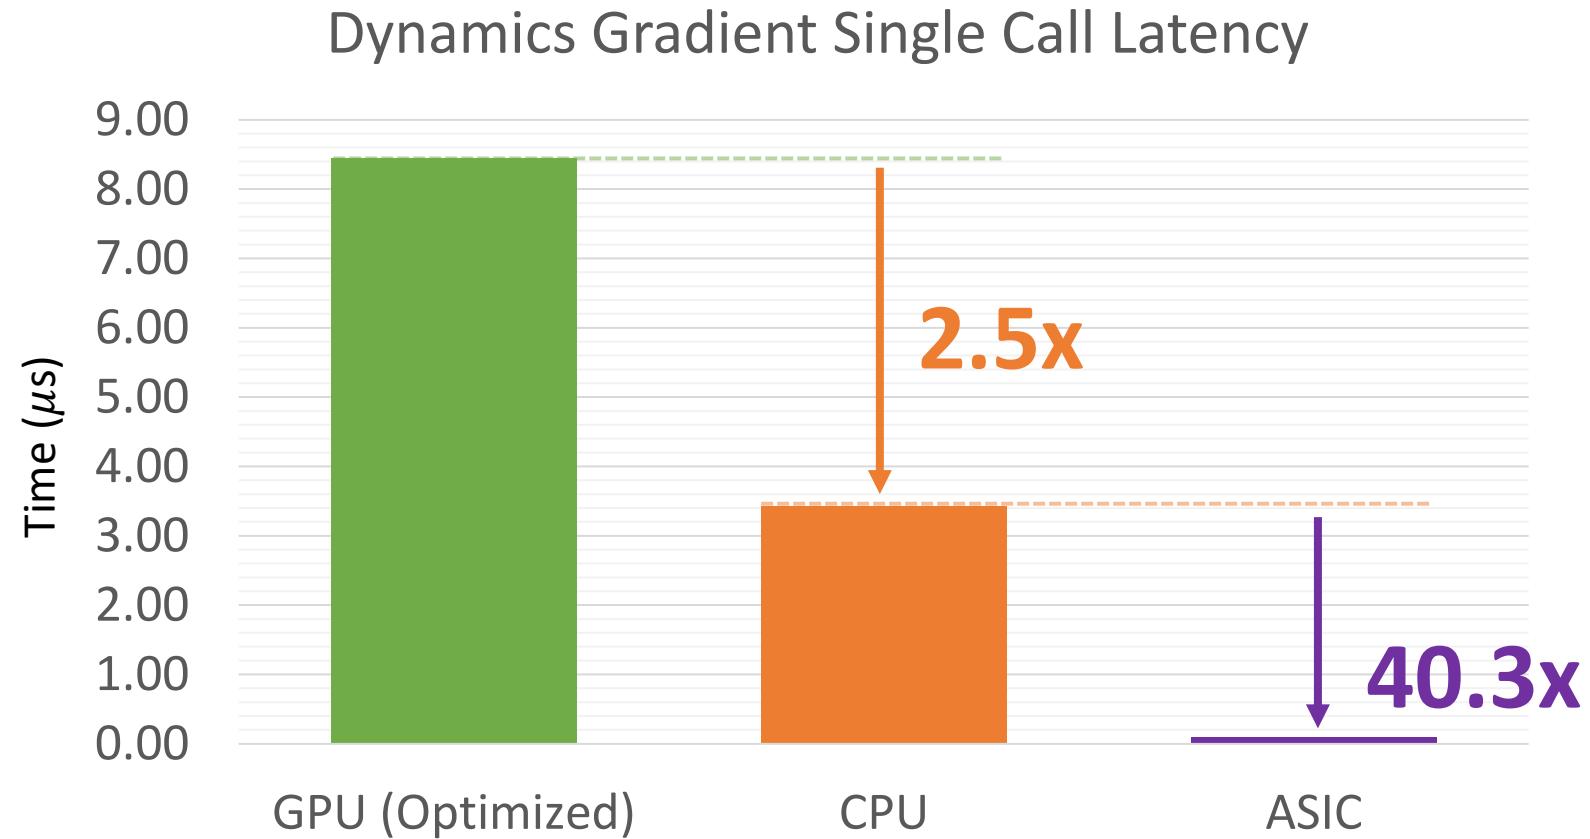
Custom Hardware Can Provide More Acceleration



B. Plancher, S. M. Neuman, T. Bourgeat, S. Kuindersma, S. Devadas, V. Janapa Reddi.
Accelerating Robot Dynamics Gradients on a CPU, GPU, and FPGA. RA-L and ICRA 2021.

S. M. Neuman, B. Plancher, T. Bourgeat, T. Tambe, S. Devadas, V. Janapa Reddi. Robomorphic Computing: A Design Methodology for Domain-Specific Accelerators Parameterized by Robot Morphology. ASPLOS 2021.

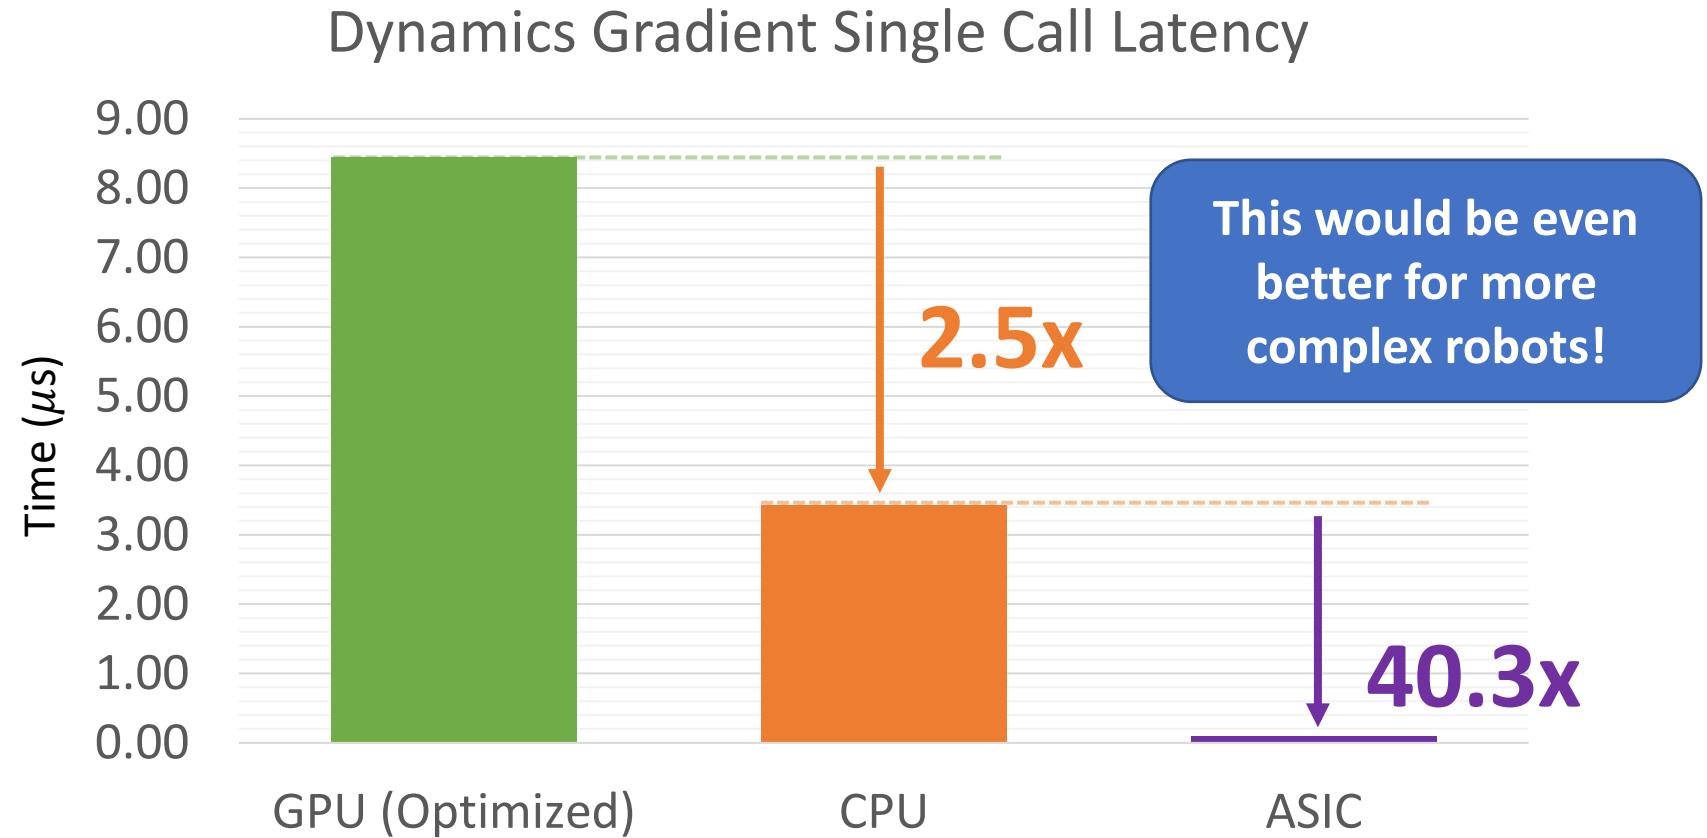
Custom Hardware Can Provide More Acceleration



B. Plancher, S. M. Neuman, T. Bourgeat, S. Kuindersma, S. Devadas, V. Janapa Reddi.
Accelerating Robot Dynamics Gradients on a CPU, GPU, and FPGA. RA-L and ICRA 2021.

S. M. Neuman, B. Plancher, T. Bourgeat, T. Tambe, S. Devadas, V. Janapa Reddi. Robomorphic Computing: A Design Methodology for Domain-Specific Accelerators Parameterized by Robot Morphology. ASPLOS 2021.

Custom Hardware Can Provide More Acceleration



B. Plancher, S. M. Neuman, T. Bourgeat, S. Kuindersma, S. Devadas, V. Janapa Reddi.
Accelerating Robot Dynamics Gradients on a CPU, GPU, and FPGA. RA-L and ICRA 2021.

S. M. Neuman, B. Plancher, T. Bourgeat, T. Tambe, S. Devadas, V. Janapa Reddi. Robomorphic Computing: A Design Methodology for Domain-Specific Accelerators Parameterized by Robot Morphology. ASPLOS 2021.

Hardware Acceleration for Realtime Robotics



A Quick Overview of (And My Approach To) Robotics

A Crash Course in (GPU) Architecture

A Case Study in Accelerating Rigid Body Dynamics

The Future of Accelerating Robotics

Enabling Robotics Research on High-Performance Parallel Architectures

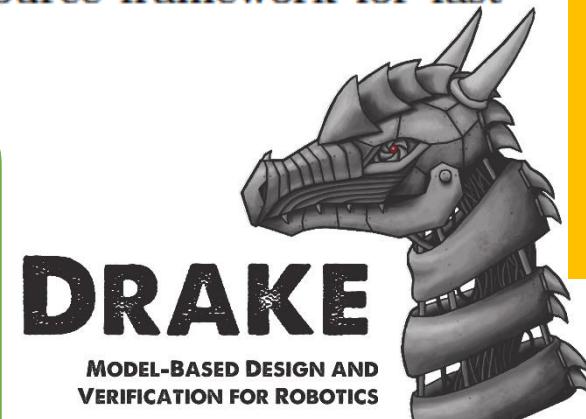


MuJoCo

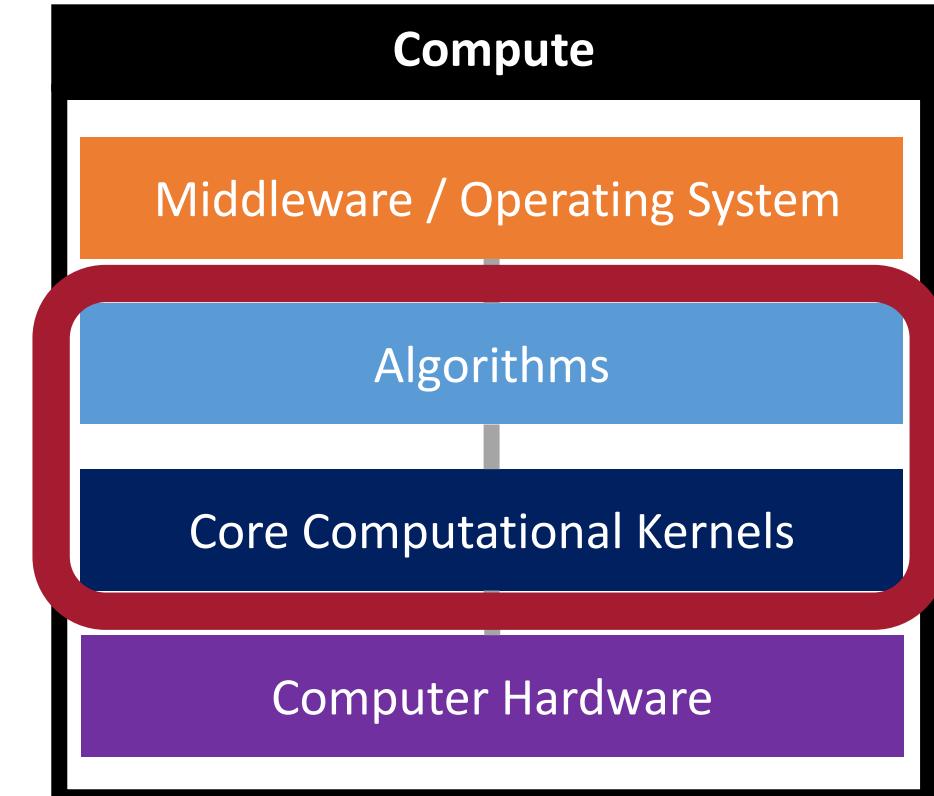


acados – a modular open-source framework for fast embedded optimal control

Only Forward Dynamics On GPU

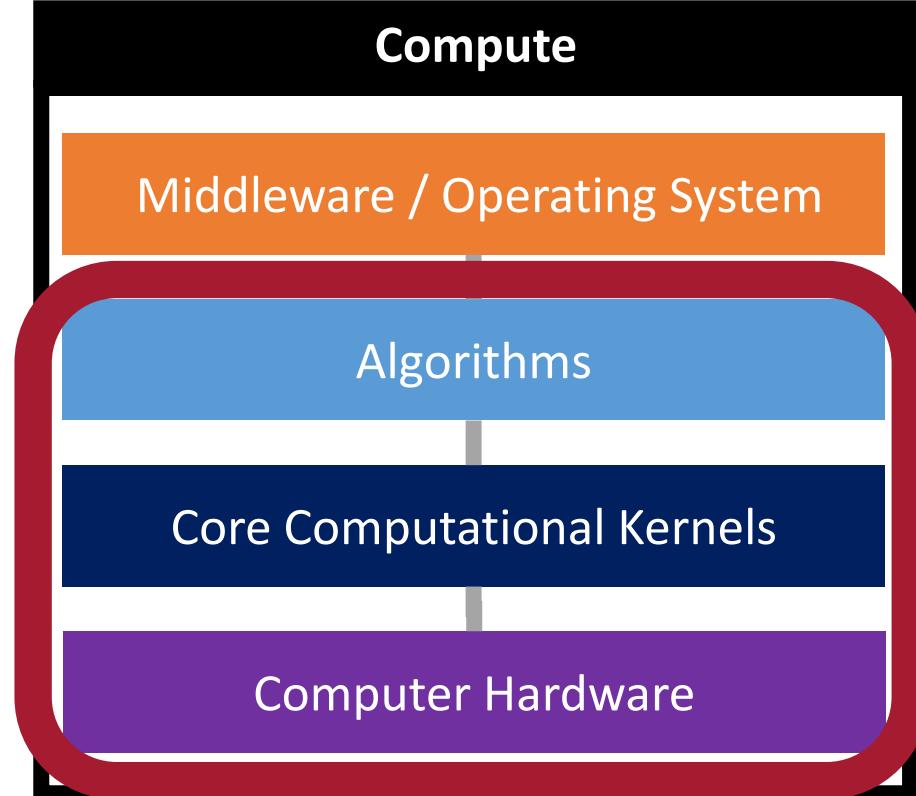
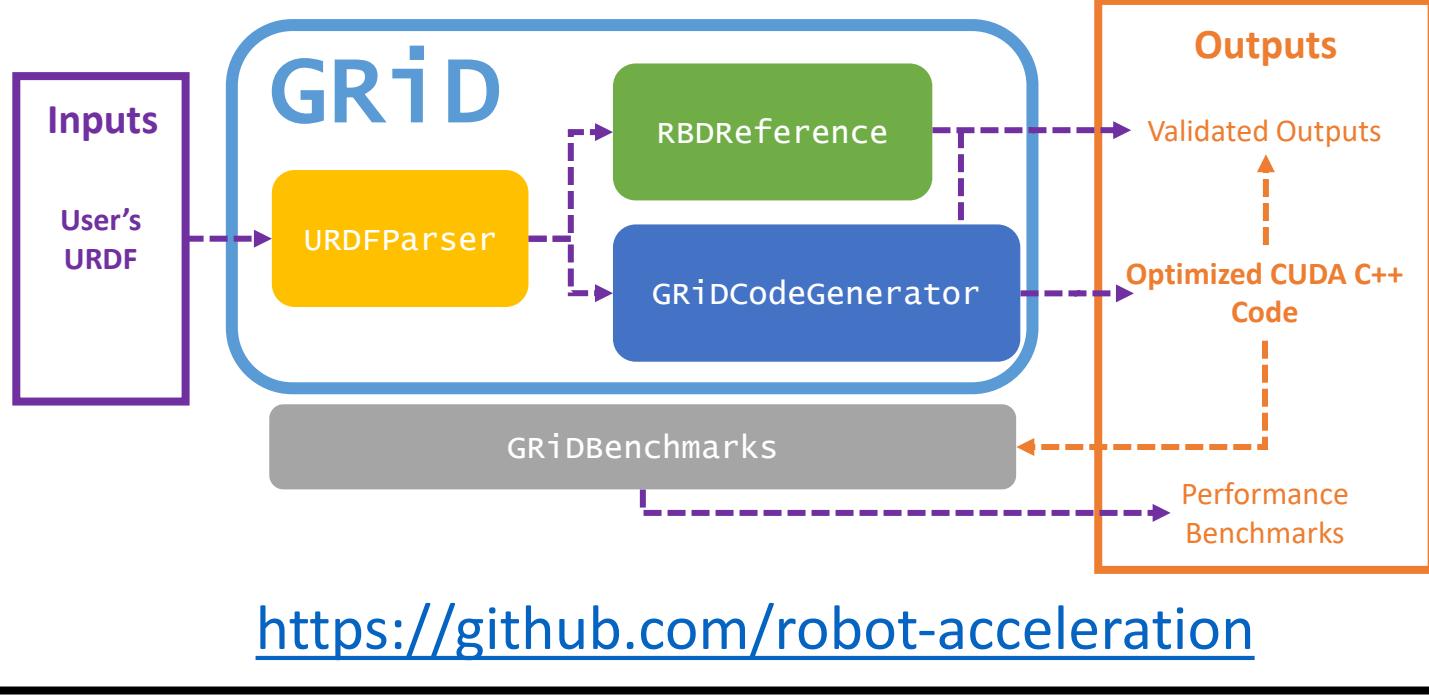


CPU
ONLY!

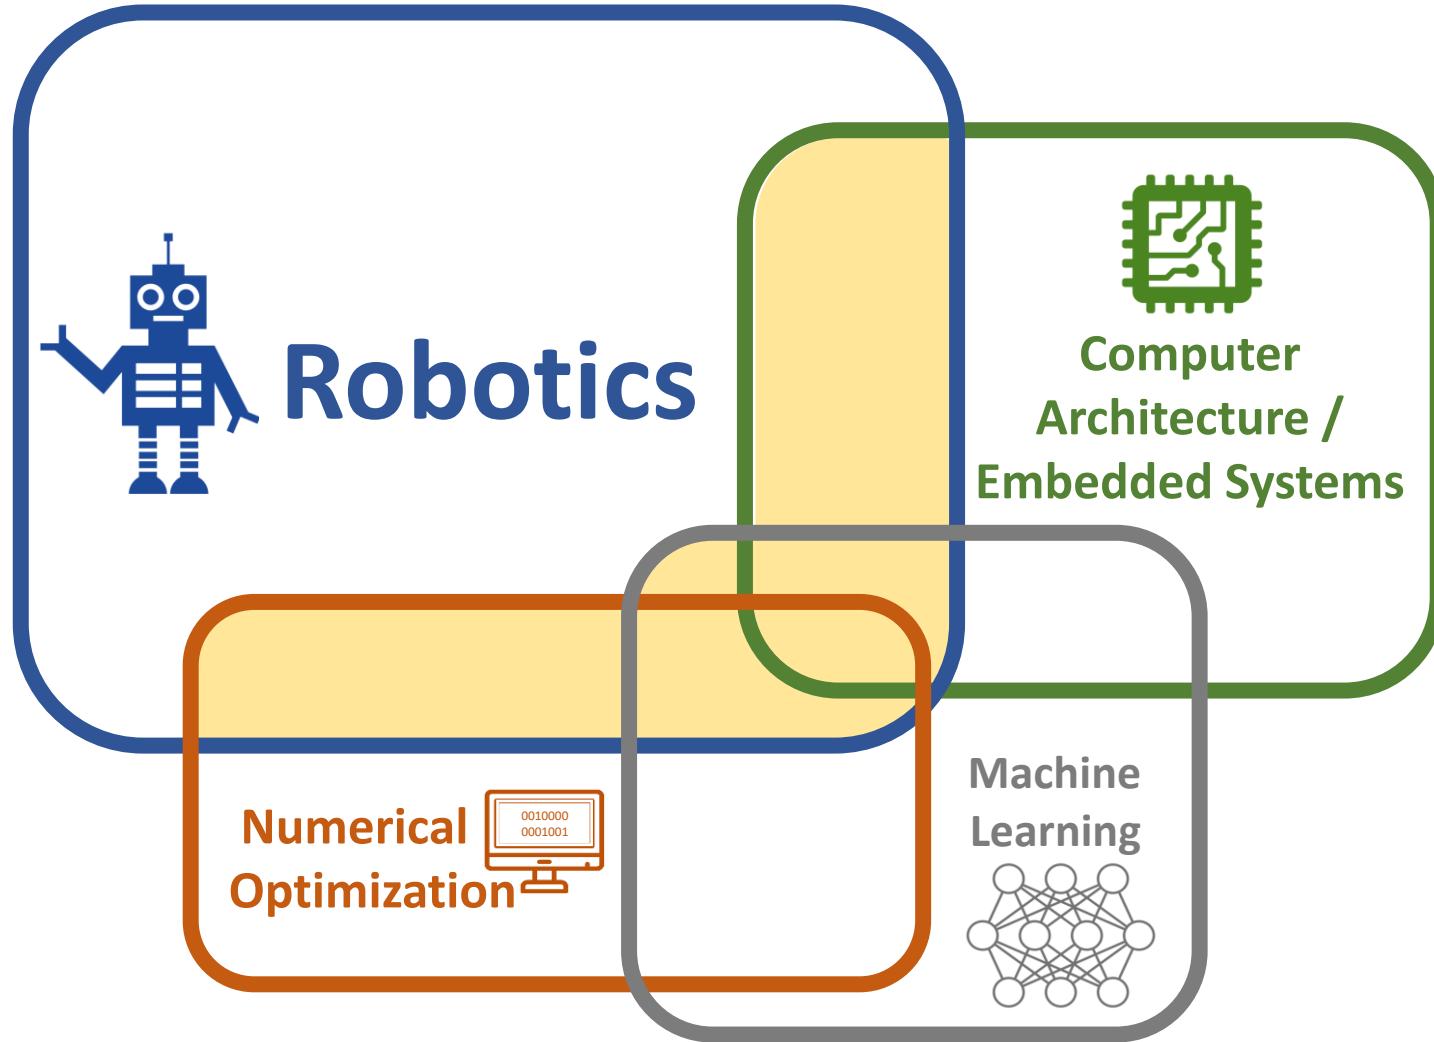


Enabling Robotics Research on High-Performance Parallel Architectures

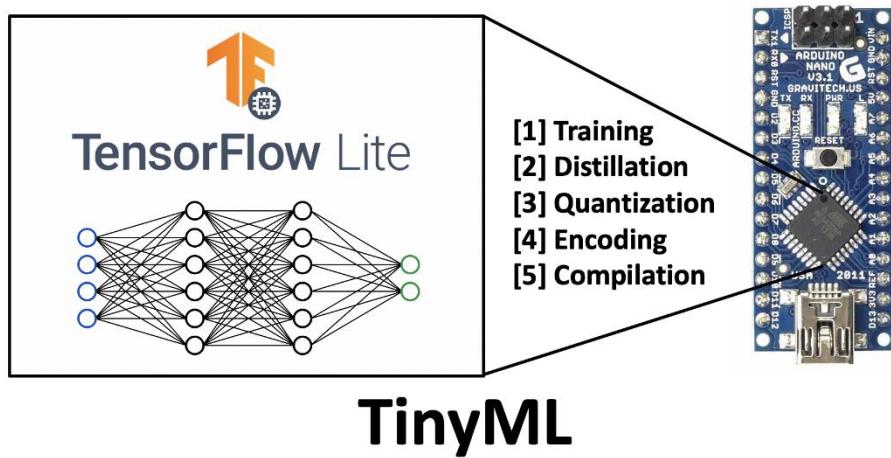
GRiD: GPU Accelerated Rigid Body Dynamics



I work at the intersection of robotics and adjacent fields

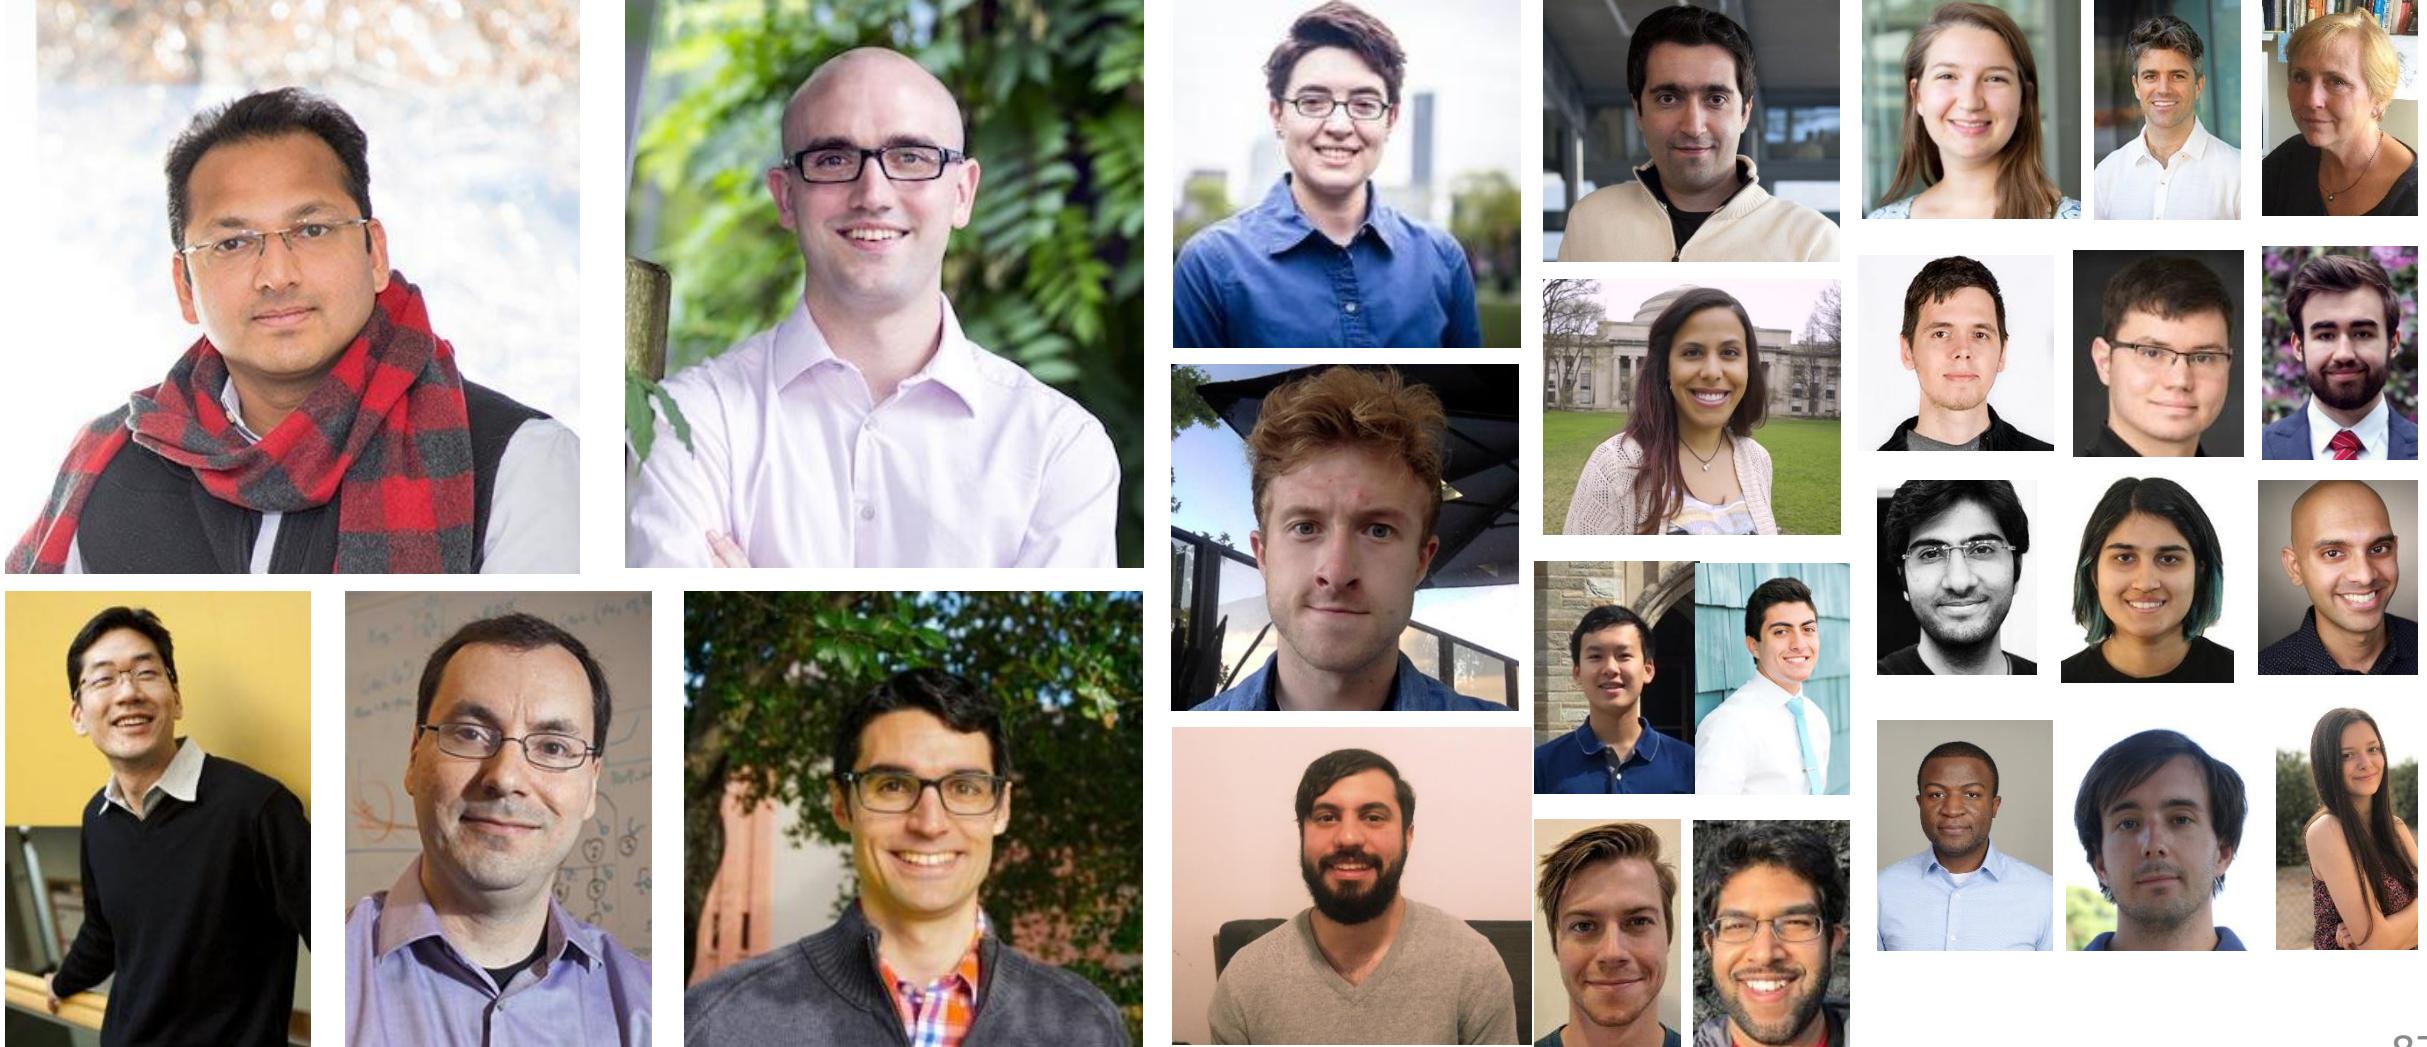


Tiny Robotics: Lowering the Cost of Robotics for Widespread Deployment



Thank You to My Funders and Collaborators

This material is based upon work supported by the National Science Foundation (under Grant DGE1745303 and Grant 2030859 to the Computing Research Association for the CIFellows Project), and the Defense Advanced Research Projects Agency (under Grant HR001118C0018). Any opinions, findings, conclusions, or recommendations expressed in this material are those of the authors and do not necessarily reflect those of the funding organizations



Hardware Acceleration for Realtime Robotics



<https://brianplancher.com>

brian_plancher@g.harvard.edu

1. Learn some of the **language of robotics**
(and computer architecture)
2. Understand the important of **parallelism**
3. Gain practice in exploring the **process** and
opportunities of **hardware acceleration**

Feedback Link: <https://bit.ly/Brian-Barnard-21>